

# COMP3221

# Distributed Systems

## Blockchain Scalability

Presented by A/Prof. Vincent Gramoli



THE UNIVERSITY OF  
SYDNEY

CRCOS 00026A

<https://www.bankrate.com/investing/best-blockchain-etfs/>

USER NOTE:  
It is optional to include  
this additional slide.

The University of Sydney's Camperdown Campus sits on the lands of the Gadigal people with campuses, teaching and research facilities on the lands of the Gamaraygal, Dharug, Wangal, Darkinyung, Burramadagal, Dharawal, Gandangara, Gamilyaraay, Barkindji, Bundjalung, Wiradjuri, Ngunawal, Gureng Gureng, and Gagadju peoples.



USER NOTE:

There are more Content Slide options available to add to your presentation by going to:  
Home > New Slide > Drop down arrow

---

# Syllabus

- Blockchain
  - History of Blockchain
- Consensus
- Security
- Robustness
- Scalability
- Other Challenges

# Full Disclosure



Australian Government  
Australian Research Council



The University of Sydney



# Blockchain Scalability and its Foundations in Distributed Systems



Springer

---

Textbook:  
*Blockchain Scalability  
and its Foundations in  
Distributed Systems*

The techniques taught herein consist of enhancing blockchain security and making blockchain scalable by relying on the observation that no blockchain can exist without solving the consensus problem.

<https://tinyurl.com/yc22he4v>



# Online Material: Blockchain Scalability MOOC

Blockchain promises to disrupt industries once it will be efficient at large scale. In this course, you will learn how to make blockchain scale. You will learn about the foundational problem of distributed computing, consensus, that is key to create blocks securely.

<https://www.coursera.org/learn/blockchain-scalability>

Browse > Computer Science > Computer Security and Networks

Offered By



## Blockchain Scalability and its Foundations in Distributed Systems

★★★★★ 4.5 74 ratings



Vincent Gramoli

Go To Course

Already enrolled

Financial aid available

4,558 already enrolled

# What to expect from these blockchain lectures

## Fundamental knowledge

- Distributed computing
- Blockchain abstractions

## Problem solving skills

- Impossibility results
- Theoretical analysis

## Not technology-specific knowledge

- No Solidity v0.8.13 syntax
- Happy to provide pointers to this

## Course taught at private and central banks

The Blockchain Scalability MOOC on Coursera is a good overview



coursera

Browse > Computer Science >  
Computer Security and Networks

Offered By



## Blockchain Scalability and its Foundations in Distributed Systems

★★★★★ 4.5 49 ratings



Vincent Gramoli

Enroll for Free  
Starts Mar 30

Financial aid available

2,517 already enrolled

---

# Blockchain

## Why Studying Blockchain?

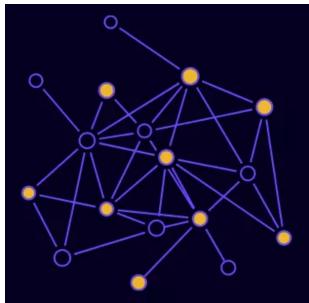
Cybersecurity is part of Australian Science and Research Priorities

Blockchain potential applications span various industries (real estate, aviation, retail, computing, etc.)

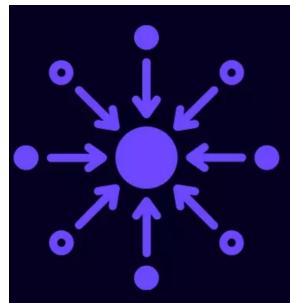
Blockchain was most desired skill on LinkedIn in 2020

# Why Studying Blockchain?

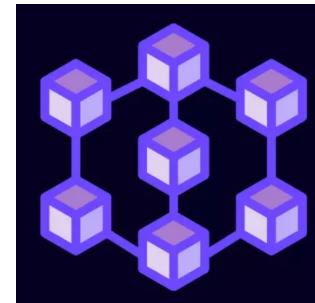
Web3, a decentralized ecosystem supported by blockchains, is a promising replacement to Web2



Web1.0 – desorganised access to information in the 90s



Web2.0 – data sharing, acquired by tech giant in the mid-2000s



Web3.0 – distributed ownership of data – 2015

*“Web 2.0 very much still exists in this very centralized model” – Gavin Wood, co-inventor of Ethereum*

*“Social networking sites that do not allow users to extract the information they put into them is the problem” Tim Berners Lee – Turing Awardee*

## Payment with Blockchain

*Bitcoin* is both a currency (lower-case b) and a payment system (big B).

It implements a **distributed ledger** that is

- Immutable: transactions can only be appended to the ledger
- Pseudonymous: participants are identified by their address (pseudonym)

It uses:

- A **Gossip-based protocol** for maintaining an overlay
- **Peer-to-peer** network where peers are both clients and servers (miners)
- **Public key cryptosystem**, participant sign their transactions
- **Consensus** for the distributed system to decide a block at index I

# What is a blockchain?

## The Blockchain abstraction

Let  $G = \langle B, P \rangle$  a directed acyclic graph (DAG) where blocks  $B$  point to each other with pointers  $P$

$\langle b_0, b_1 \rangle \in P$  is a pointer from current block  $b_1$  to previous block  $b_0$



The *pointer* is a representation of a hash of the destination block that the source block contains

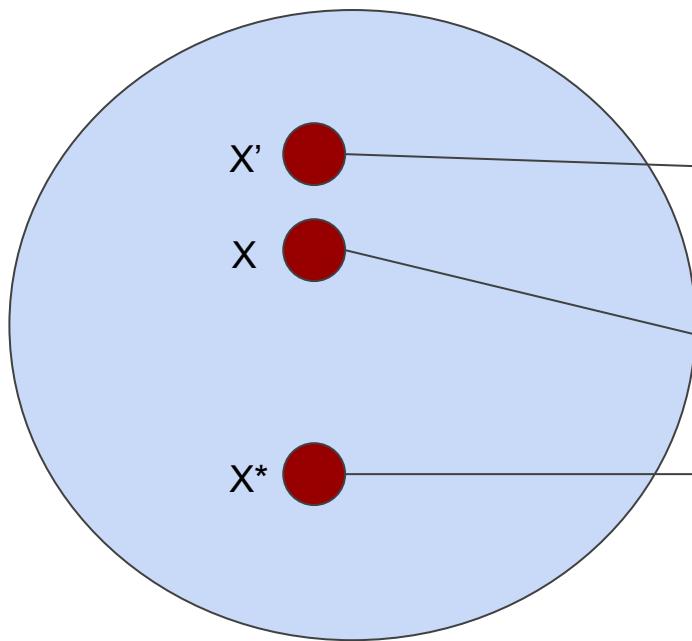
The *genesis block* is a special block known initially by all participants

# Hash Function

# HASH FUNCTION

---

Original String



N-bit Hash Value

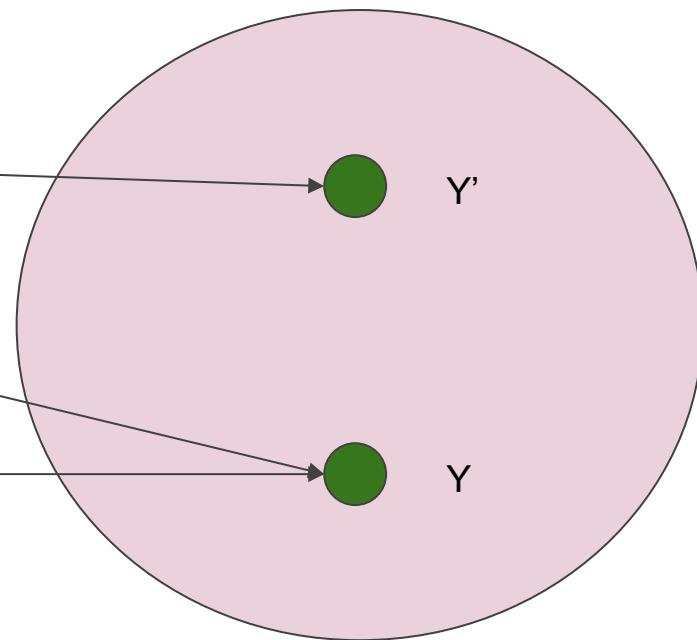


Figure 3.2: Illustration of Hash Function

Year 2023

# Cryptographic Hash Function - Ideal Properties

1. **Deterministic**: always same output hash for the same input
2. **Fast**: quick to obtain the hash value
3. **Secure**: hard to find an input that will generate a given ouput hash
4. **Discontinuous**: a small change in input should change hash substantially
5. **Low Collision**: hard to find two inputs that give the same output hash

E.g., SHA256

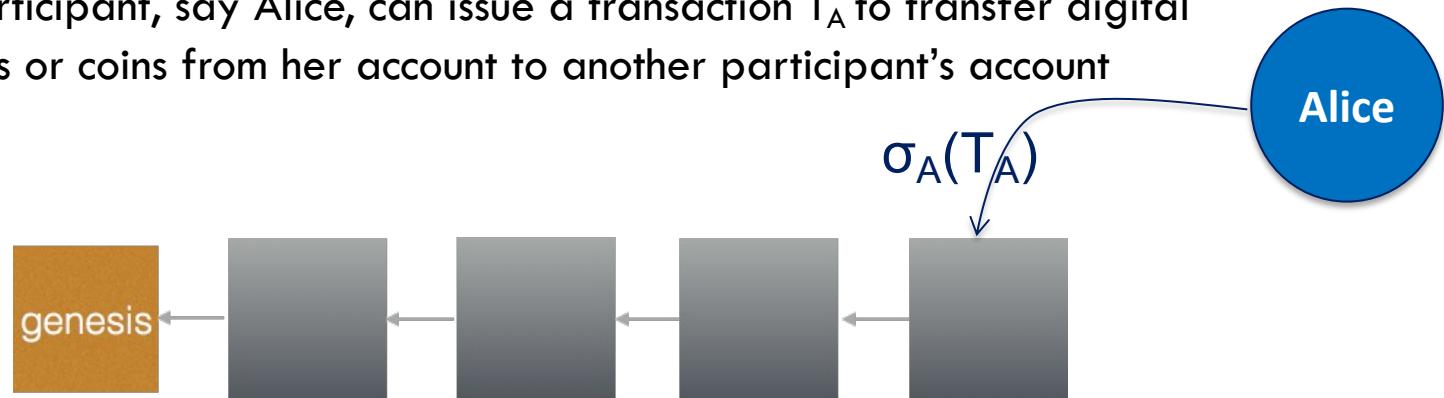
# Applications of hash functions

- Github commit identifier
  - You have used a version of a software that works but the most up-to-date version does no longer work.
  - You check the commit hash of the version you use and roll-back the code to the working hash
- Software version check
  - You download an large image and want to know whether it is valid copy and not a version infected with malware or trojan.
  - You ask the developer to provide an MD5 hash of the image and compare it to the result of applying the hash function to the image.
- Message digest
  - Someone sends you a message along with a signed hash.
  - You compute the hash of the message content and compare it to the signed hash to detect tampering.

# Signatures

## Blocks contain transactions

A participant, say Alice, can issue a transaction  $T_A$  to transfer digital assets or coins from her account to another participant's account



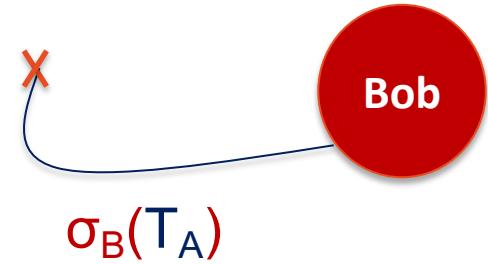
She simply signs her transaction  $T_A$  with her signature  $\sigma_A$ , then the signed transaction  $\sigma_A(T_A)$  has to be stored in a block of the blockchain

## Transactions are securely signed

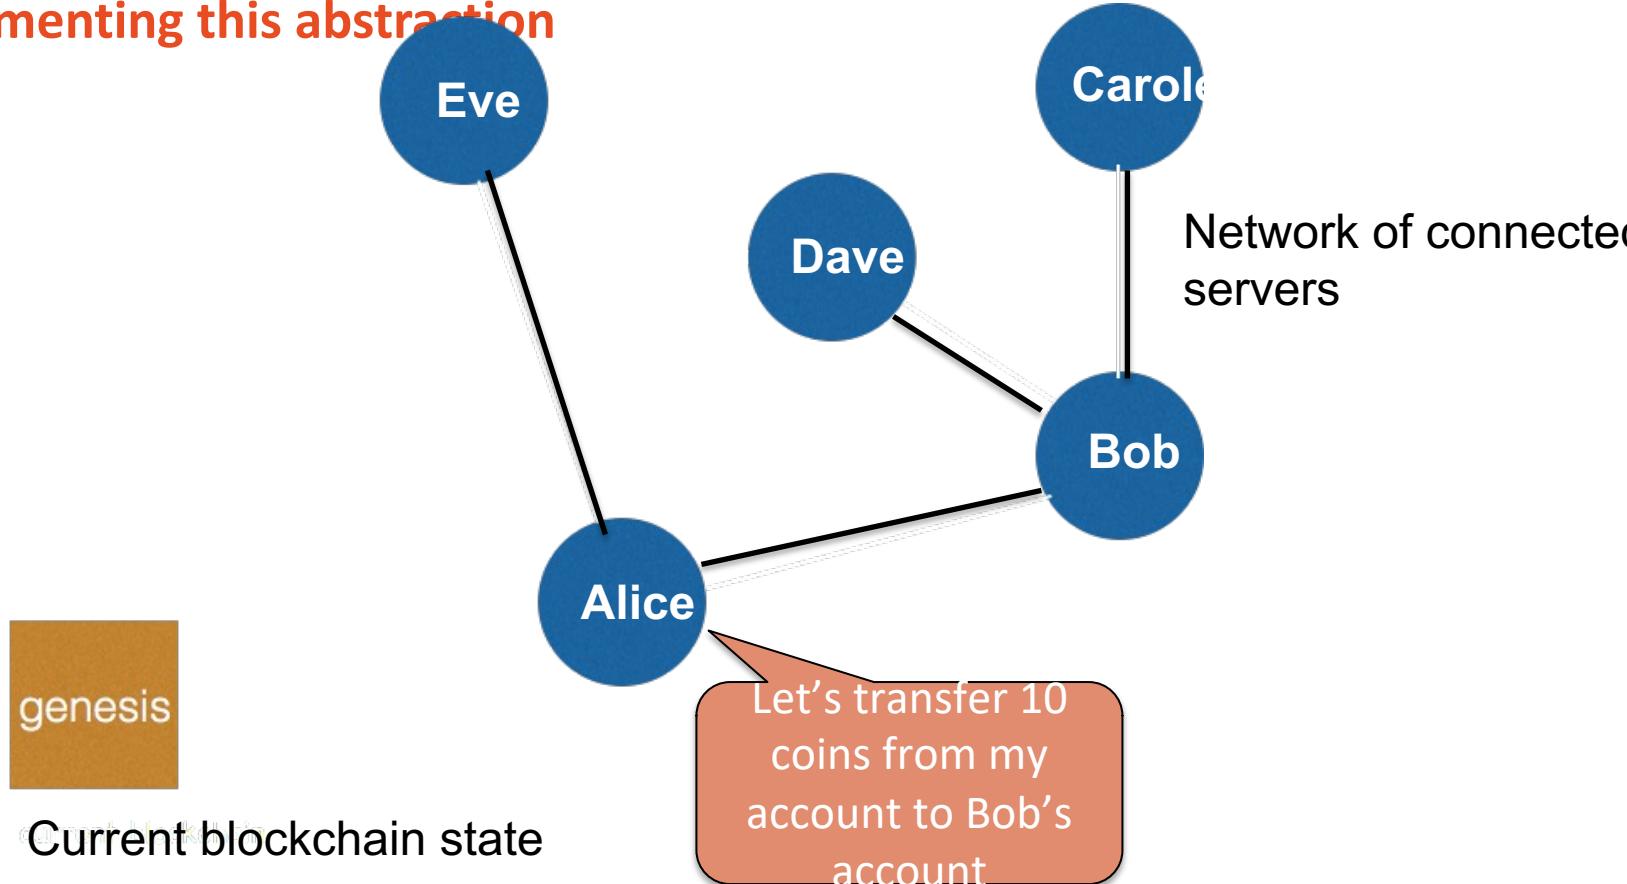
Because the issuer has to sign its transaction with its private key no other participant can withdraw his coins without his consent



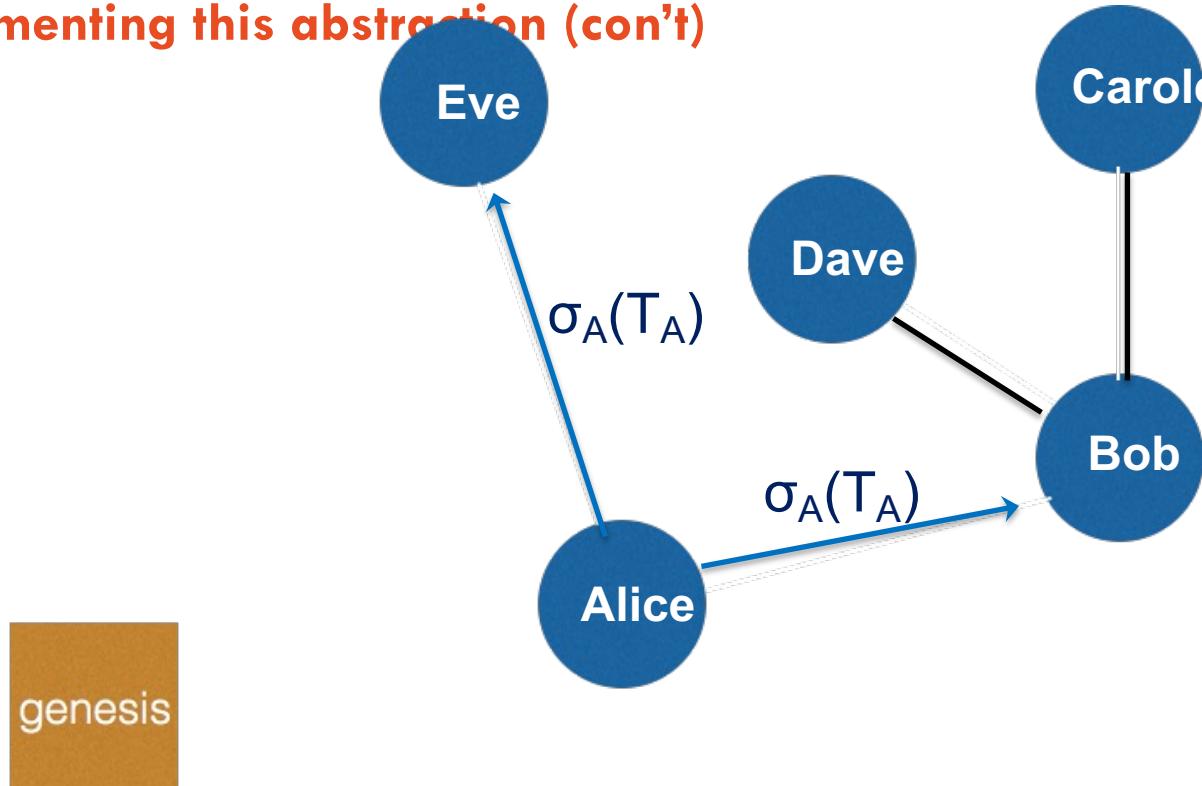
Any transaction not signed by the issuer will not be stored in a block



## Implementing this abstraction

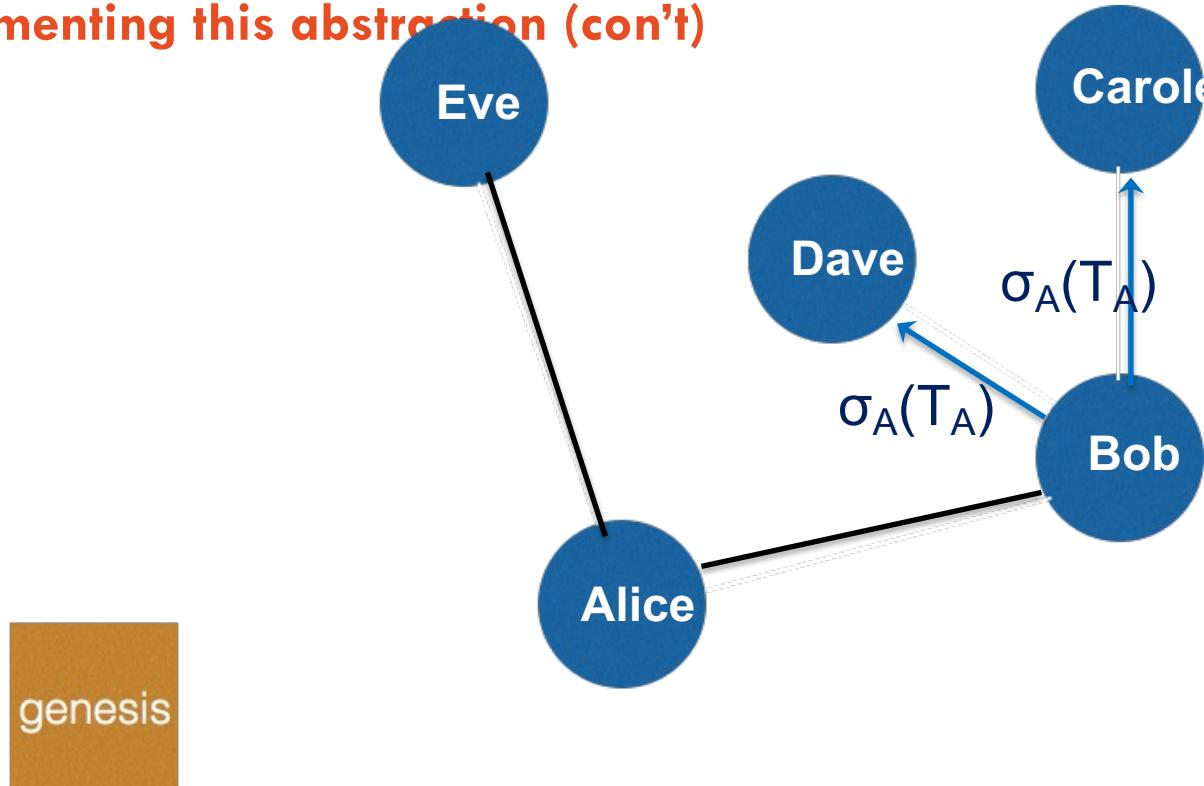


## Implementing this abstraction (con't)



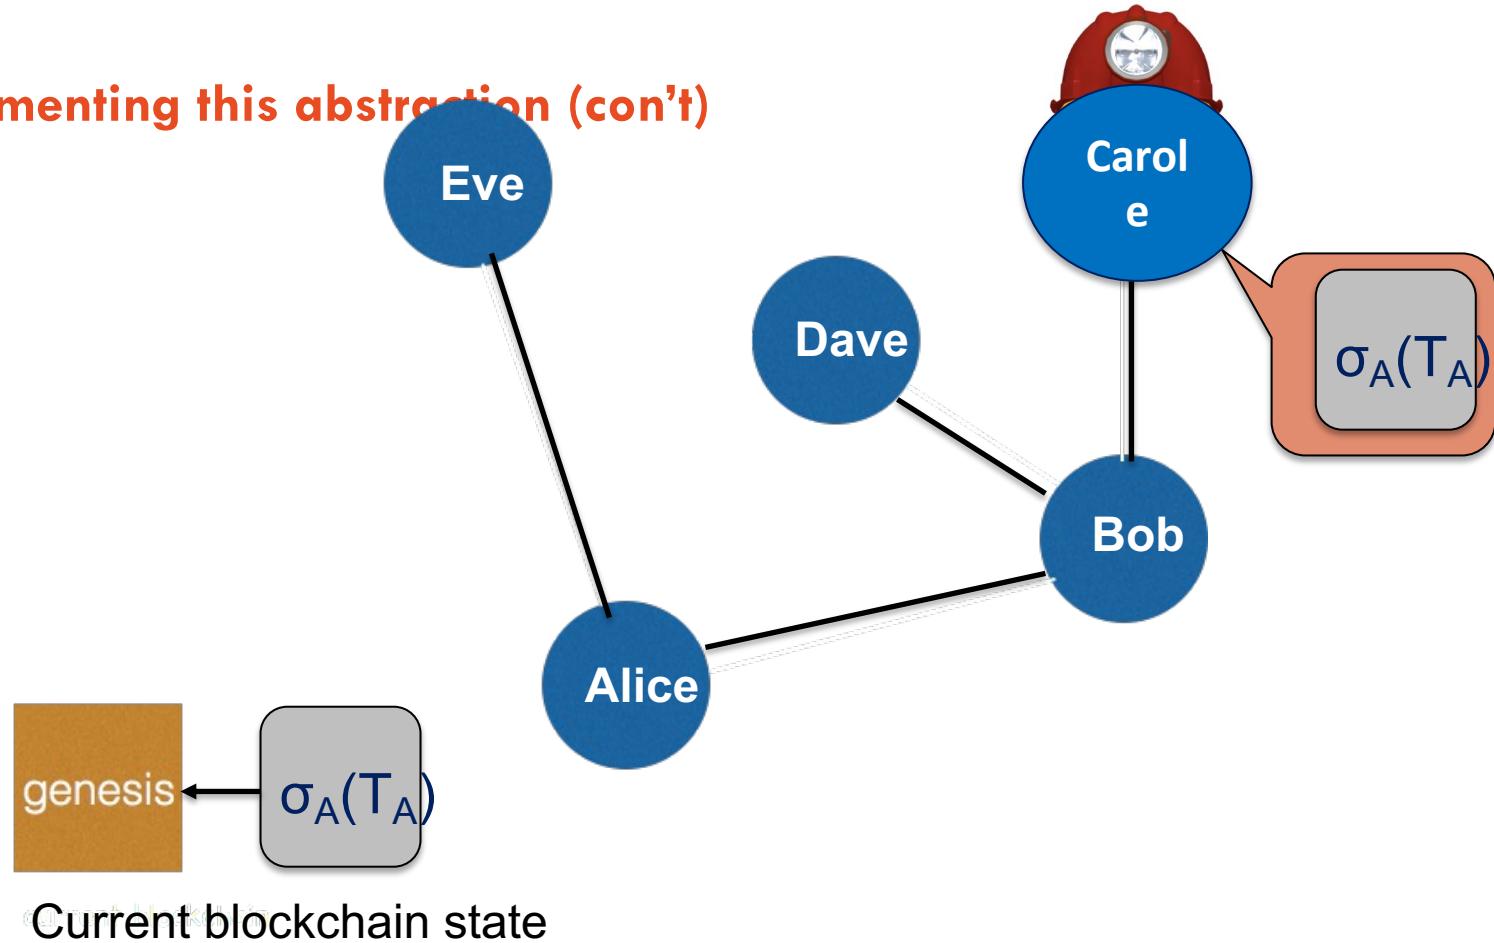
Current blockchain state

## Implementing this abstraction (con't)

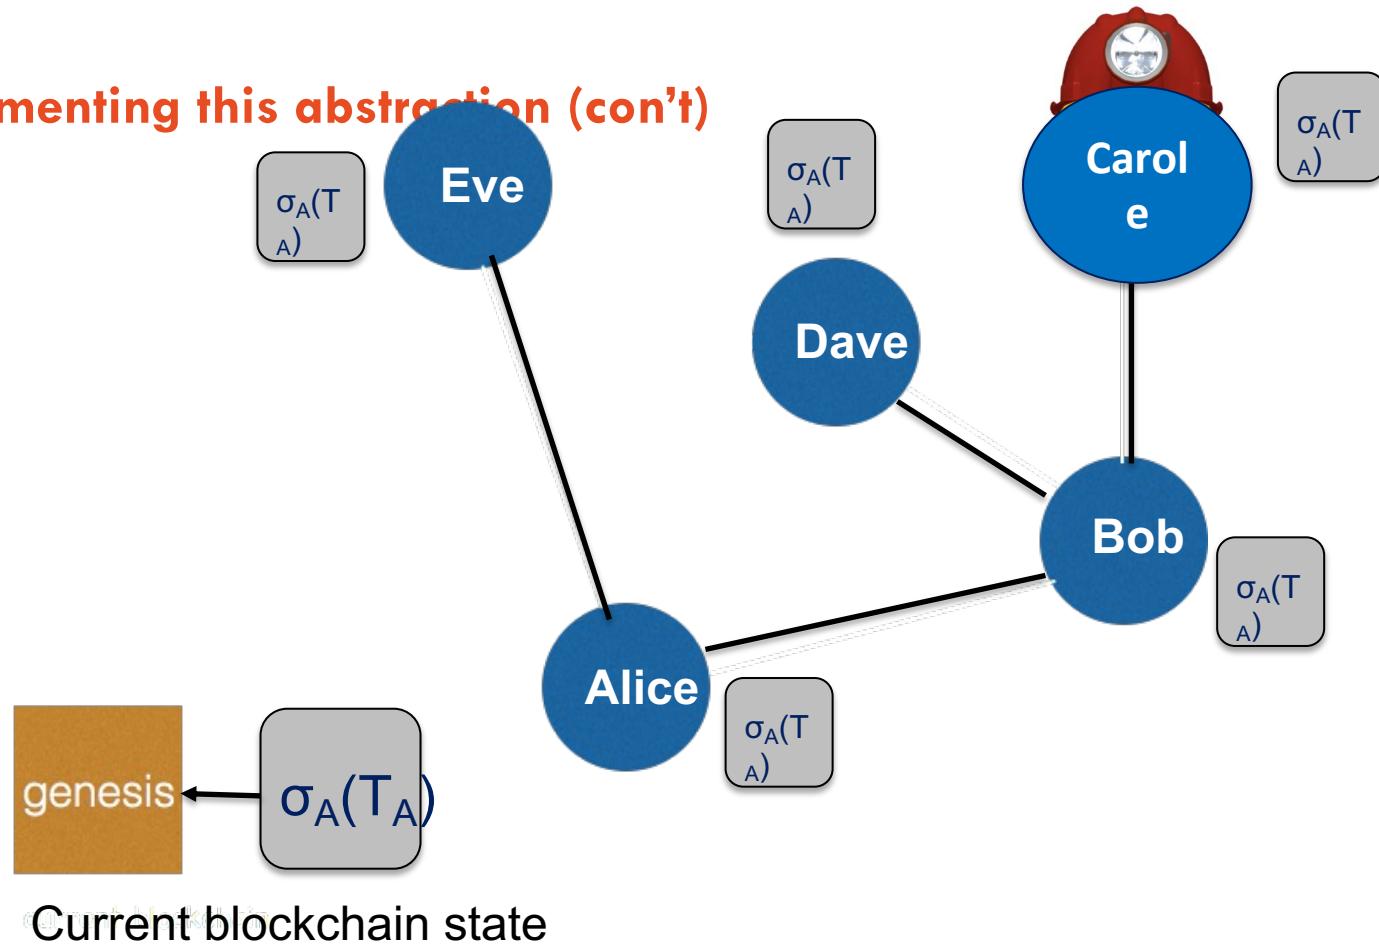


Current blockchain state

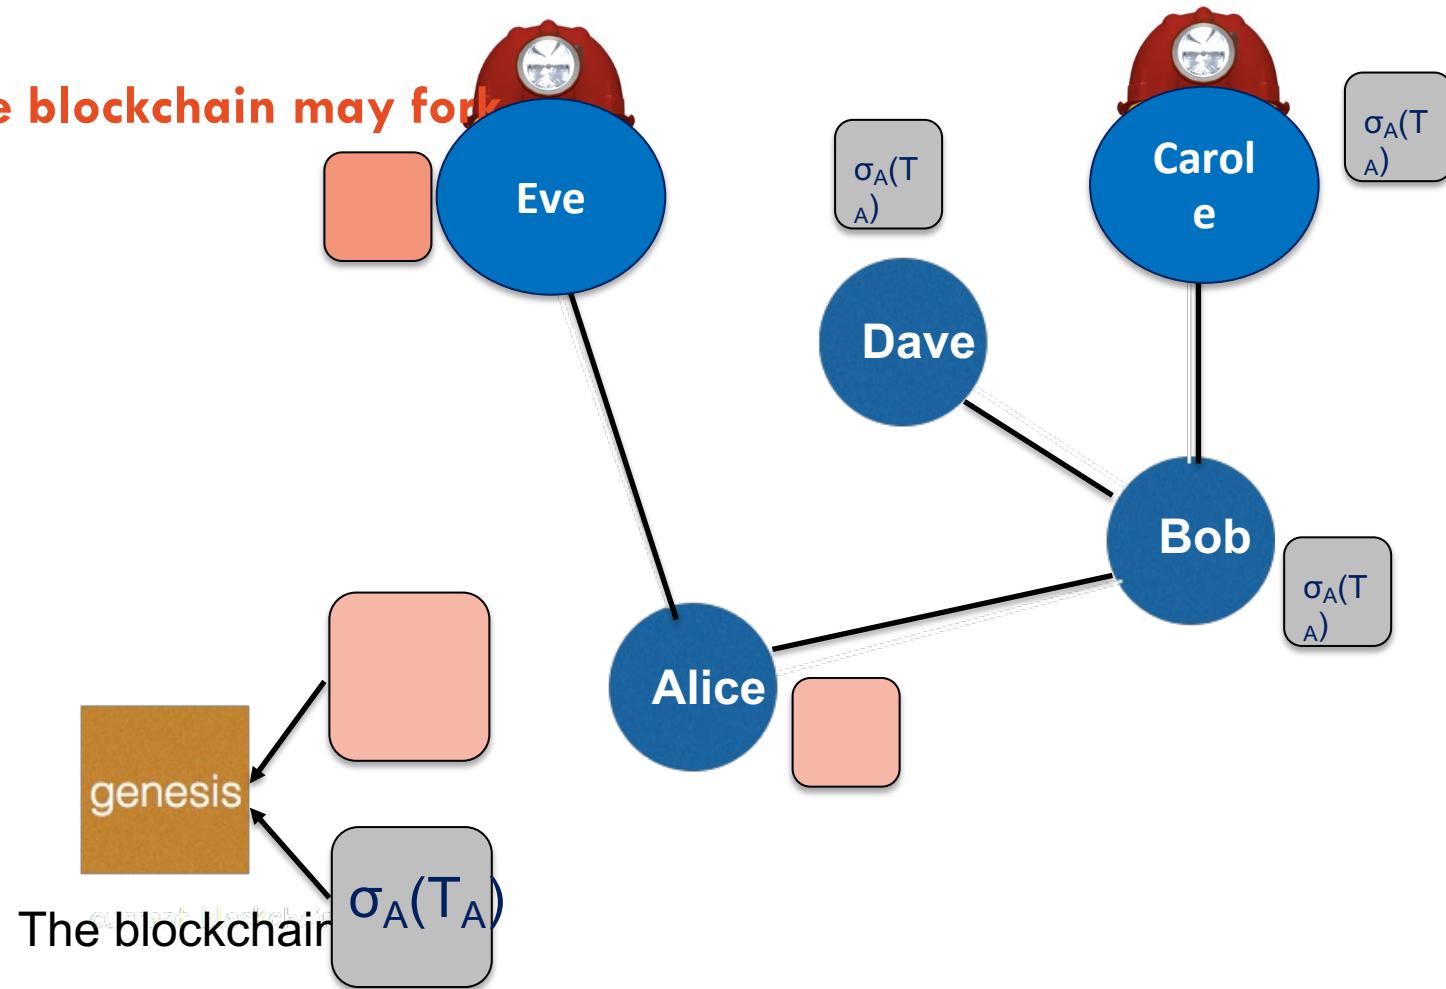
## Implementing this abstraction (con't)



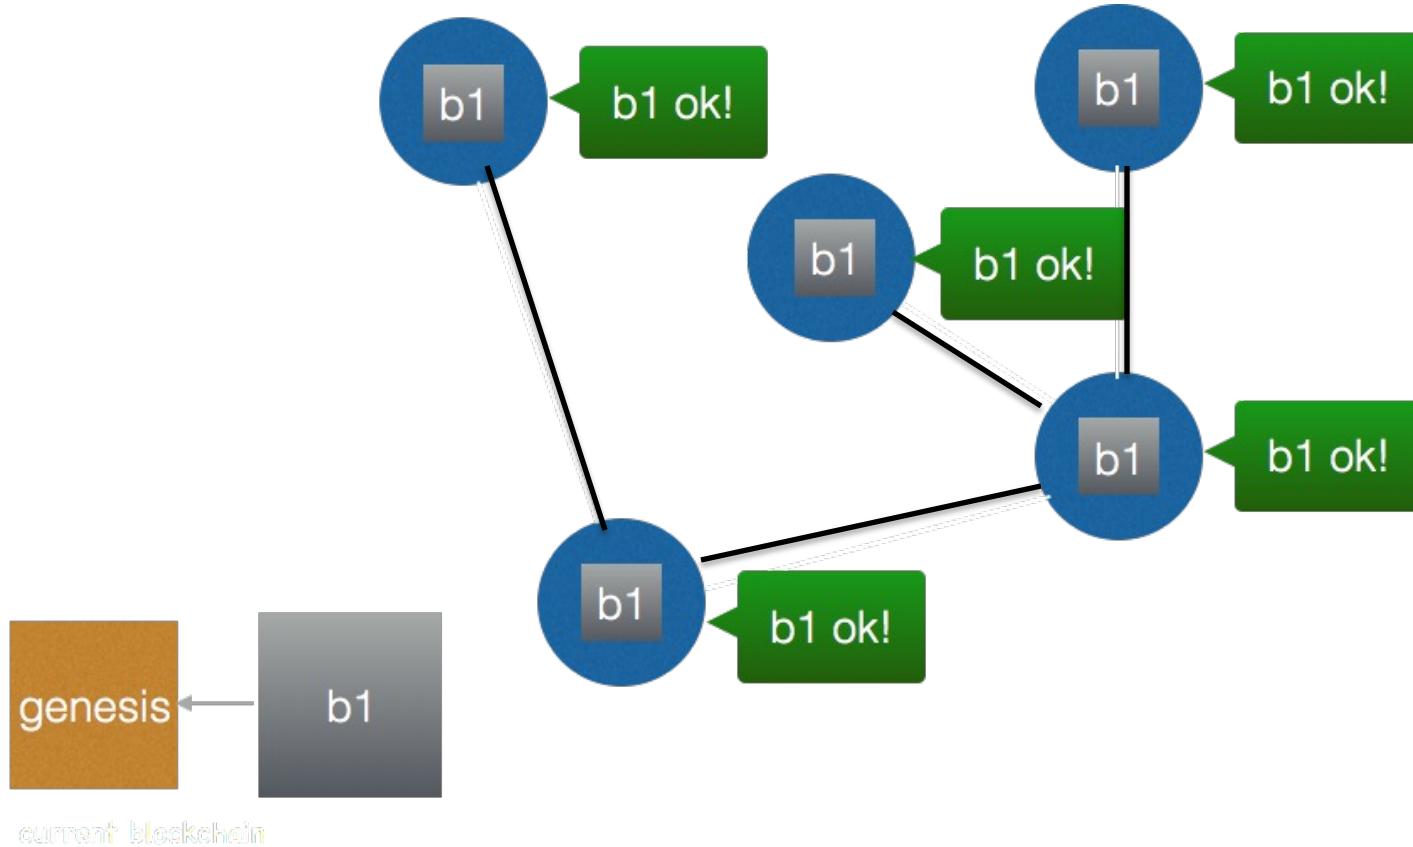
## Implementing this abstraction (con't)



The blockchain may fork



# Distributed implementation of a blockchain



# History of Blockchains

## Before Bitcoin

### Cypherpunks and digital currency

- **1975** – A symmetric key cryptosystem is published as the Data Encryption Standard.
- **1983** – David Chaum proposes blind signatures for a payment system that offers auditability and privacy.
- **1989** – He founds DigiCash Inc., an electronic money corporation, that sold its assets to eCash.
- **1992** – Eric Hughes, Timothy May and John Gilmore meet on a monthly basis in San Francisco and the cypherpunk mailing list is created.

[IBM75] IBM, “Data Encryption Standard”, US Federal Register, 1975.

[Cha83] D. Chaum, “Blind Signatures for Untraceable Payments”, Advances in Cryptology, Springer 1993. [https://link.springer.com/chapter/10.1007/978-1-4757-0602-4\\_18](https://link.springer.com/chapter/10.1007/978-1-4757-0602-4_18)

## Before Bitcoin

*“Privacy in an open society requires anonymous transaction systems. Until now, cash has been the primary such system. An anonymous transaction system is not a secret transaction system.”*

Eric Hugues, 1993



<https://medium.com/coinmonks/a-cypherpunks-manifesto-a-definition-of-privacy-66c36f99e940>

## Before Bitcoin

### Proof-of-work and double spending

- **1992** – Cynthia Dwork and Moni Naor require a user to compute a moderately complex function in order to avoid abuse of resource usage.
- **1997** – Adam Black proposes a hashcash mint, a partial hash collision generator, on the cypherpunk mailing list. The idea is to use partial hashes to be made arbitrarily expensive to compute but verified instantaneously to detect double spending: <http://www.hashcash.org/papers/announce.txt>
- **2002** – Adam Black publishes the Hashcash technical report.
- **2004** - Hal Finney introduced “reusable proofs of work” and offers a public implementation of a transparent server whose integrity can be checked. <https://cryptome.org/rpow.htm>

[DN92] Dwork and Naor, “Pricing via Processing or Combatting Junk Mail”, CRYPTO’92.

<https://dl.acm.org/citation.cfm?id=705669>

[Bla02] Black, “Hashcash - a denial of service counter-measure”, Cypherspace, TR 2002.

<http://www.hashcash.org/papers/hashcash.pdf>

## Before Bitcoin

### Bit gold and B-money

- **1998** – Wei Dai proposes b-money a system where users can transfer money between accounts represented as public keys by broadcasting messages
- **2005** – Nick Szabo proposed Bit Gold crypto-currency that applies a chain of hash to guarantee immutability <https://nakamotoinstitute.org/bit-gold/>
- **2008** – Satoshi Nakamoto writes Bitcoin, a protocol to transfer a cryptocurrency peer-to-peer without the need of central intermediary.

[Sza05] Szabo. Bit gold. Technical report, 2005.

[Nak08] Nakamoto. ["Bitcoin: A Peer-to-Peer Electronic Cash System"](#) (PDF). March 2014.

# Before Bitcoin

Bit gold

# Bitcoin: A Peer-to-Peer Electronic Cash System

Satoshi Nakamoto  
satoshin@gmx.com  
www.bitcoin.org

**Abstract.** A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network requires minimal structure. Messages are broadcast on a best effort basis; it requires no central authority or identities beyond what happened while they were gone.

[S  
[HS  
[Nak

financial institutions serving as  
works well enough for  
trust based model.  
ations cannot

✓ between  
hash to  
peer-to-

## After Bitcoin

- **2009** – Blockchain v1.0: bitcoin is released as an open source software.
- **2013** - Vitalik Buterin describes Ethereum (Blockchain v2.0) based on generic crypto-puzzle and offers services beyond cryptocurrency.
- **2015** – Gavin Wood writes the Ethereum yellow paper.
- **2016** – A Decentralized Autonomous Organization (DAO) was proposed as a smart contract, raised \$150M and got hacked.
- **2017** – More secure models of blockchains appear (e.g., AlgoRand, Red Belly Blockchain).

[Woo15] – Ethereum: A Secure and Decentralized Generalized Transaction Ledger

[GHM17] – Gilad, Hemo, Micali, Vlachos, Zeldovich. AlgoRand: Scaling Byzantine Agreements for Cryptocurrencies. SOSP'17.

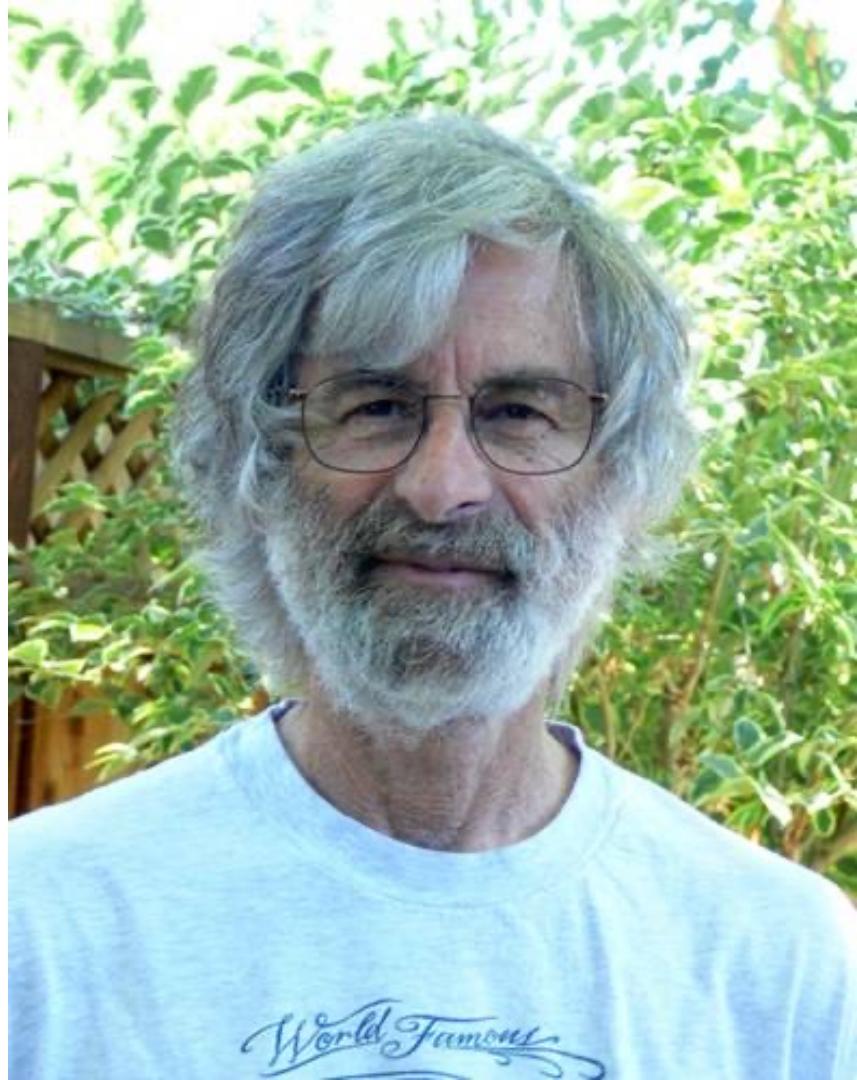
[Gra17] – The Red Belly Blockchain. Personal Presentation. MIT, USA.

# Failures

## Failures

*"A distributed system is a system where I can't get my work done because a computer has failed that I've never even heard of."*

Leslie Lamport



## Failure taxonomy

- *Failure*: a system *fails* when it cannot meet its promises
- *Crash failure*: a machine halts but is working correctly until it halts
- *Arbitrary (Byzantine) failure*: a machine may produce arbitrary responses at arbitrary times or may not send responses
- The most serious failure is arbitrary (a.k.a. **Byzantine**), it can be of any type.

## Failures

Motivations for arbitrary (Byzantine) failures:

- Malfunctioning hardware
- Buggy software
- Malicious attacks

In blockchains, there is an incentive to steal digital assets, so malicious attacks are expected.



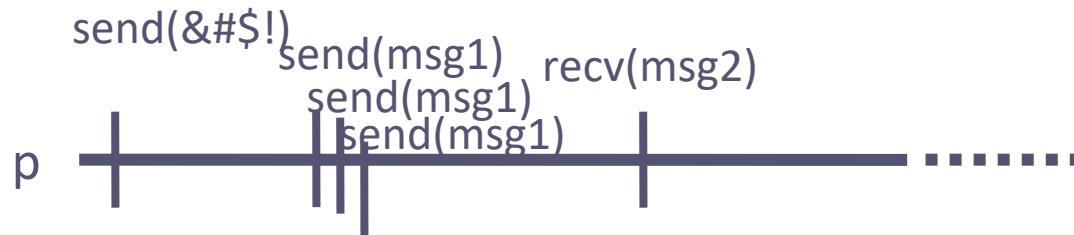
## Failures

We will focus on

- Crash failures:



- Byzantine failures:



- A correct process never fails

# Coping with failures

# Consensus

## Why Consensus?

The consensus problem has a lot of applications:

- With consensus you can decide on a leader
- Coordinate machines to do an attack

Similarities with other problems of distributed computing:

- Totally ordered broadcast
- Atomic commit
- Terminating reliable broadcast
- ...

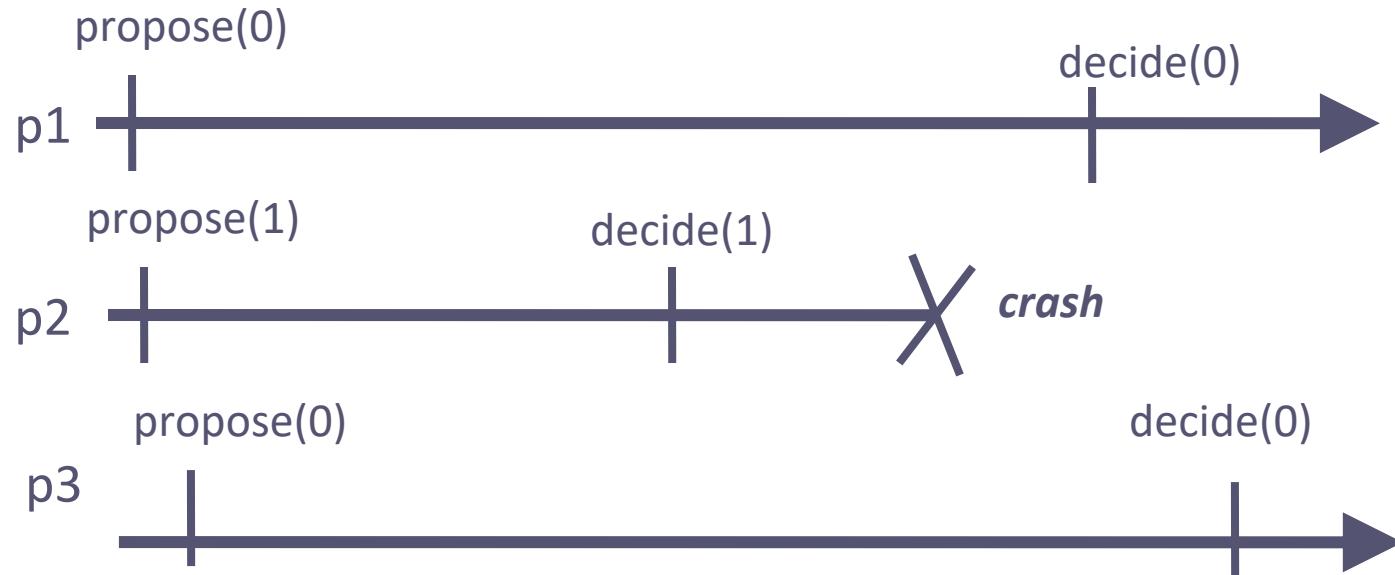
## The Consensus Problem

In the consensus problem, the processes propose values and have to agree by deciding one among these values

### Consensus Definition 1:

1. *Agreement:* No two correct processes decide differently
2. *Termination:* Every correct process eventually decides
3. *Validity:* Any decided value is a value proposed

## The Consensus Problem (con't)



# Consensus and Blockchain

## Why Consensus in blockchain?

To construct a blockchain, distributed processes have to agree on a unique block at a given index

Assume all correct processes know the genesis block at index 0:



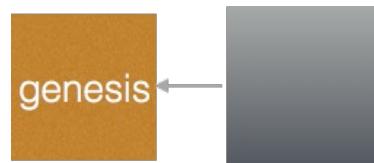
## Why Consensus in blockchain?

To construct a blockchain, distributed processes have to agree on a unique block at a given index

Assume all correct processes know the genesis block at index 0:

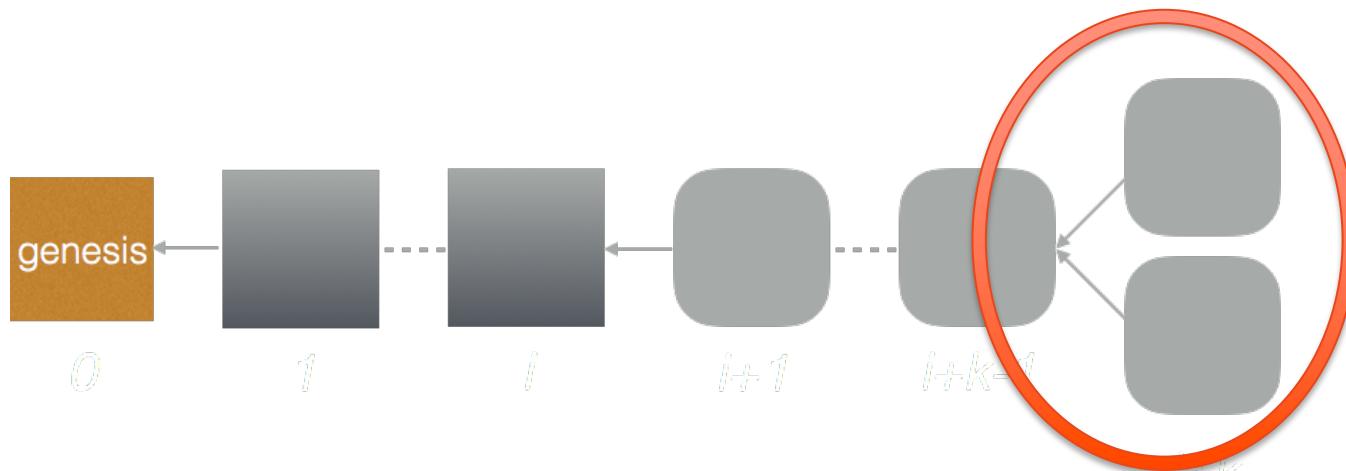


Then, they run consensus at index 1 to decide upon a new block:



## Why Consensus in blockchain?

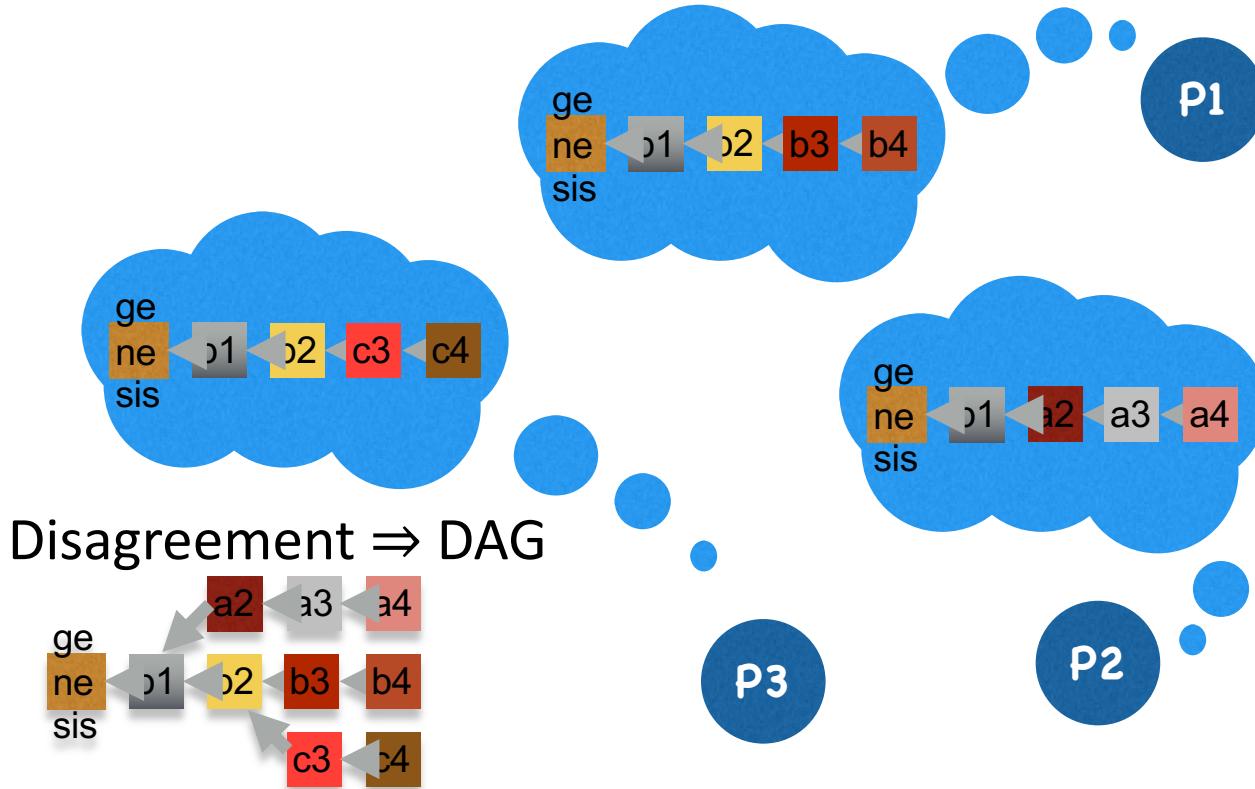
Consensus is necessary to totally order the blocks, maintaining the *chain*



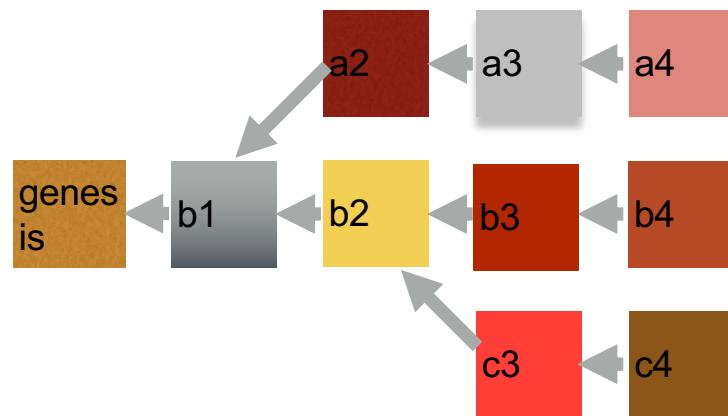
Disagreement ⇒  
no chain but a  
DAG

# Double Spending

## Disagreement



## Directed Acyclic Graph (DAG)



## Double Spending

Let Alice be a Byzantine process willing to steal assets (coins)

Let Bob and Carole be two merchants selling real products in exchange of coins

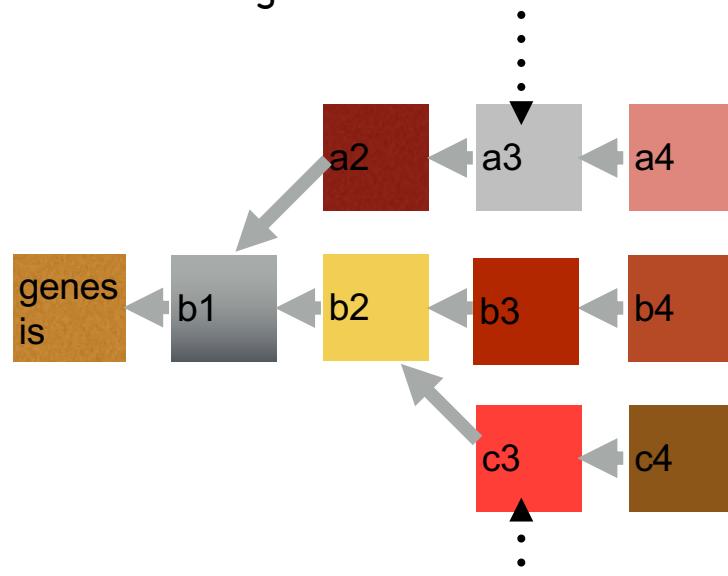
Alice issues:

- a transaction  $t_A$  spending **all** her coins with Bob
- a transaction  $t_A'$  spending **all** her coins with Carole

If both Bob and Carole observe their transaction in the blockchain, they will ship the “supposedly paid” product to Alice.

## Double Spending (con't)

$t_A$ : Alice gives all her coins to Bob



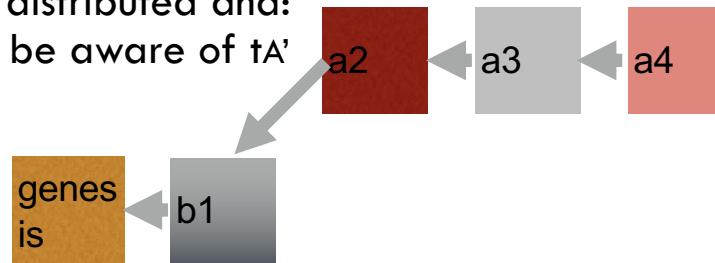
$t_A'$ : Alice gives all her coins to Carole

## Double Spending (con't)

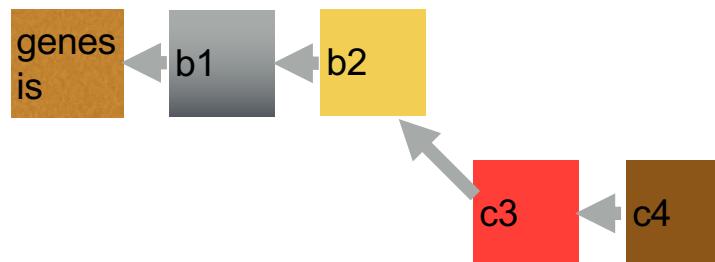
The problem is that  $tA$  and  $tA'$  conflict

But the system is distributed and:

- Bob may not be aware of  $tA'$



- while Carole may not be aware of  $tA$



# Blockchain Failures

## Failure model

### Model

- $n$  nodes in the system
- $f$  are faulty

Because blockchains protect ownership, there is an incentive for an attacker to steal the goods of others

- The fault model is **Byzantine** (i.e., arbitrary)



## Solution for Byzantine consensus

Limiting the number  $f$  of failures is key to solving consensus

There are solutions when  $f < n/3$  [CL02]



[CL02] M. Castro and B. Liskov. Practical byzantine fault tolerance and proactive recovery. ACM Trans. Comput. Syst., 20(4):398-461, Nov. 2002.

## Theorem

There is no solution to the Byzantine consensus if  $f=1$  and  $n=3$ .

### Proof intuition.

Consider three processes  $p_1$ ,  $p_2$ ,  $p_3$ . Byzantine process  $p_2$  can play  $p_1$  and  $p_3$  off against each other, telling  $p_1$  that  $p_3$  is Byzantine and  $p_3$  that  $p_1$  is Byzantine. Since  $p_1$  is telling  $p_3$  the same thing about  $p_2$  that  $p_2$  is saying about  $p_1$ ,  $p_3$  cannot tell the difference and does not know who to believe.

## Theorem

There is no solution to the Byzantine consensus if  $f \geq n/3$ .

### Proof.

By reduction of the impossibility result of when  $f \geq n/3$  to when  $f=1$  and  $n=3$ .

Suppose that there were an algorithm that solved Byzantine agreement with  $n=3f$  processes. Group the processes into groups of size  $f$ . Let each of the  $n=3$  processes simulate one group of processes receiving the same input. Then we get a protocol for  $n=3$  and  $f=1$ , which is a contradiction.

# Proof of Work

## Sybil attack

A *Sybil attack* is an attack where a malicious user forges identities. It is named after the subject of the book *Sybil*, a case study of a woman diagnosed with dissociative identity disorder.



Some solutions [CL02] are prone to Sybil attacks where an adversary generates fake faulty nodes to have  $f \geq n/3 \Rightarrow$ consensus impossible.



## Miners

Specialised peers, called *miners*, receive a reward for verifying transactions provably solving a cryptopuzzle [Bla02] to append a new transaction block to the blockchain.

Cryptopuzzle: given a **block** and a **threshold**, a miner repeatedly:

- selects a nonce and
  - applies a pseudo-random function to this block and the selected nonce
- ...until it obtains a result lower than the threshold.

[Bla02] A. Black, “Hashcash - a denial of service counter-measure”,  
Cypherspace, TR 2002. <http://www.hashcash.org/papers/hashcash.pdf>



## Proof-of-Work

The nonce is included in the block, this is the *proof-of-work* [DN93]:

- finding the nonce takes time, but
- validating that the nonce is correct is easy.

Everyone can verify that someone lied about having solved the puzzle

[DN93] C. Dwork and M. Naor. Pricing via processing or combatting junk mail. In Proceedings of the 12th Annual International Cryptology Conference on Advances in Cryptology, CRYPTO '92, pages 139-147, 1993.

# Cryptopuzzle

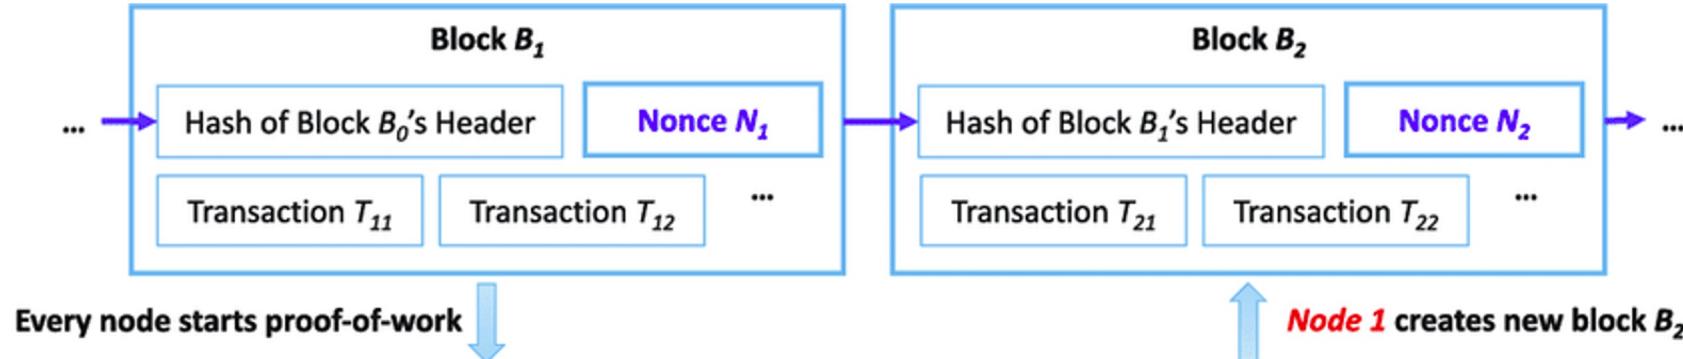
# The Bitcoin Protocol - Proof of Work

All nodes on the network try to create a block. There is a “lottery” that selects one node every ten minutes or so to be the “miner” of the next block.

The lottery is implemented by having all computers try, in parallel, to solve a hard cryptographic puzzle called the “Proof of Work” puzzle.

The difficulty of the puzzle is automatically adjusted over time so that it is harder to solve when there are more nodes - on average 10 minutes to pick a winner.

# PROOF OF WORK



**Node 1**

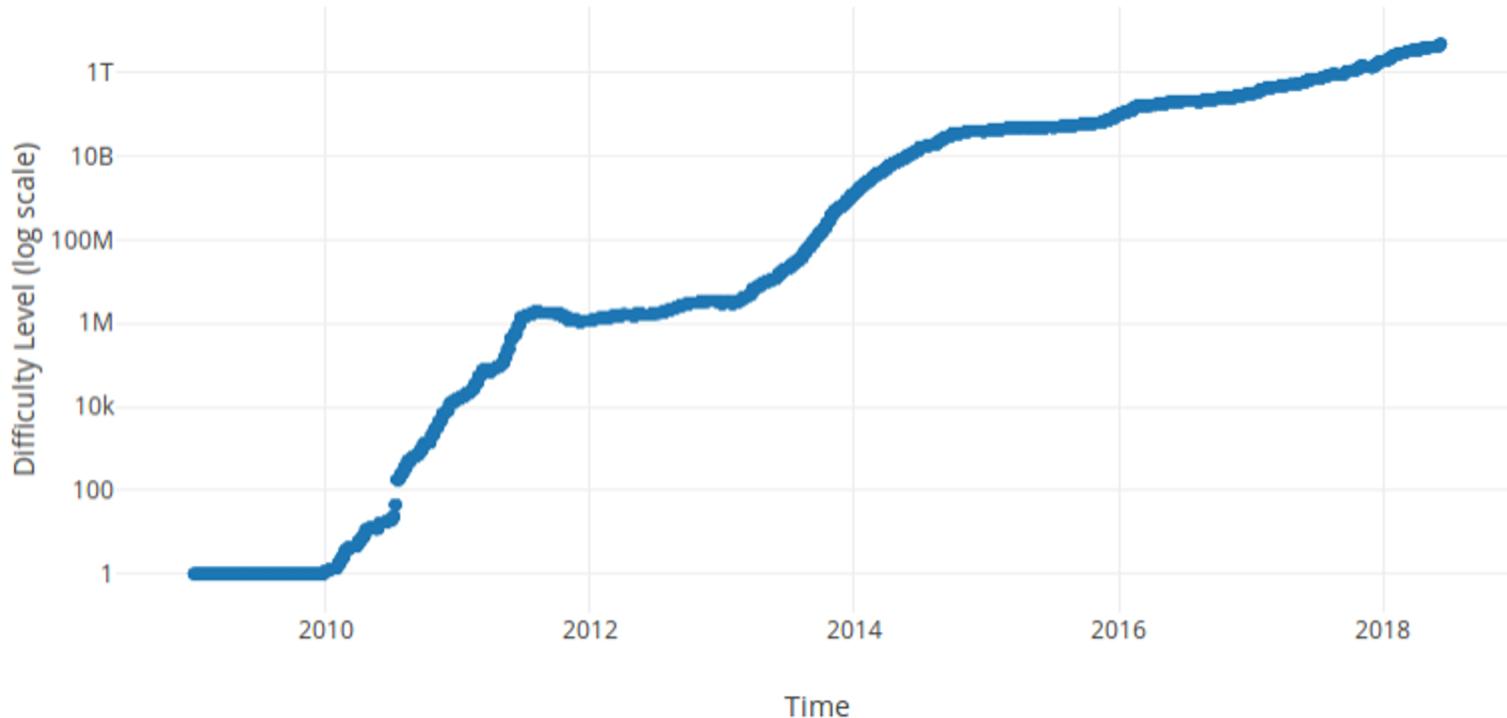
...  
Nonce  $N_2 = "7C\ 4D\ DB\ 29"$  → Hash of  $B_2$ 's header = "2D\ F8\ 8E\ 32\ ... 10\ 9A\ FE\ 1C", Failed (time = 10:14:20)  
Nonce  $N_2 = "7C\ 4D\ DB\ 30"$  → Hash of  $B_2$ 's header = "41\ 2A\ B3\ DC\ ... 94\ 29\ AB\ B5", Failed (time = 10:14:25)  
Nonce  $N_2 = "7C\ 4D\ DB\ 31"$  → Hash of  $B_2$ 's header = "00\ 00\ 4F\ 65\ ... 2F\ ED\ 31\ 09", Succeeded (time = 10:14:30)

**Node 2**

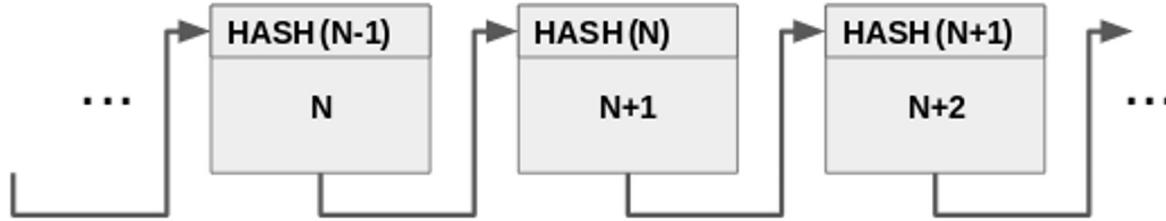
...  
Nonce  $N_2 = "61\ 0A\ 3F\ 3A"$  → Hash of  $B_2$ 's header = "A8\ C7\ 08\ C9\ ... 3D\ F1\ A2\ F9", Failed (time = 10:14:23)  
Nonce  $N_2 = "61\ 0A\ 3F\ 3B"$  → Hash of  $B_2$ 's header = "2A\ E9\ 84\ 66\ ... 91\ B4\ 58\ CE", Failed (time = 10:14:28)  
Stopped after identifying that **Node 1** has completed proof-of-work at time = 10:14:30

**Node 3**

...  
Nonce  $N_2 = "99\ 06\ 10\ 13"$  → Hash of  $B_2$ 's header = "FB\ 2F\ 26\ D9\ ... 39\ F5\ C1\ 0B", Failed (time = 10:14:21)  
Nonce  $N_2 = "99\ 06\ 10\ 14"$  → Hash of  $B_2$ 's header = "E2\ 1C\ 09\ 05\ ... 25\ 3E\ AA\ CF", Failed (time = 10:14:26)  
Stopped after identifying that **Node 1** has completed proof-of-work at time = 10:14:30



# Proof of Work also helps with immutability

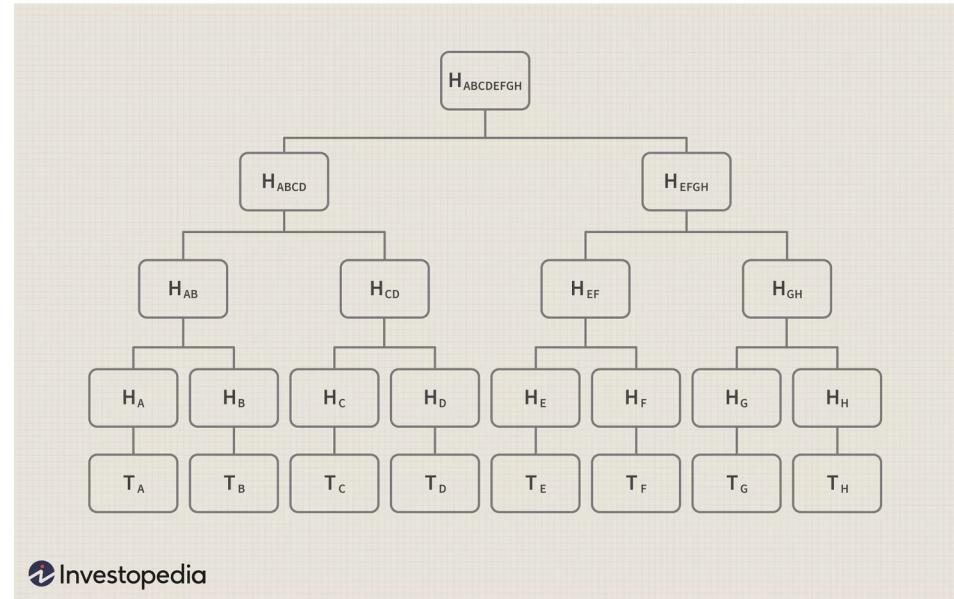


- Each new block added to the chain incorporates the hash (summary) of the previous block
- If someone wants to modify / change a block, they would have to change the hash in the next block, which changes that block, so they would have to change the hash in the block after and so on...
- The proof of work puzzle makes it hard to recreate a block after modification.
- And thus it makes it hard to modify older blocks
- For all intents and purposes, practically, any block that is 4-6 blocks “deep” into the chain cannot be modified, i.e. the blockchain makes transactions in those blocks **immutable**

# The Bitcoin Protocol - Merkle Tree

The transaction stored in each block are stored in a data structure called the Merkle tree that makes it computationally efficient to confirm membership.

Verifying membership of a leaf node requires a log number of computations



 Investopedia

E.g. given  $H_{root}$ ; to verify  $T_D$  is present, only need  $H_C$ ,  $H_{AB}$ ,  $H_{EFGH}$

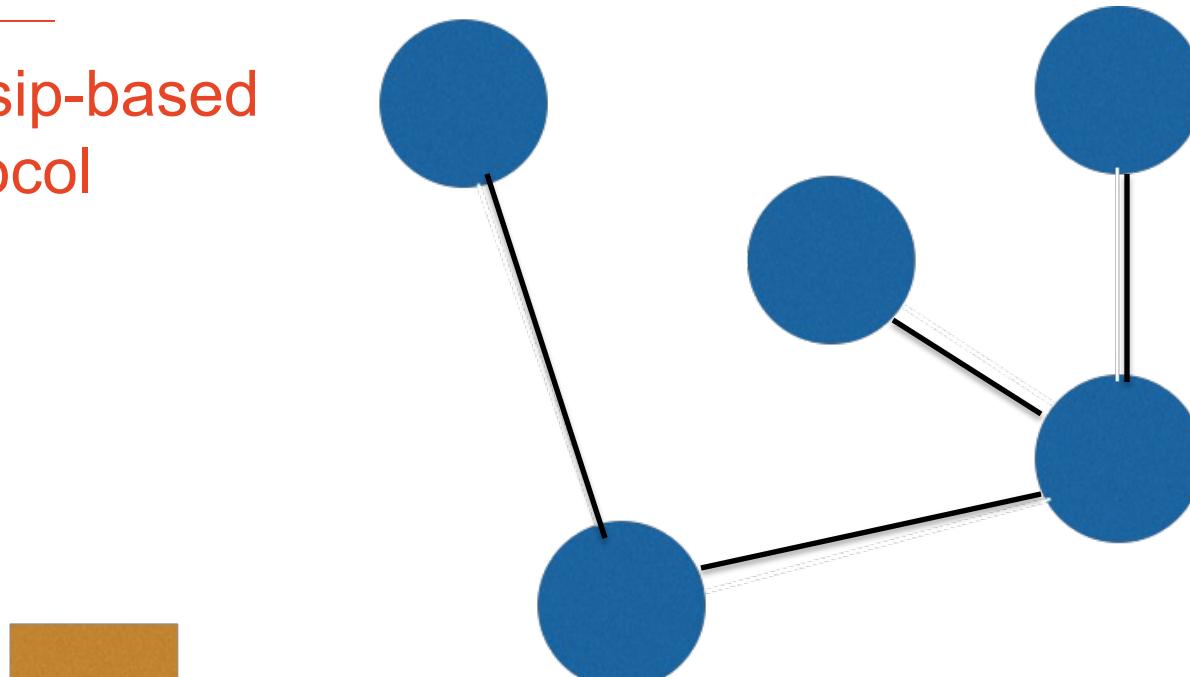
The University of Sydney

Year 2023

# Execution

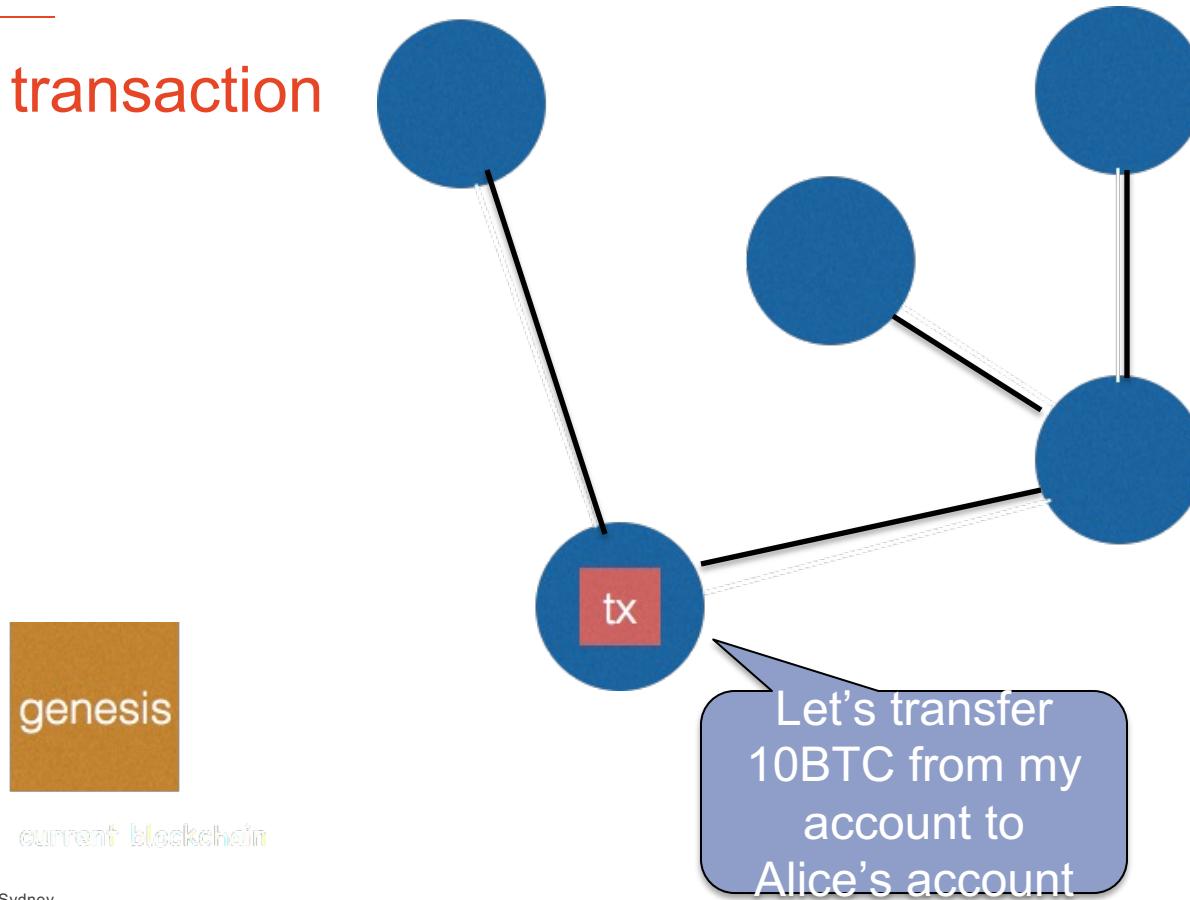
---

## Gossip-based protocol



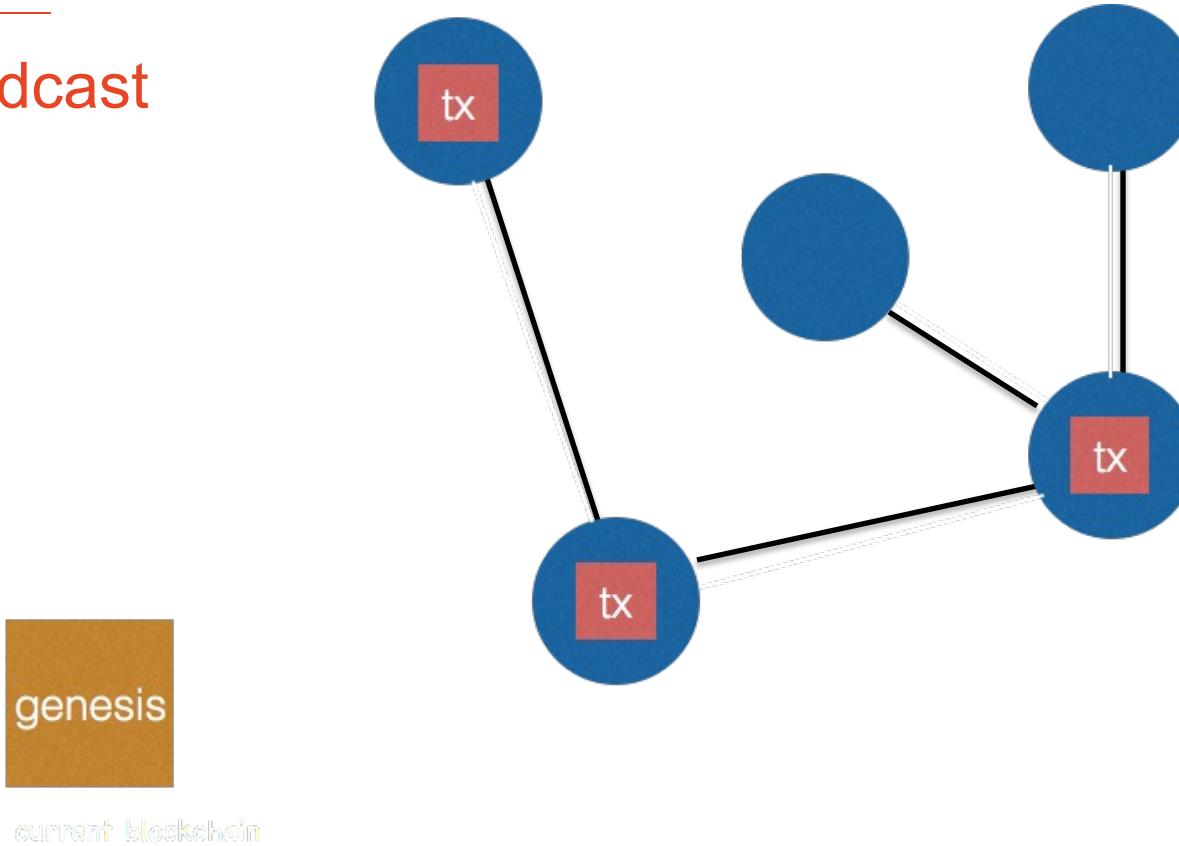
Current blockchain state

## New transaction

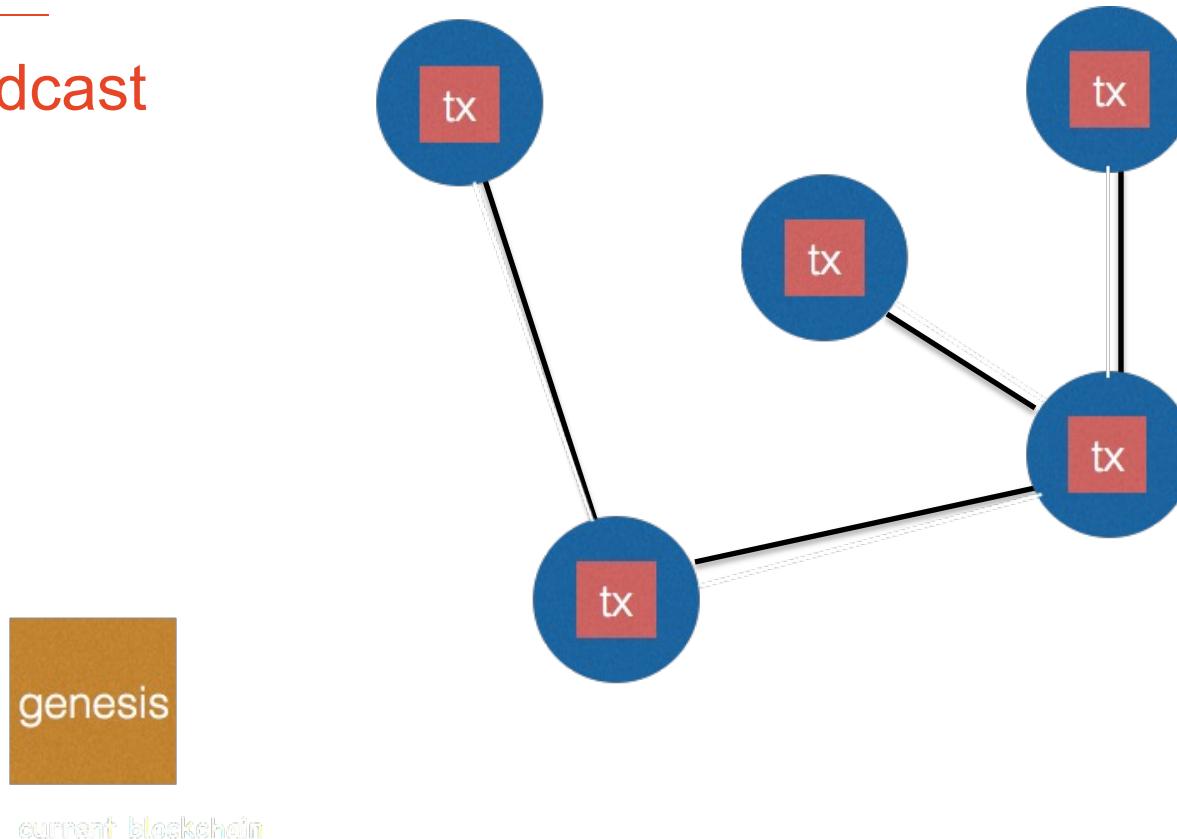


---

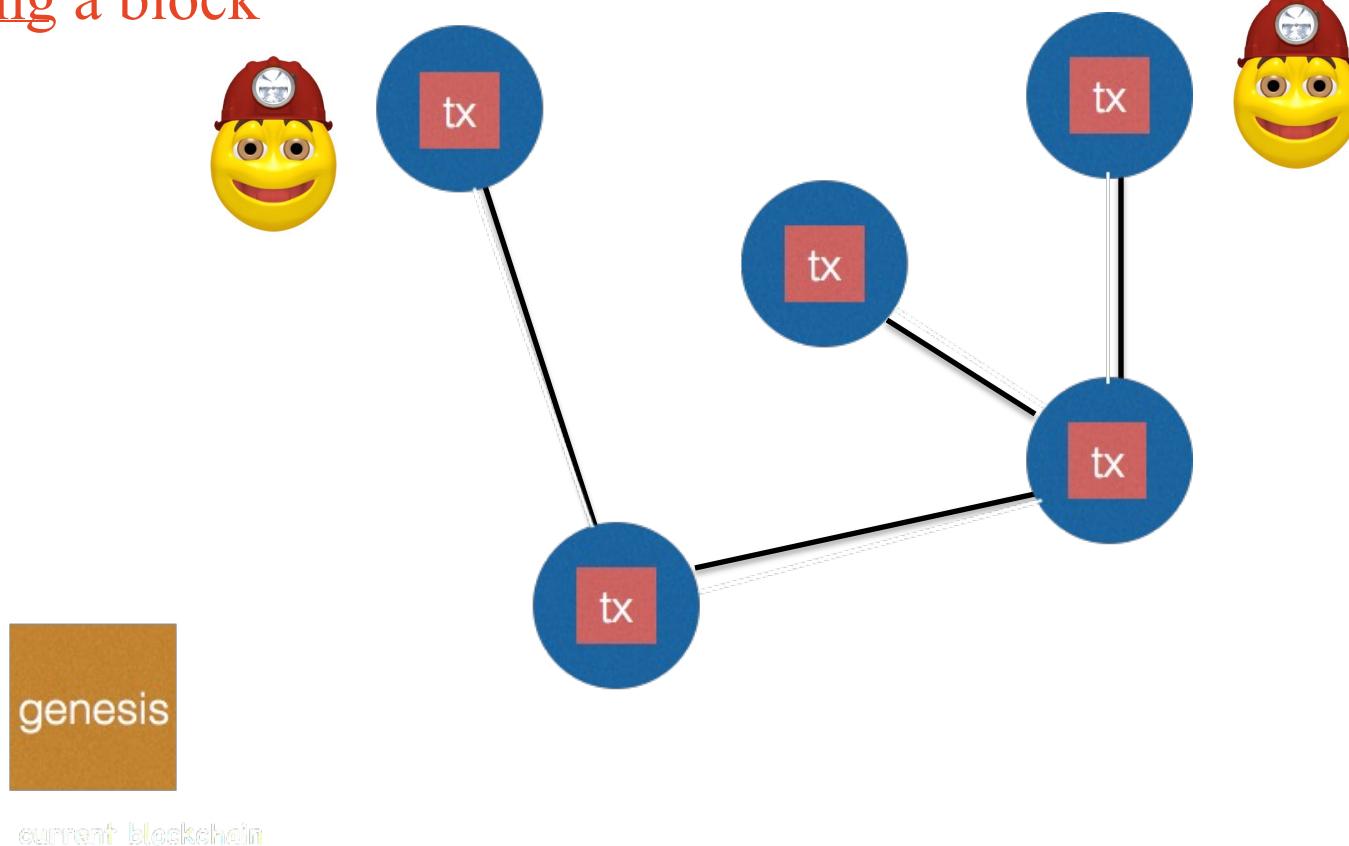
## Broadcast



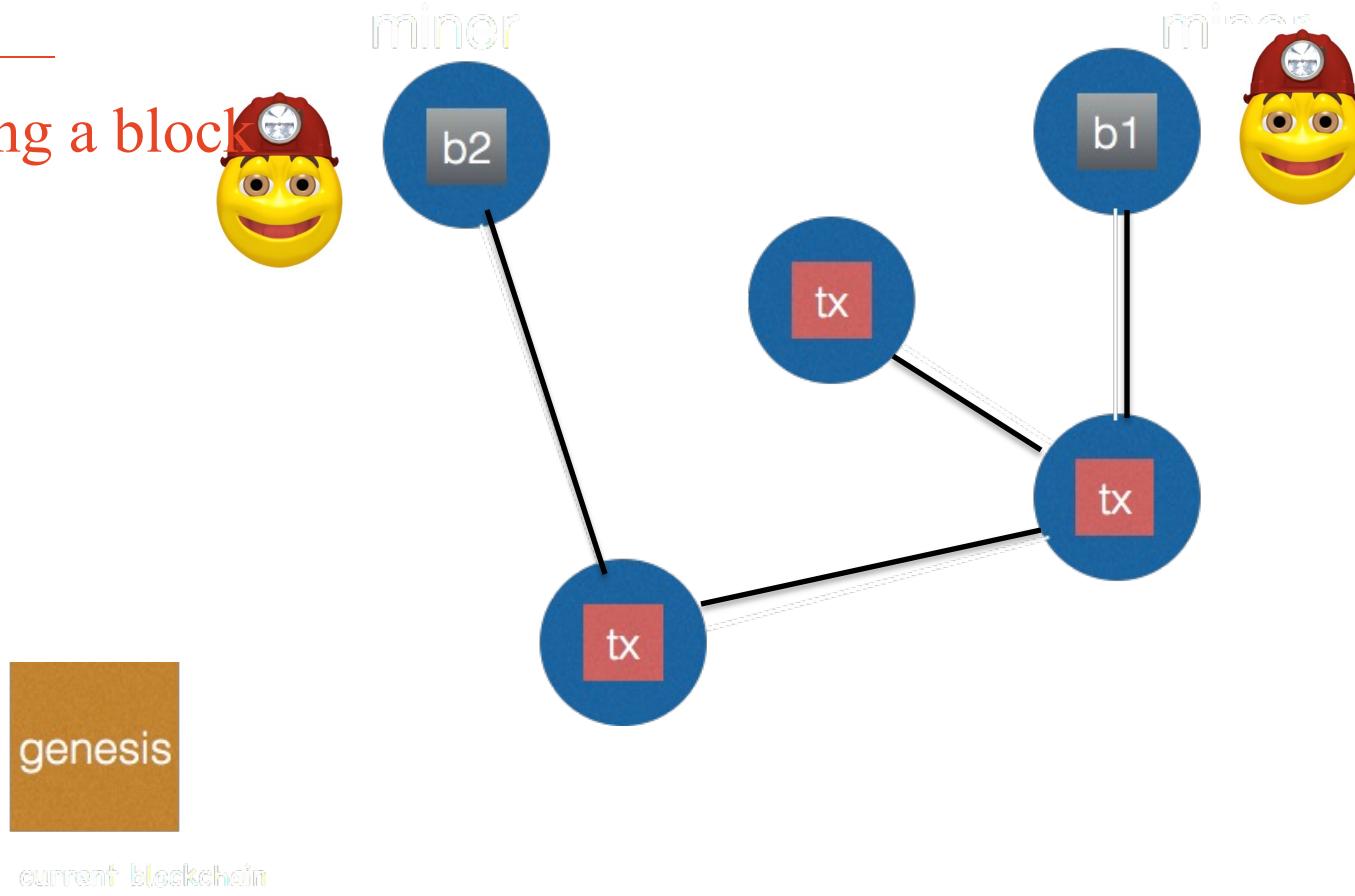
# Broadcast



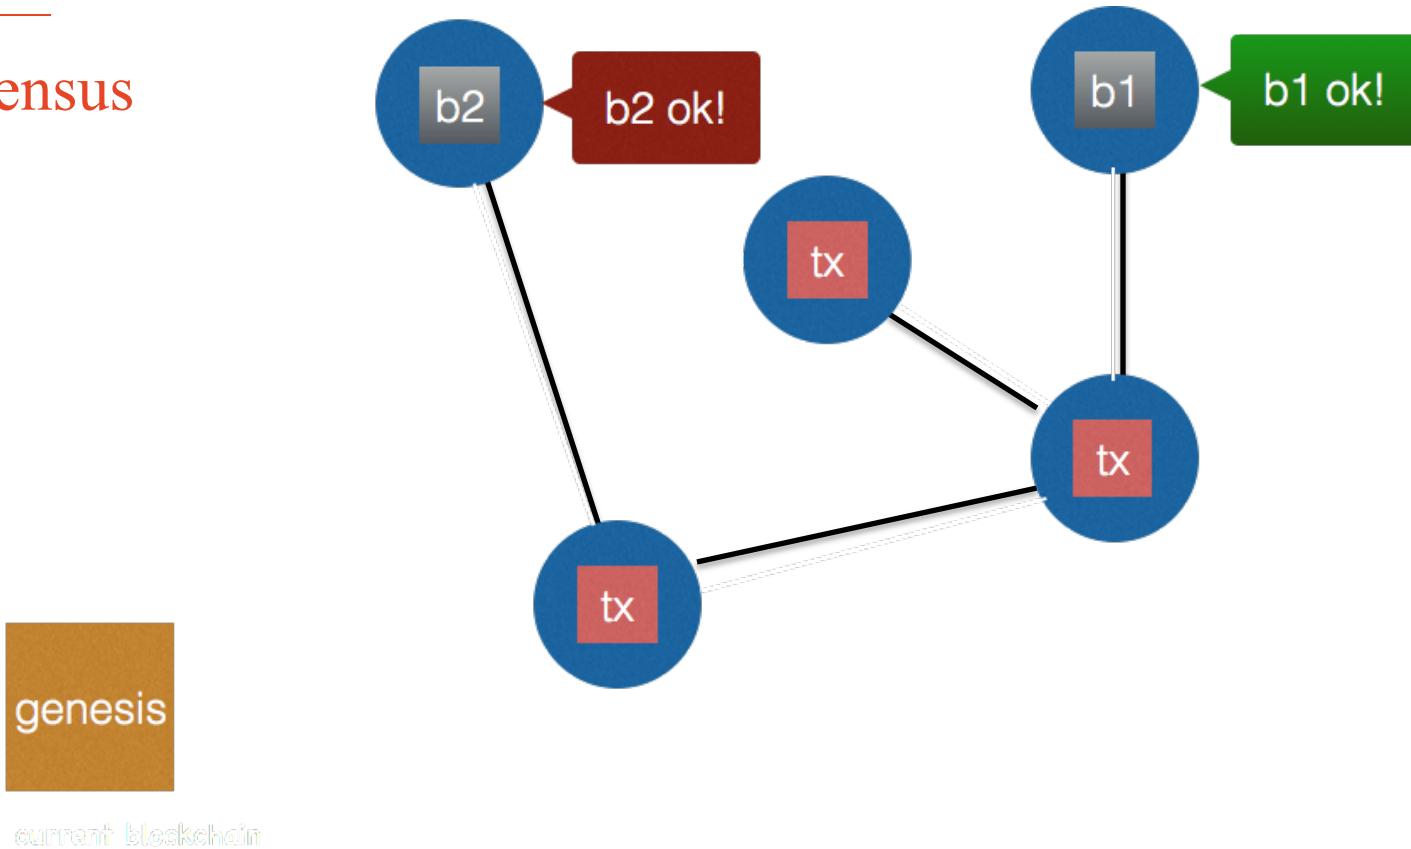
# Mining a block



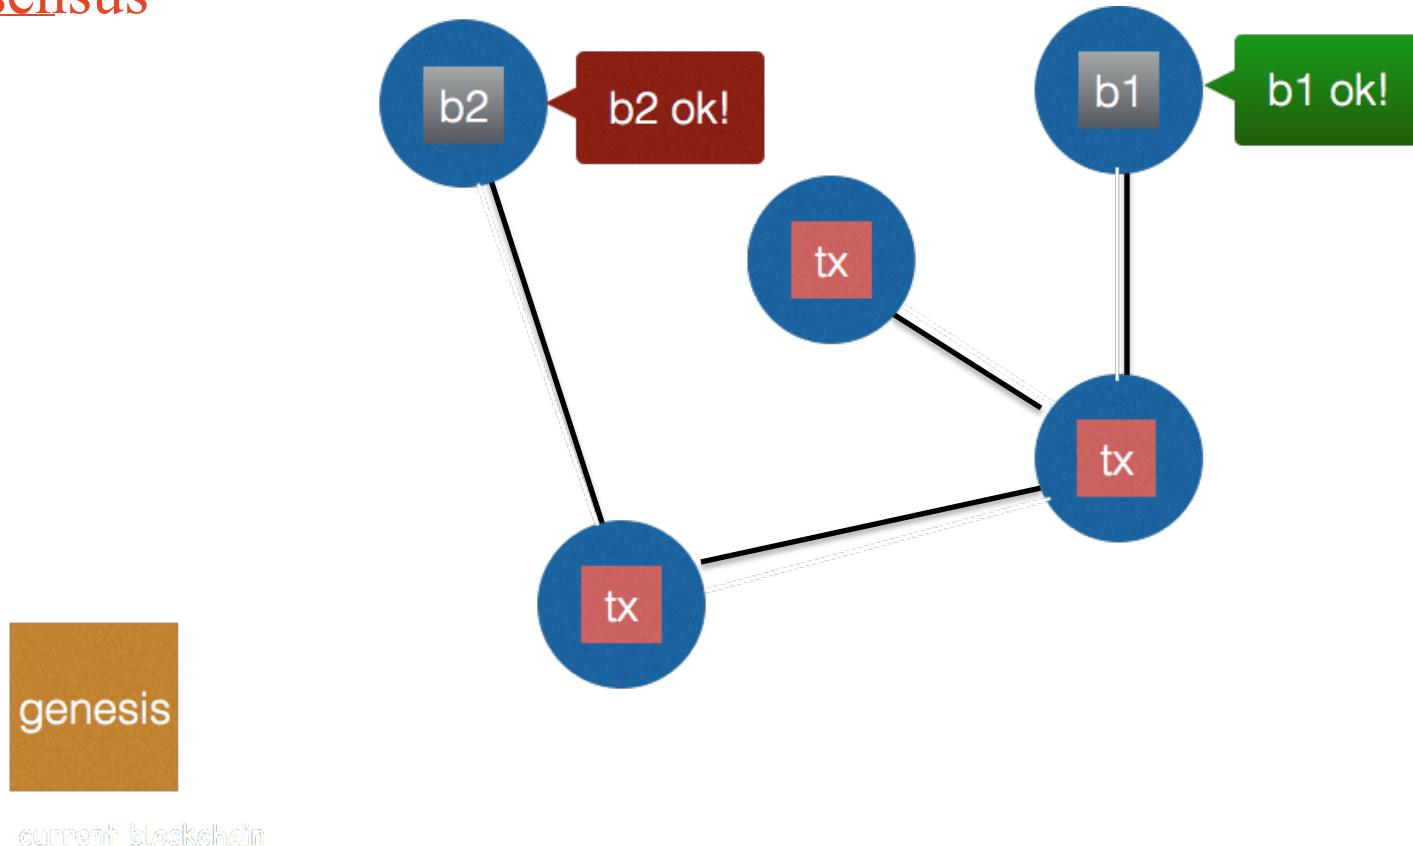
## Mining a block



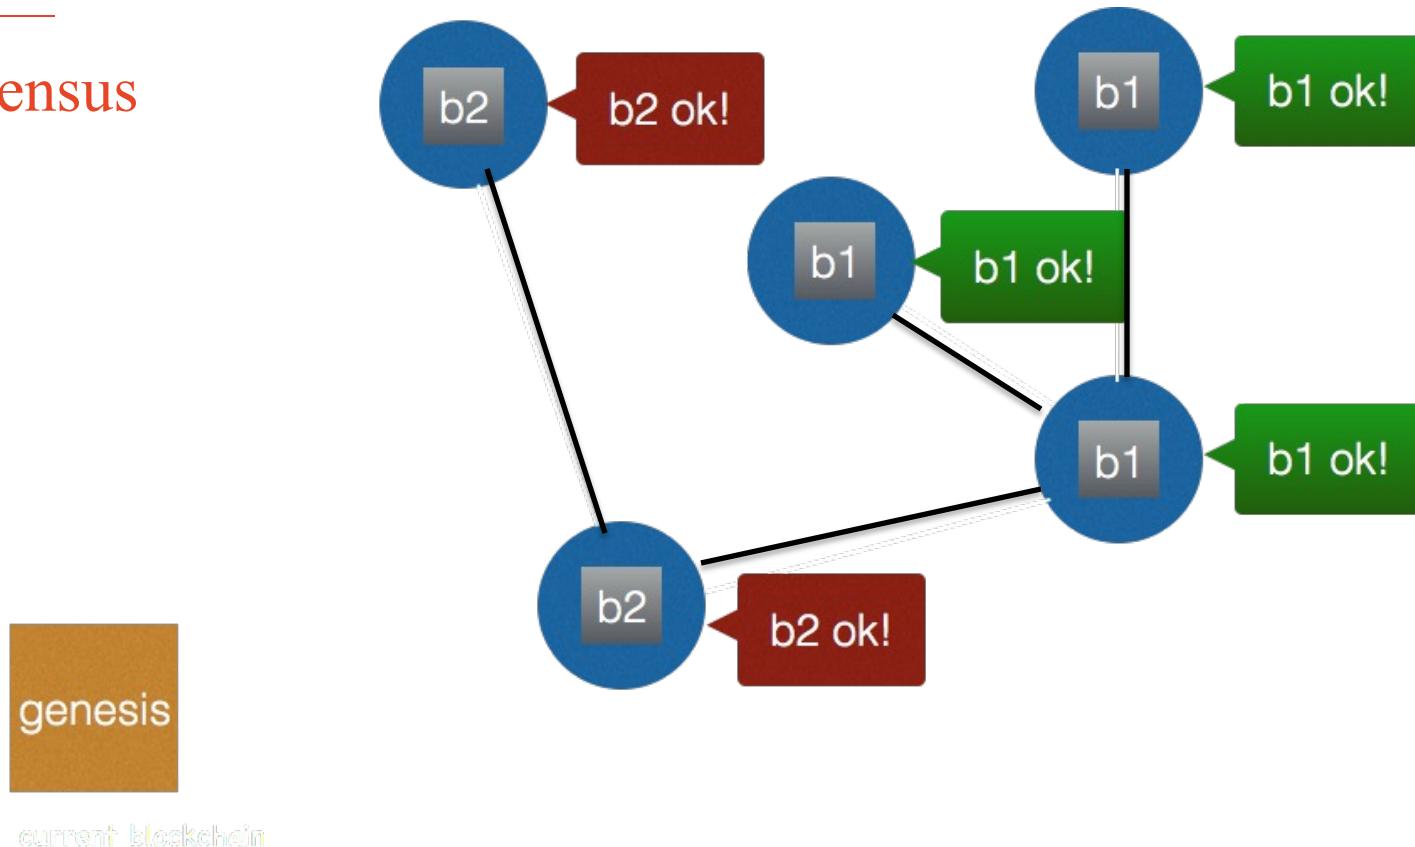
# Consensus



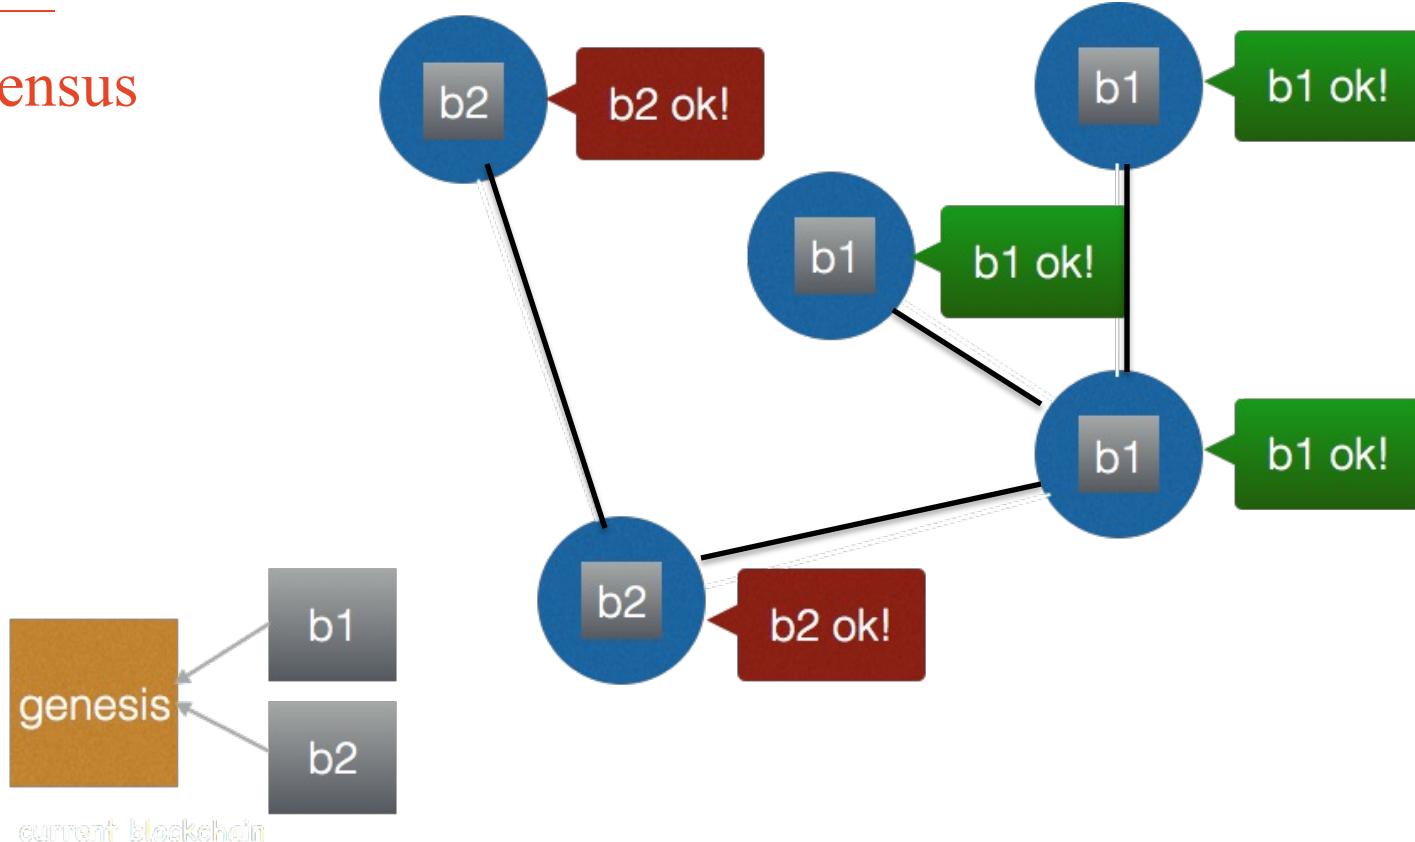
# Consensus



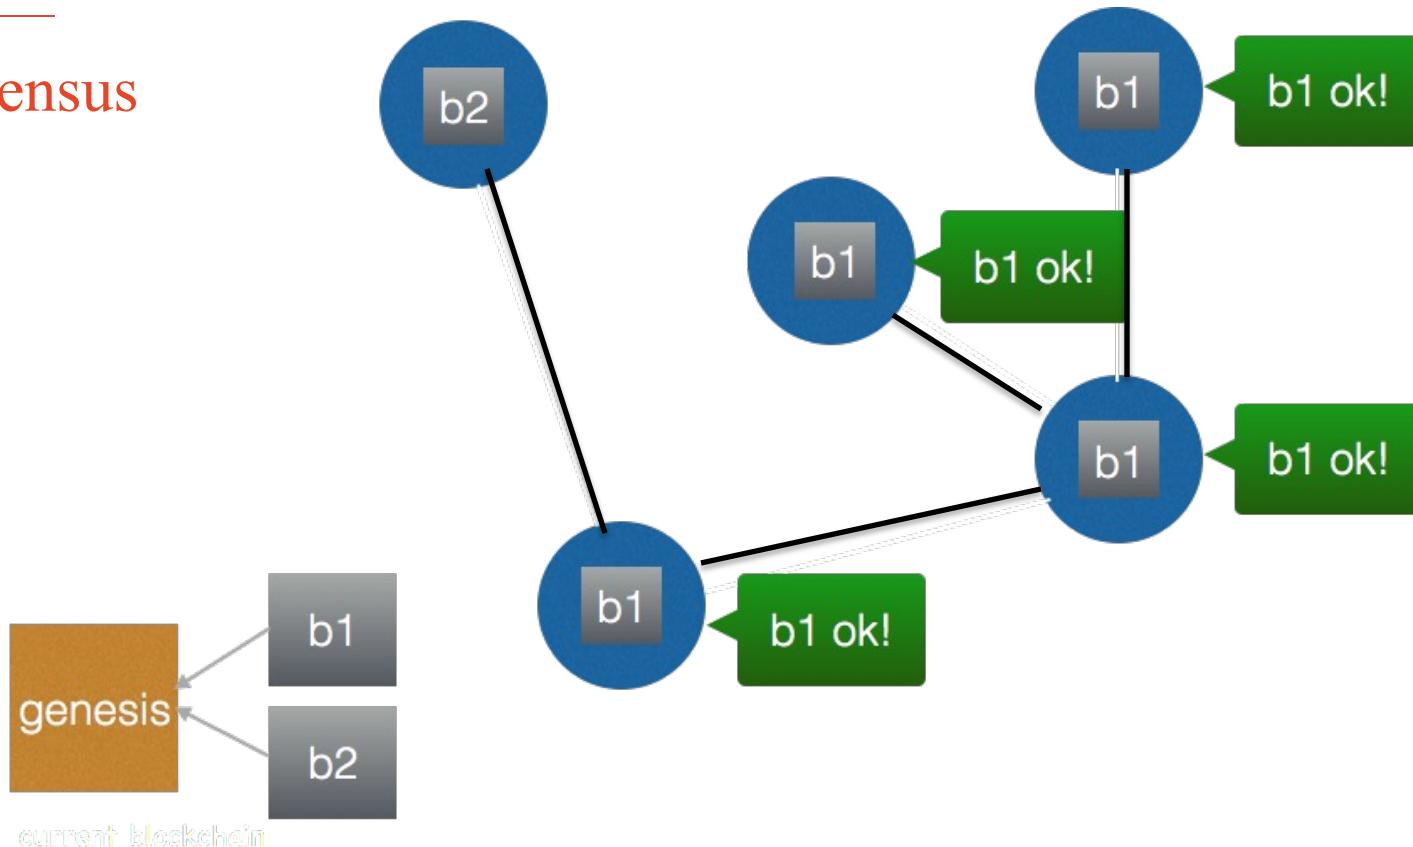
# Consensus



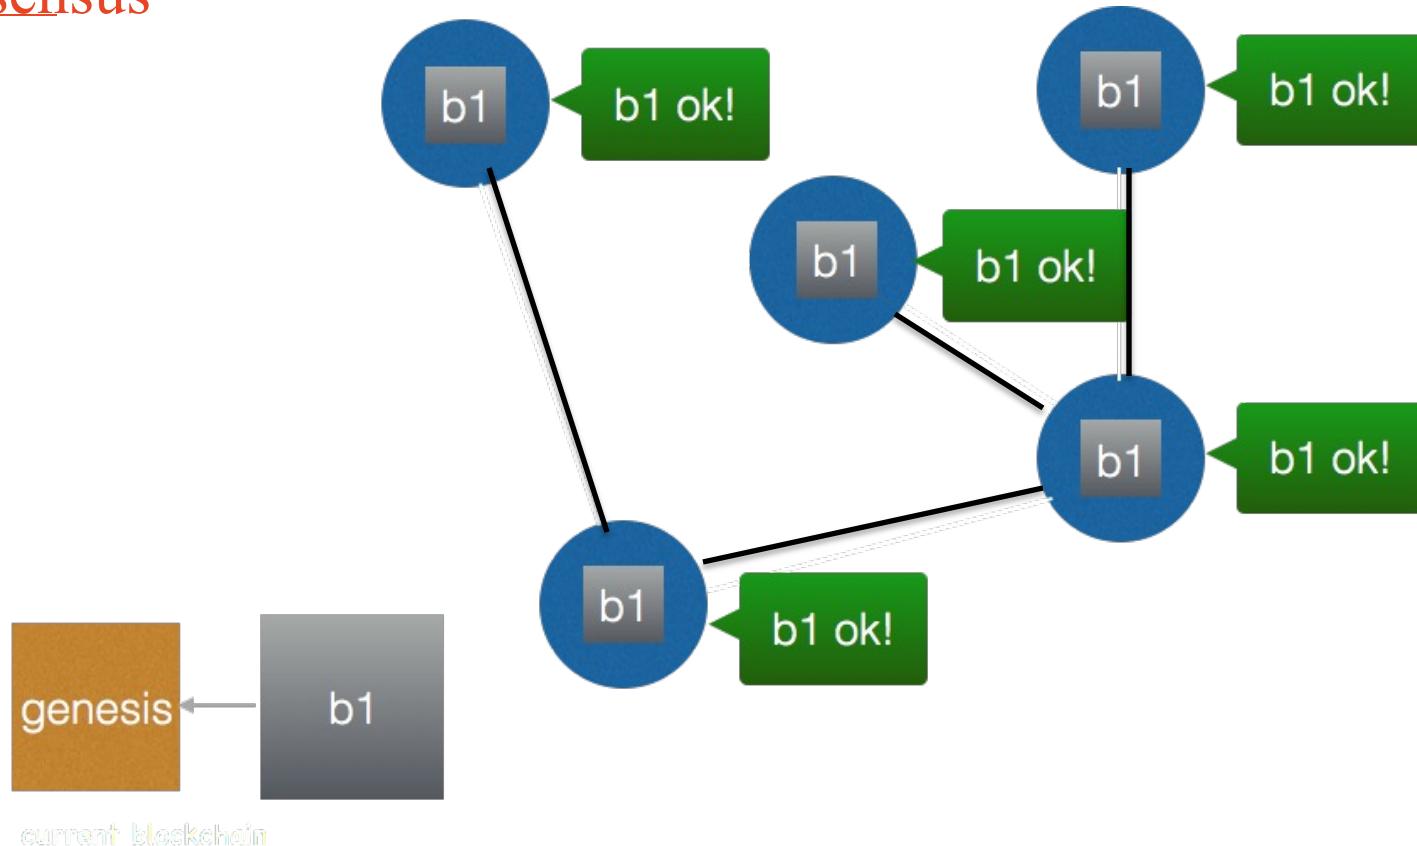
## Consensus



## Consensus



# Consensus



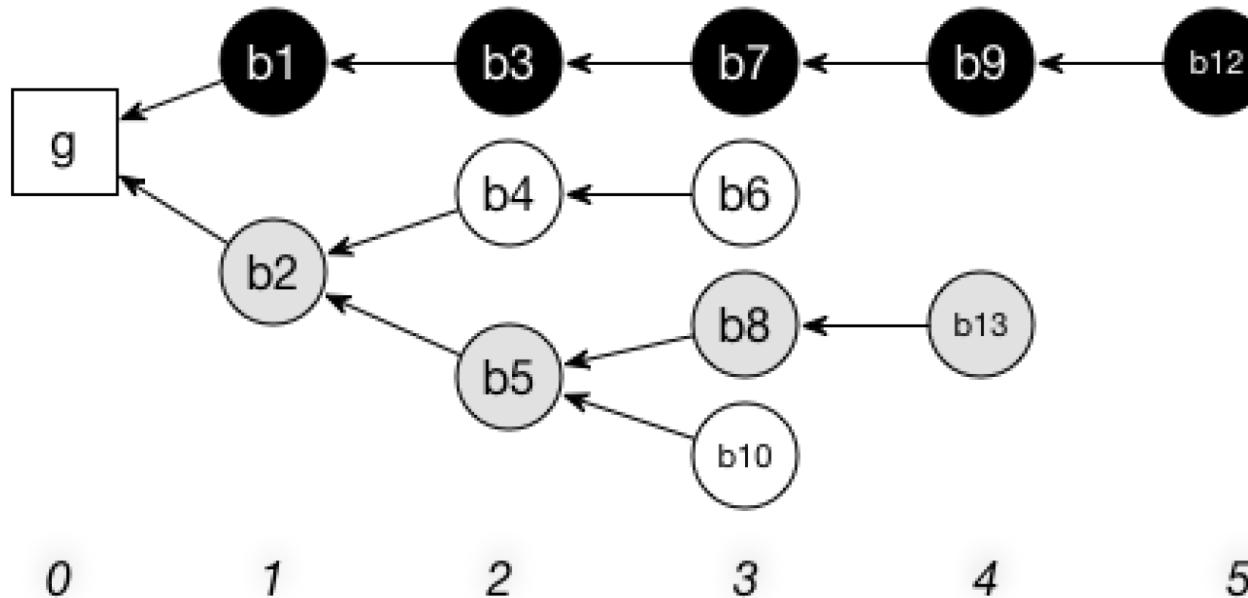
---

# Recovering from disagreements

# Resolving a fork

---

Assume that the current blockchain state is this DAG



# Bitcoin's consensus chooses the deepest branch

## State

$\langle B_i, P_i \rangle$  the local blockchain view

## Each peer of the blockchain executes:

Receive blocks  $\langle B_j, P_j \rangle$  from j

$$B_i = B_i \cup B_j$$

$$P_i = P_i \cup P_j$$

## depth(b):

if  $\text{children}(b) = \emptyset$  then return 1

else return  $1 + \max_{c \in \text{children}(b)} \text{depth}(c)$

## Prune shortest branches at i

$b = \text{genesis-block}(B_i)$

while  $b.\text{next} \neq \perp$

block =  $\operatorname{argmax}_{c \in \text{children}(b)} \{ \text{depth}(c) \}$   
}

$$B = B \cup \{\text{block}\}$$

$$P = P \cup \{\langle \text{block}, b \rangle\}$$

$$b = \text{block}$$

$$\langle B_i, P_i \rangle = \langle B, P \rangle$$

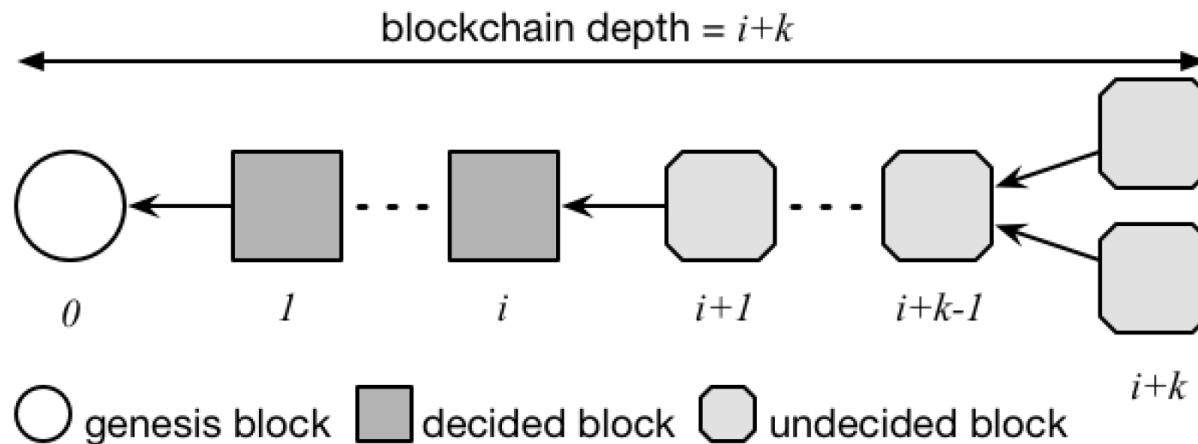
[Nak08] S. Nakamoto, "Bitcoin: a peer-to-peer electronic cash system," 2008,  
<http://www.bitcoin.org>.



# When is a transaction committed?

# Committed transaction

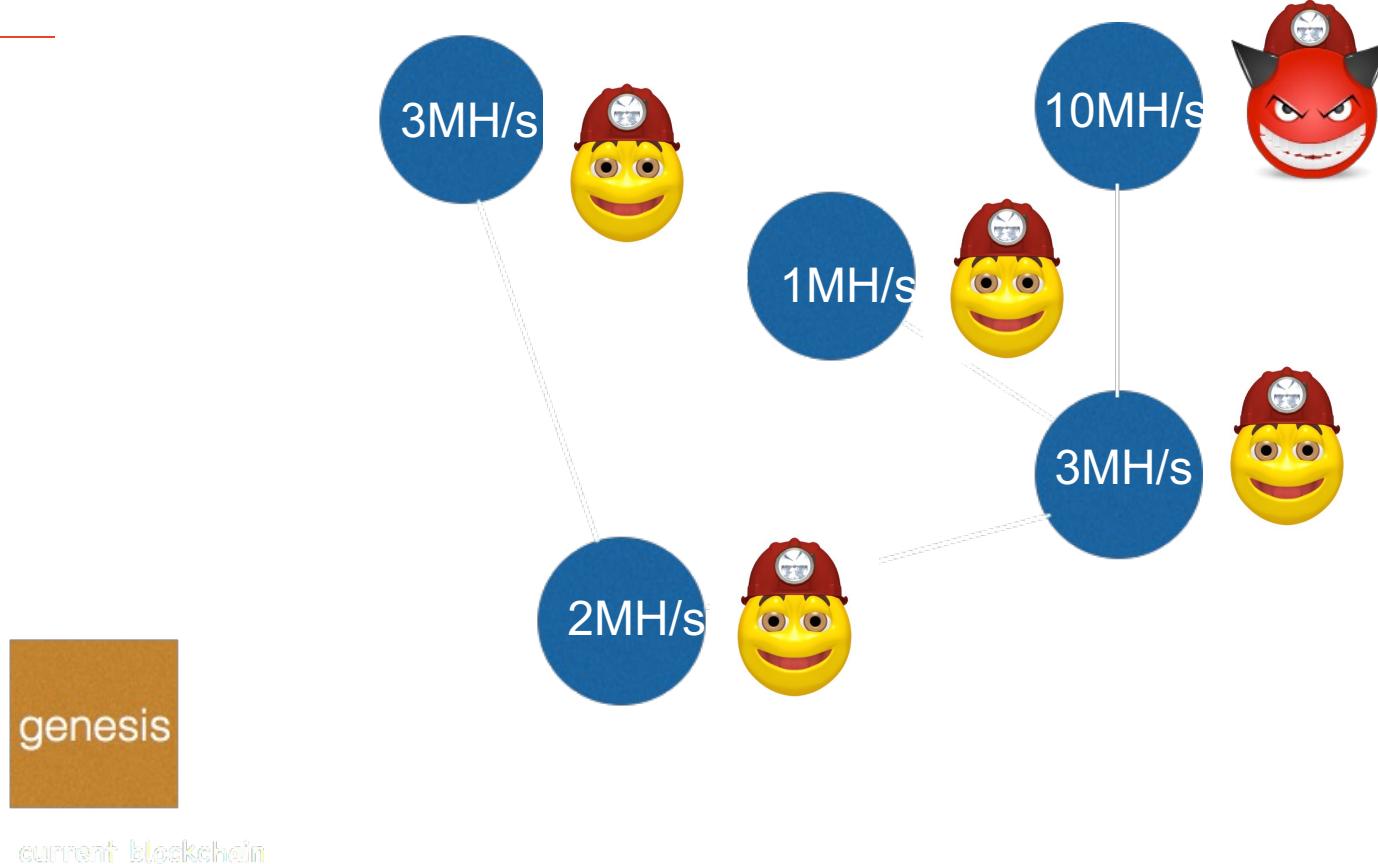
- Given a blockchain with parameter  $k$ , a block at index  $i$  is *decided* when the chain depth reaches  $i+k$
- A transaction is *committed* if it belongs to a decided block



# The 51% attack

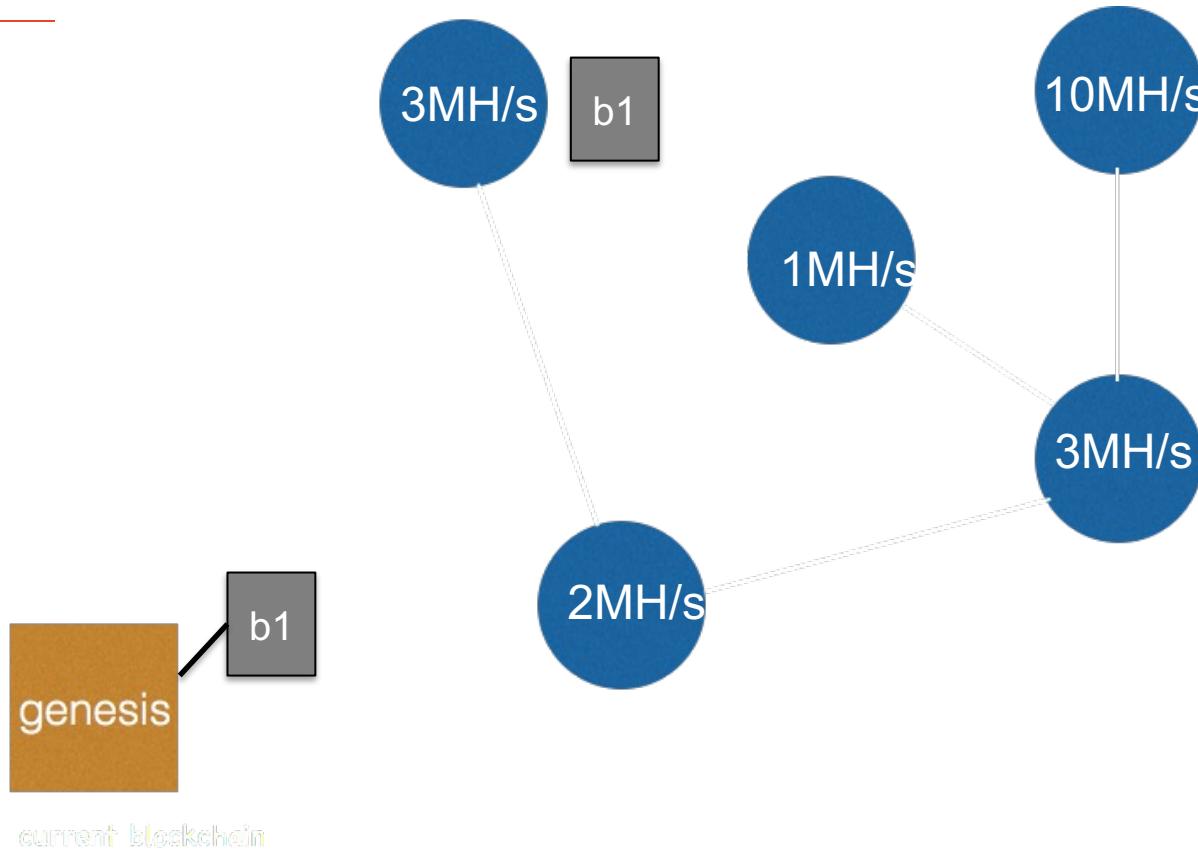
# Impact of mining power

---



# Impact of mining power

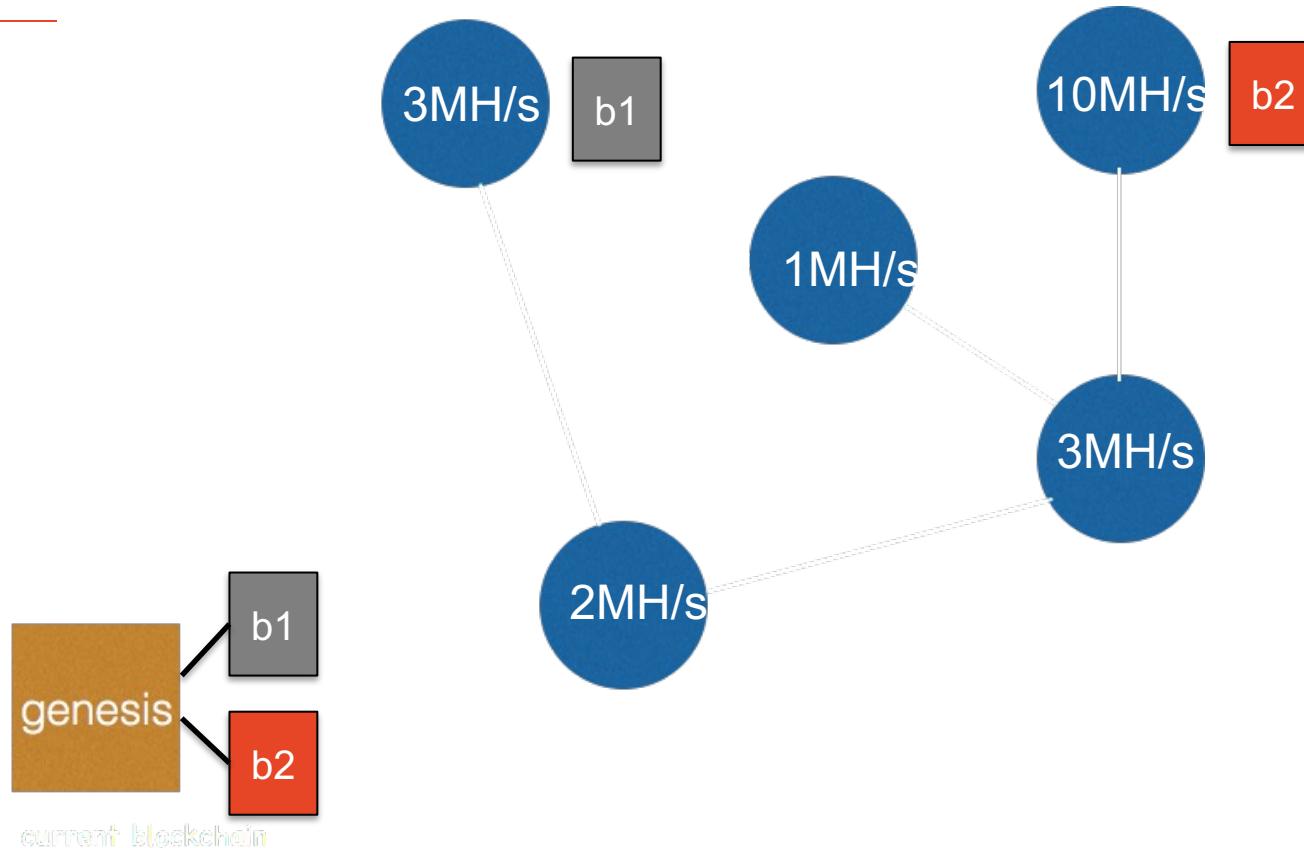
---



current blockchain

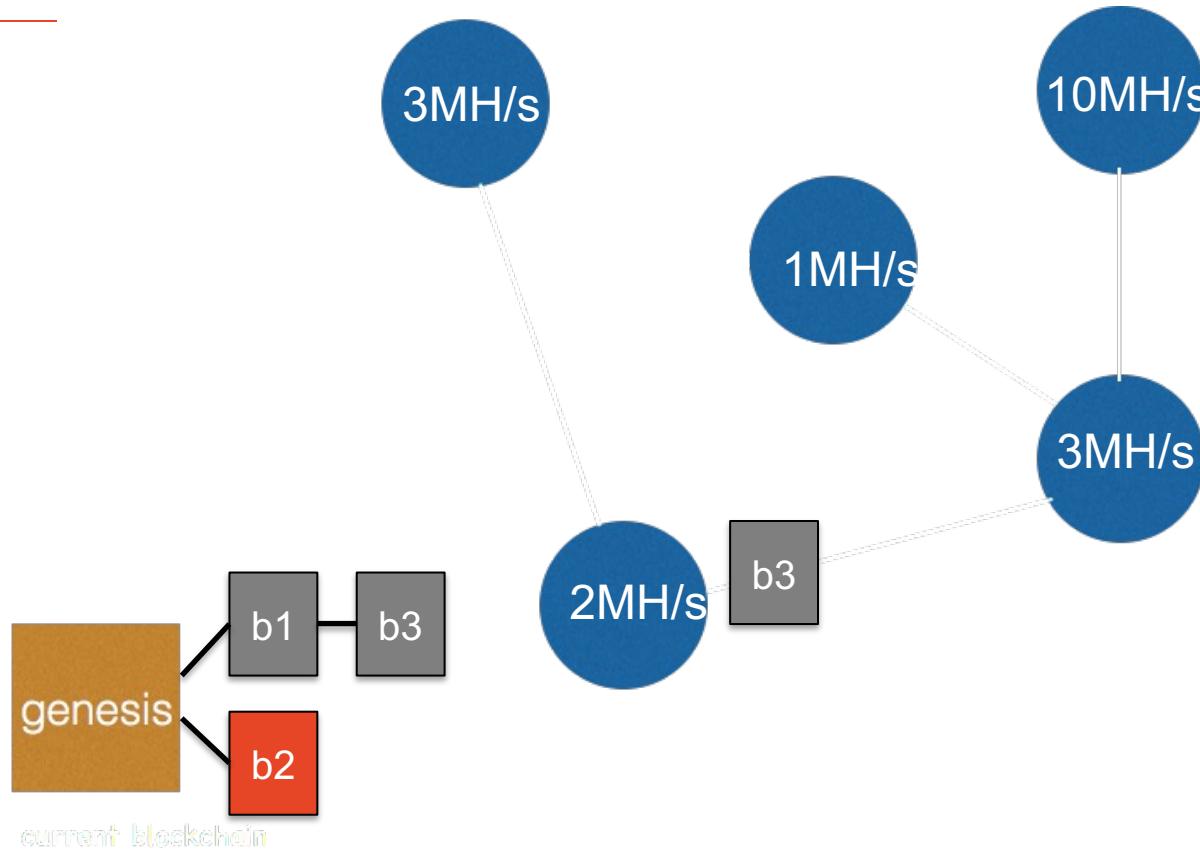
# Impact of mining power

---

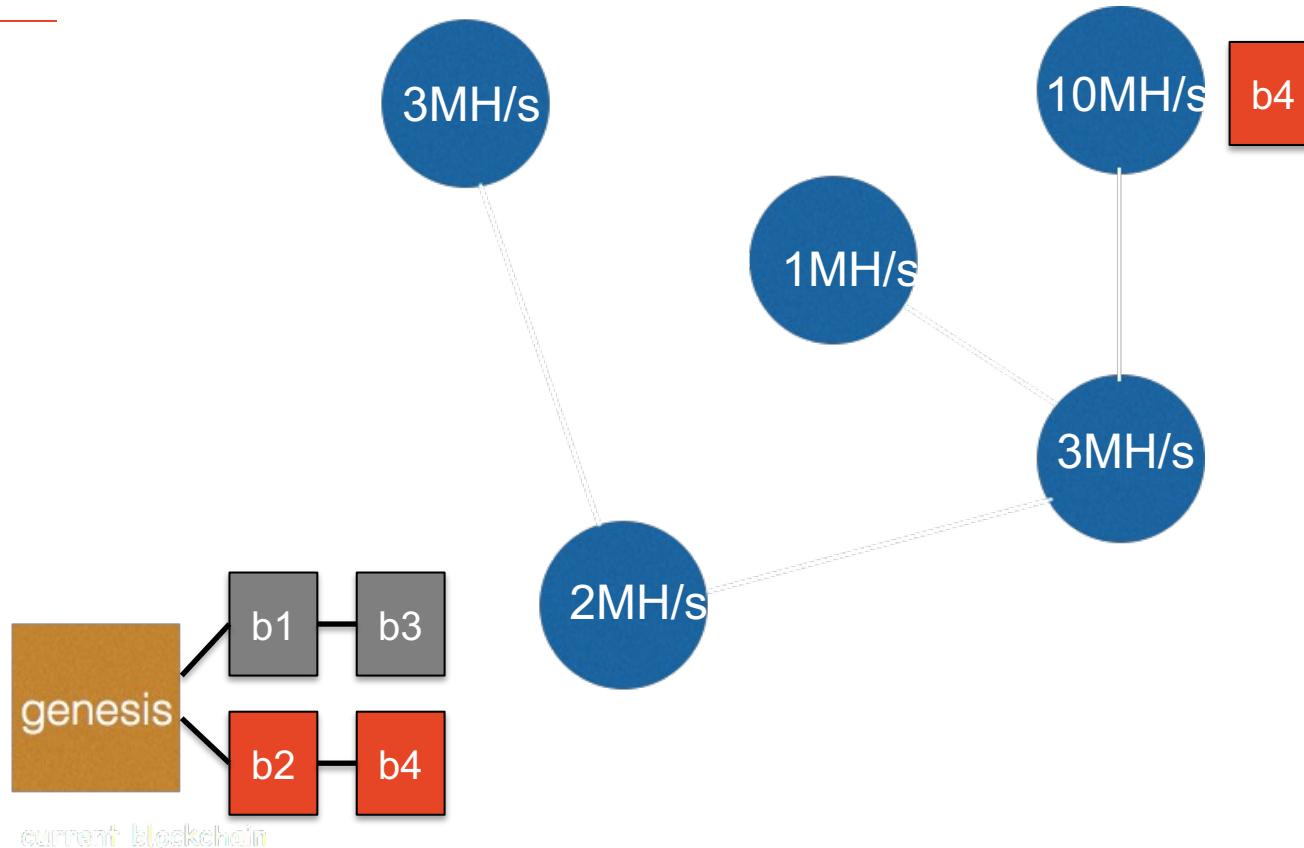


# Impact of mining power

---

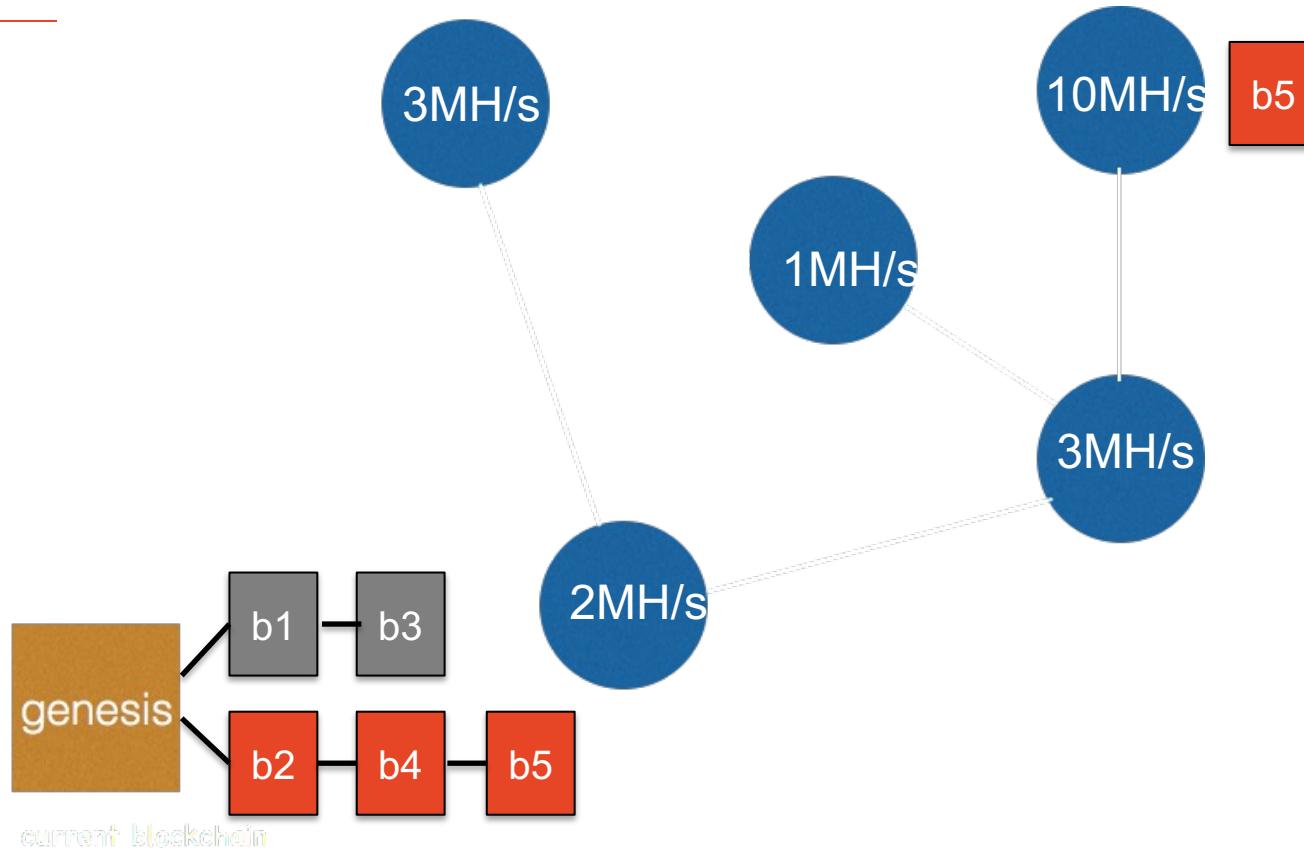


# Impact of mining power

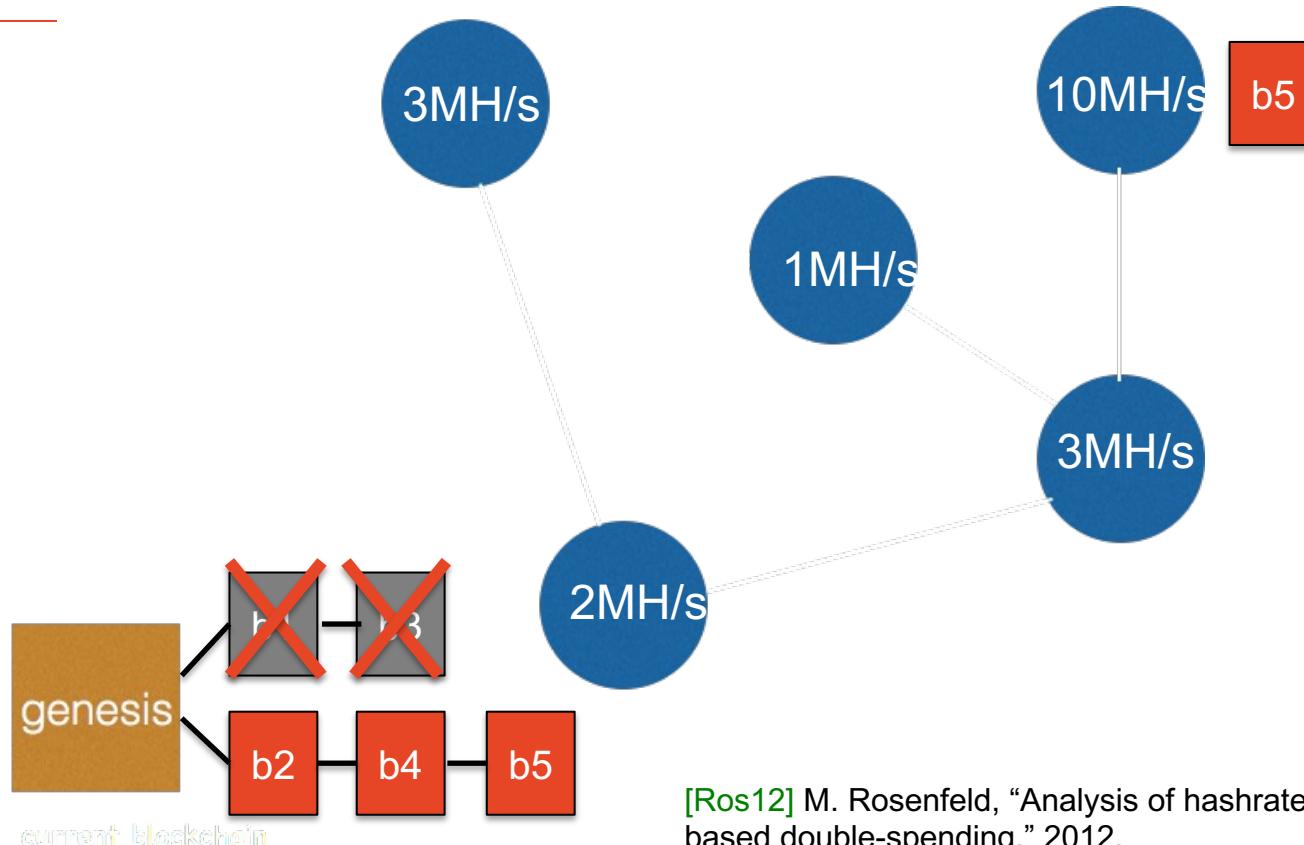


# Impact of mining power

---



# Impact of mining power

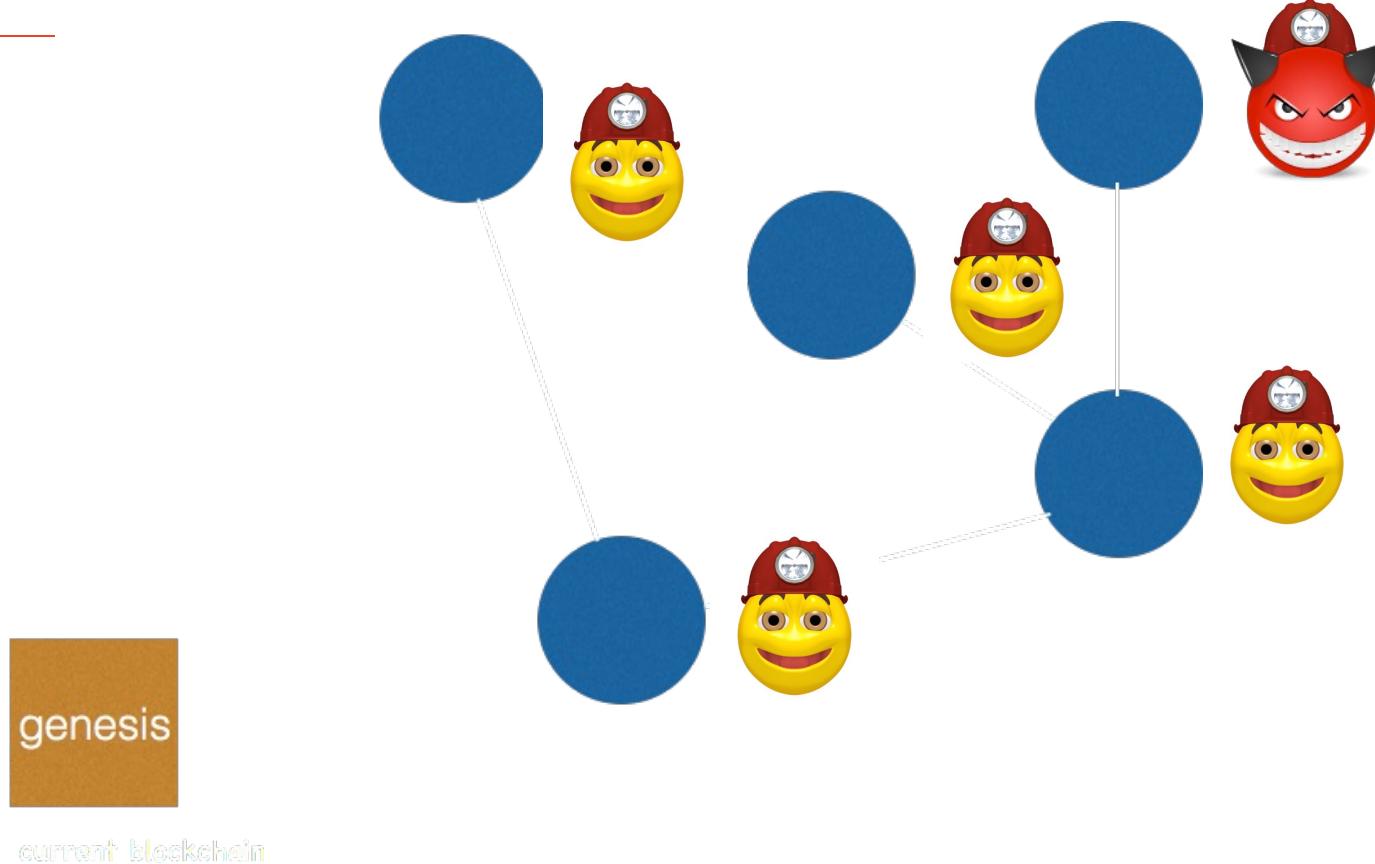


[Ros12] M. Rosenfeld, “Analysis of hashrate-based double-spending,” 2012.

# The impact of network delays

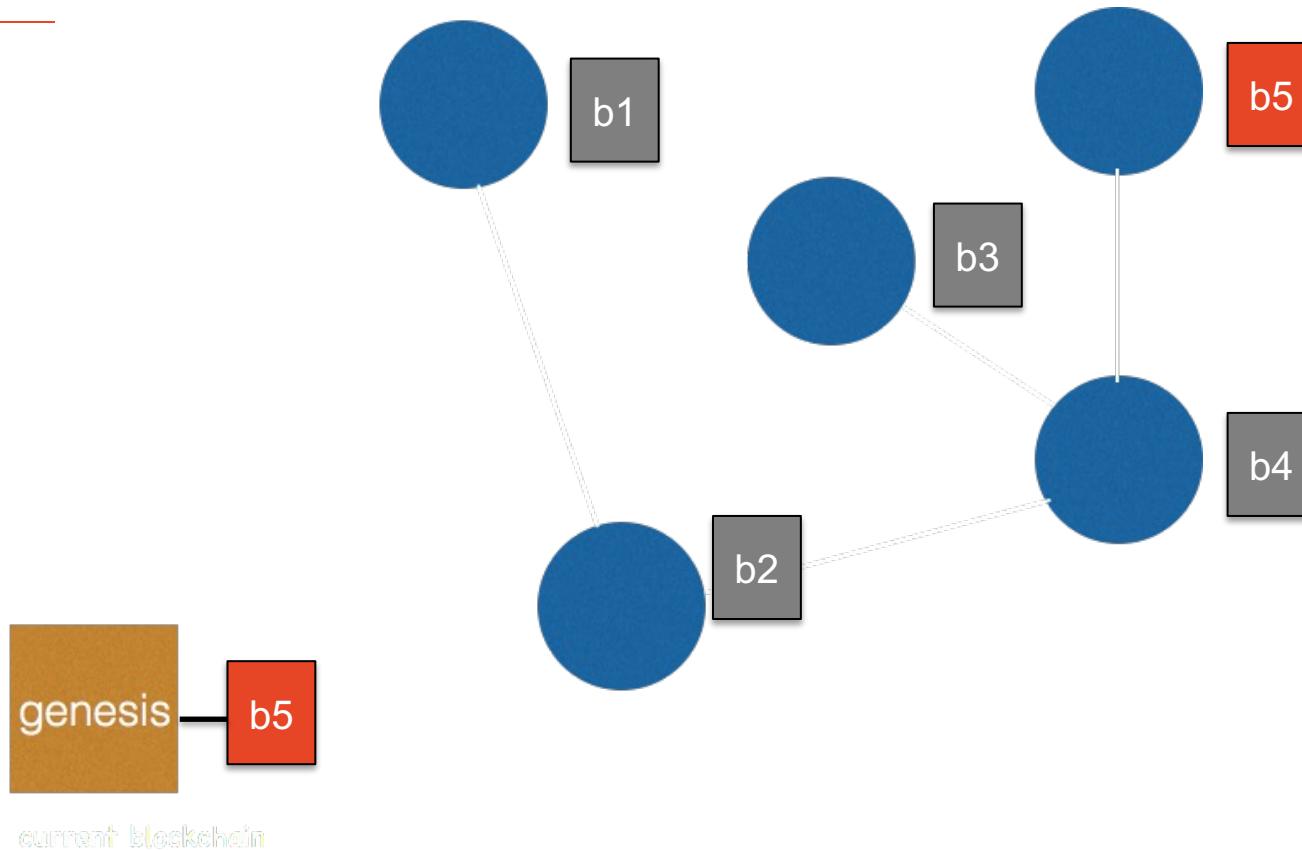
# Impact of communication delay

---



# Impact of communication delay

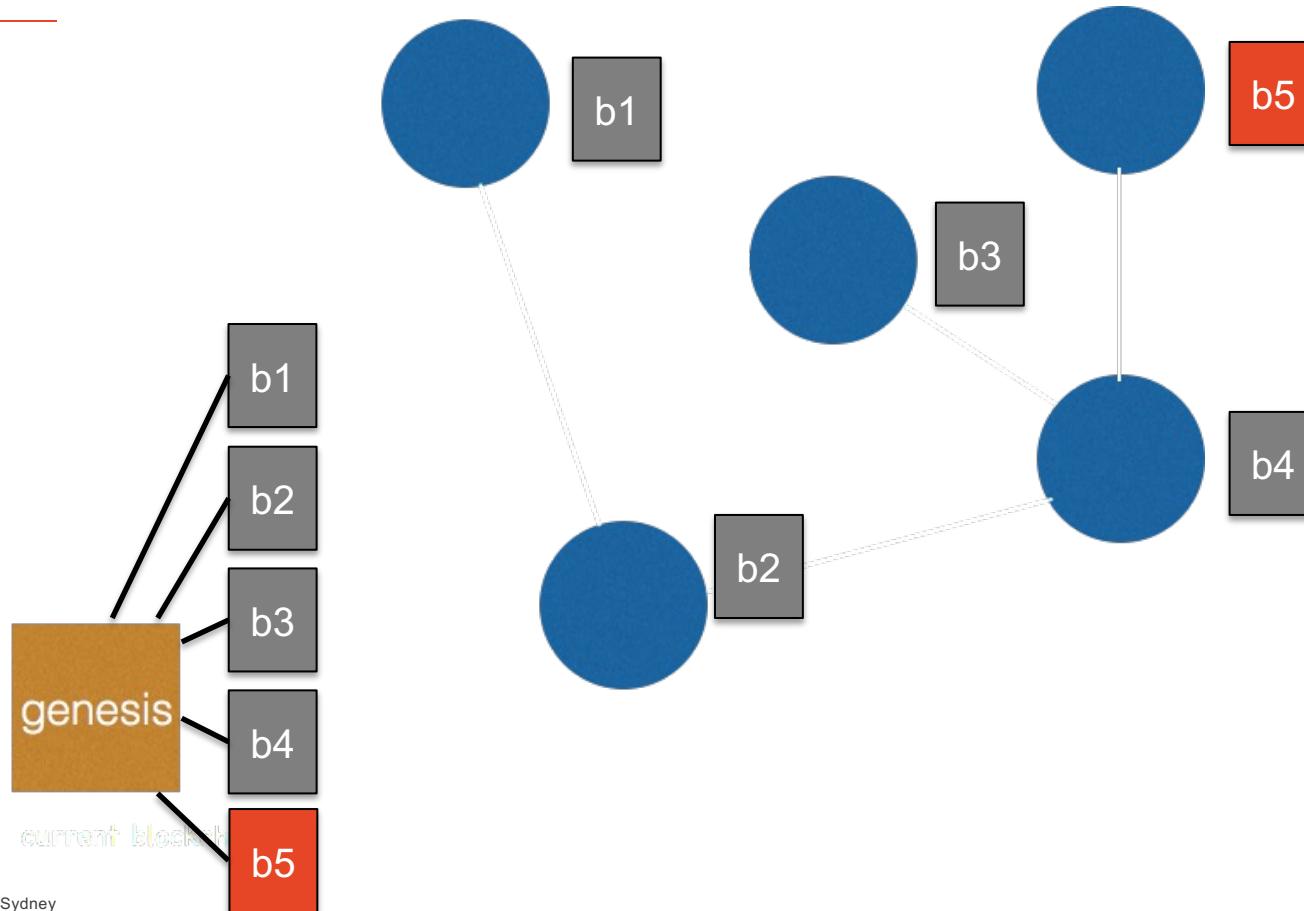
---



current blockchain

# Impact of communication delay

---



# Ethereum and GHOST

# Ethereum consensus was inspired by the choice of the greedily heaviest observable subtree (GHOST)

## State

$\langle B_i, P_i \rangle$  the local blockchain view

## Each peer of the blockchain executes:

Receive blocks  $\langle B_j, P_j \rangle$  from j

$$B_i = B_i \cup B_j$$

$$P_i = P_i \cup P_j$$

## num-desc(b):

if  $\text{children}(b) = \emptyset$  then return 1  
else return  $1 + \sum_{c \in \text{children}(b)} \text{num-desc}(c)$

## Prune lightest branches at i

$$b = \text{genesis-block}(B_i)$$

while  $b.\text{next} \neq \perp$

$$\text{block} = \text{argmax}_{c \in \text{children}(b)} \{\text{num-desc}(c)\}$$

$$B = B \cup \{\text{block}\}$$

$$P = P \cup \{\langle \text{block}, b \rangle\}$$

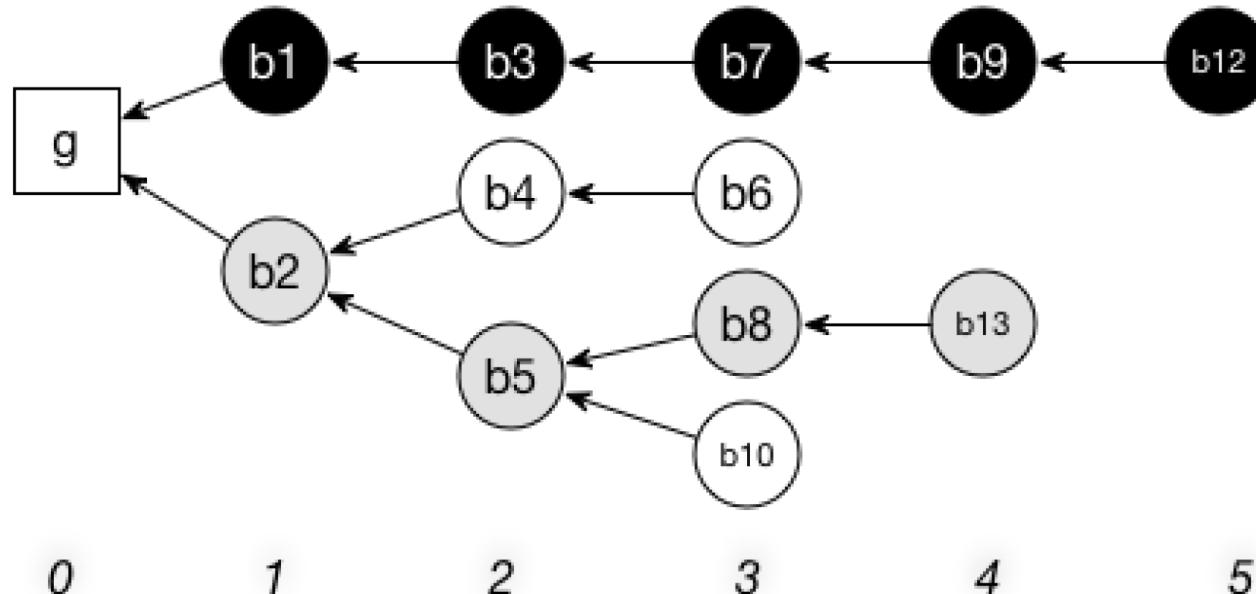
$$b = \text{block}$$

$$\langle B_i, P_i \rangle = \langle B, P \rangle$$



[SZ15] Y. Sompolinsky and A. Zohar, “Secure high-rate transaction processing in bitcoin,” in Financial Cryptography and Data Security FC’15, 2015, pp. 507–527.

## Bitcoin vs. Ethereum/GHOST



---

# Ethereum branch selection

Ethereum codebase evolved dramatically over the past two years

While originally stated as inspired by GHOST, Ethereum's branch selection algorithm is currently different:

Ethereum selects the branch with the highest total difficulty, which is the sum of the difficulties of all the cryptopuzzles of its last block and its ancestors.



# Quiz

---

# Question 1

What is the PoW mechanism useful for?

---

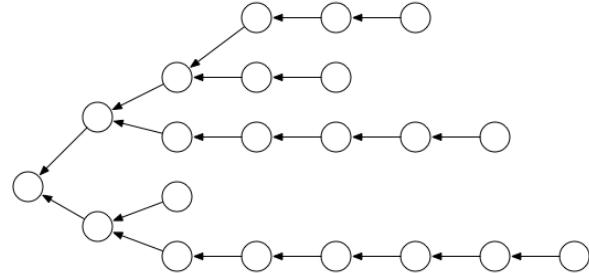
## Question 2

What is the PoW mechanism useful for?

What is its major drawback?

## Question 3

What are the blocks of the branch selected by Bitcoin?



---

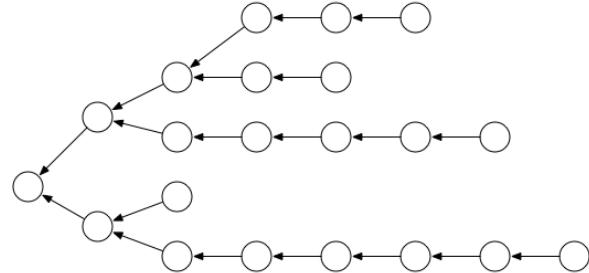
# Consensus

# Consensus without failures

---

## Question 4

What are the blocks of the branch selected by Bitcoin?



What are the blocks of the branch selected by the GHOST protocol?

# Consensus without failures

Each process:

1. Broadcasts its own value
2. Decides on the minimum of all received values

**State of process  $p_i$ :**

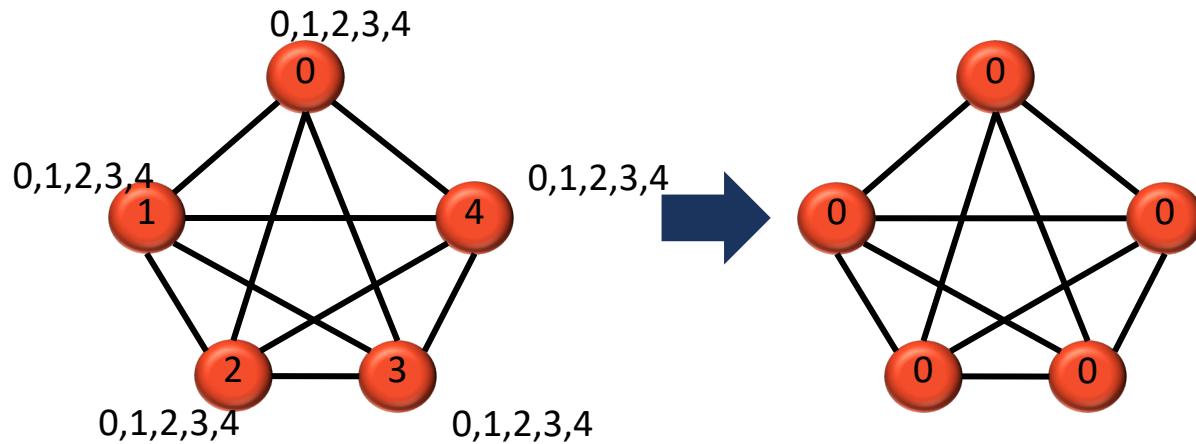
$V$ , the set of known values, initially  $\{x\}$  the proposed value of  $p_i$

**Basic consensus algorithm for  $p_i$ :**

```
send {v in V that  $p_i$  has not sent yet} to all  $p_j$       // broadcast values just found
receive  $S_j$  from all  $p_j$ ,  $j \neq i$                       // receive values
 $V = V \cup S_j$  for all  $j$                                 // records the
values
 $y := \min(V)$                                          // at the end, output min
```

# Consensus without failures

- Broadcast values and decide on minimum → Consensus!
- Validity is satisfied because the decided value is a proposed one
- Agreement is satisfied because all decide the same



# Quiz

# Complexity of Consensus w/o failure algorithm?

- Number of messages needed to complete?
- Number of bits exchanged needed to complete?
- Time needed to complete?

**State of process  $p_i$ :**

$V$ , the set of known values, initially  $\{x\}$  the proposed value of  $p_i$

**Basic consensus algorithm for  $p_i$ :**

send  $\{v \in V \text{ that } p_i \text{ has not sent yet}\}$  to all  $p_j$  // broadcast values just found

receive  $S_j$  from all  $p_j j \neq i$  // receive values

$V = V \cup S_j$  for all  $j$  // records the values

$y := \min(V)$  // at the end, output min

# Complexity of Consensus w/o failure algorithm?

**Message complexity.** Note that only one communication round is needed!

Each process:

1. Broadcast own value →  $O(n^2)$
2. Decide on the minimum of all received values

The message complexity is thus  $O(n^2)$ .

**Communication complexity.**

Each message is an integer, leading to a communication complexity of  $O(bn^2)$ , where b is the number of bits necessary to encode any value

**Time complexity.**

As all messages can be exchanged in parallel, it only takes  $O(1)$  message delay to terminate.

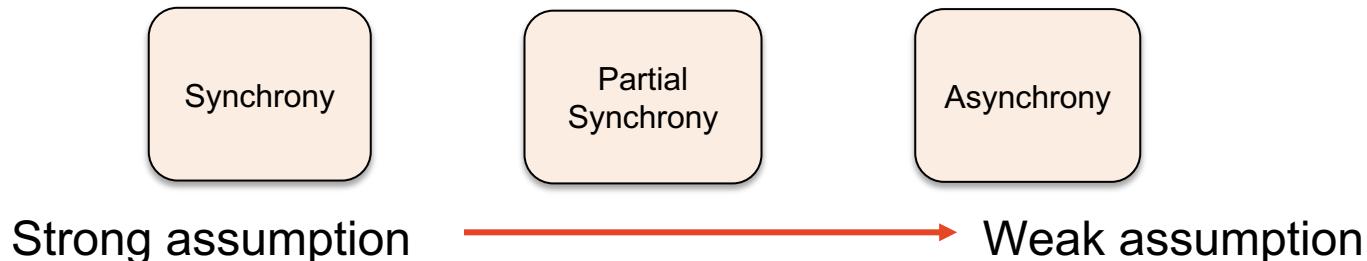
# Synchrony and asynchrony

# Synchrony vs. Asynchrony

So far, we assumed *synchrony*: that all messages are received within a bounded time that is known from the algorithms

What will happen to the previous algorithms if we relax this assumption?

1. *Partial synchrony*: the bound on message delay is unknown
2. *Asynchrony*: message delays are unbounded



# Synchrony vs. Asynchrony

Synchrony

Synchrony helps a lot actually:

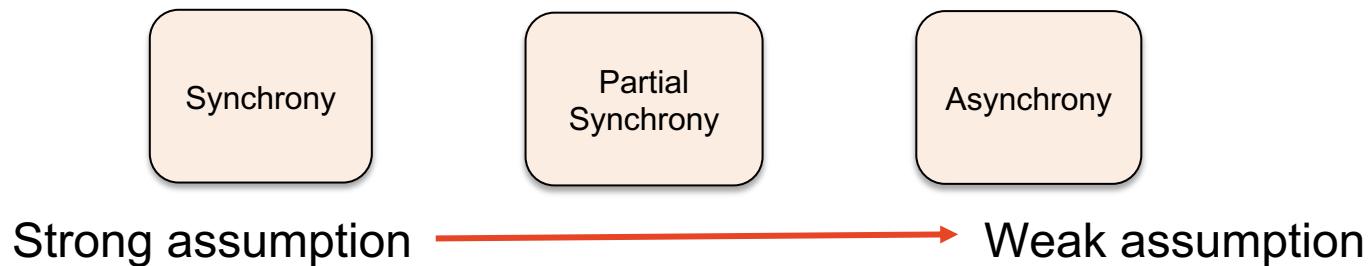
- If you know how much time a message takes to be transmitted and how much time a process needs to treat a message and sends a response, then one can detect:
  - Crashed processes or
  - Message losses
- Otherwise, you cannot distinguish whether a message is in-transit and will arrive eventually, it will never arrive.

Partial  
Synchrony

Asynchrony

## Partial synchrony is a stronger assumption than asynchrony

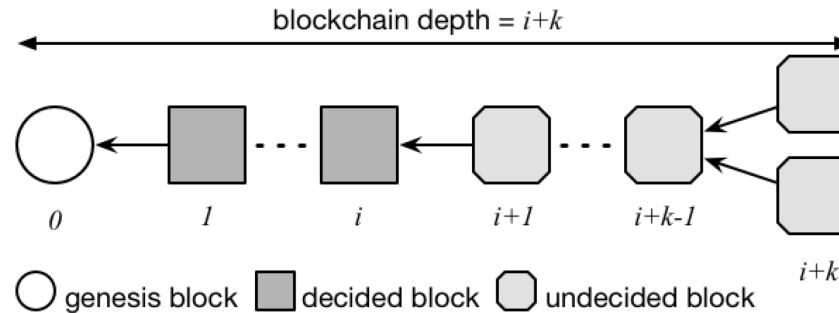
If a protocol does not work under partial synchrony then it cannot work under asynchrony as partial synchrony is a stronger assumption



# Quiz

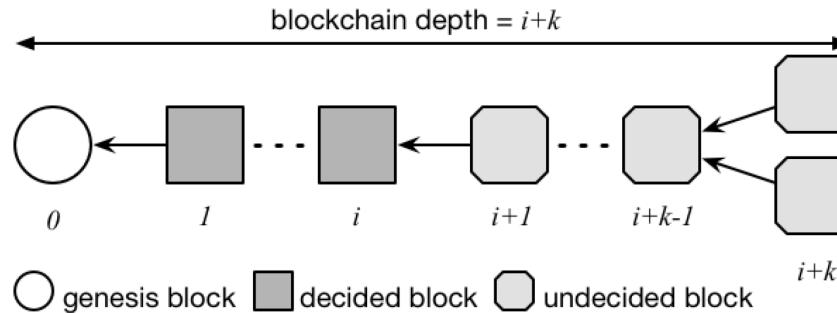
## Question 1

What assumption is needed here for Ethereum to work?



# Question 1

What assumption is needed here for Ethereum to work?



## Synchrony

- Without synchrony, no one can distinguish a committed transaction from a transaction that is not yet committed and that can possibly be rolled back later on

---

# Robustness

# Consensus with crash failures

# Crash Tolerant Consensus

- Assumptions
  - Synchronous communication using unicast
  - Crash failures
  - A maximum of  $f$  processes can fail

## **State of process $p_i$ :**

$V$ , the set of known values, initially  $\{x\}$  the proposed value of  $p_i$   
 $f$ , the maximum number of processors that can fail

## **Basic consensus algorithm for $p_i$ :**

```
for (k=1; k ≤ f+1; k++) {                                // in each  $f+1$  round
    send {v in V that  $p_i$  has not sent yet} to all  $p_j$ . // broadcast values just found
    receive  $S_j$  from  $p_j$   $j \neq i$                                 // receive values
     $V = V \cup S_j$  for all  $j$                                 // records the values
}
 $y := \min(V)$                                             // at the end, output min
```

# Crash Tolerant Consensus

- Intuition of the proof: at least one iteration (i.e., *round*) without failure
  - Termination: the loop has a **bounded** number of iterations, so  $y$  is assigned
  - Agreement: at the end we have  $V_i = V_j$  for any  $i$  and  $j$  so  $y_i = y_j$
  - Validity: there exists a  $p_j$  such that its input  $x_j$  is the decided value

## State of process $p_i$ :

$V$ , the set of known values, initially  $\{x\}$  the proposed value of  $p_i$   
 $f$ , the maximum number of processors that can fail

## Basic consensus algorithm for $p_i$ :

```
for (k=1; k ≤ f+1; k++) {                                // in each f+1 round
    send {v in V that  $p_i$  has not sent yet} to all  $p_j$   // broadcast values just
    found
    receive  $S_j$  from  $p_j$   $j \neq i$                                 // receive values
     $V = V \cup S_j$  for all  $j$                                 // records the
    values
}
 $y := \min(V)$                                             // at the end, output
```

# Complexity of Crash Tolerant Consensus?

# Complexity of Crash Tolerant Consensus?

**Message complexity.** There are  $f+1$  rounds where  $f < n$ .

The number of messages is at most  $O(n^2)$  in each round.

The total number of messages is thus  $O((f+1) n^2)$ .

**Communication complexity.** Each message contains up to  $n$  values, and each value is represented with  $b$  bits, the communication complexity is  $O(b (f+1) n^3)$  bits.

**Time complexity.** There are  $f+1$  rounds, in each of these rounds, messages can be exchanged in parallel. This leads to a time complexity of  $O(f+1)$  message delay.

# Theorem

At least  $f+1$  rounds are necessary to solve consensus where  $f < n$ .

## **Proof.**

In the worst case scenario, one process fails in each round sending to only one process but not the others.

With  $f+1$  rounds, there is at least one round without failure, in which all messages are successfully delivered, so that all correct processes can decide upon the same value.

# Quiz

# Question

What are the message, communication and time complexities of the Crash Tolerant Consensus algorithm?

## **State of process $p_i$ :**

$V$ , the set of known values, initially  $\{x\}$  the proposed value of  $p_i$   
 $f$ , the maximum number of processors that can fail

## **Basic consensus algorithm for $p_i$ :**

```
for ( $k=1$ ;  $k \leq f+1$ ;  $k++$ ) {  
    send  $\{v \in V \text{ that } p_i \text{ has not sent yet}\}$  to all  $p_j$  // broadcast values just  
    found  
    receive  $S_j$  from  $p_j$   $j \neq i$  // receive values  
     $V = V \cup S_j$  for all  $j$  // records the  
    values  
}  
y := min( $V$ ) // at the end, output
```

# Complexity of Crash Tolerant Consensus?

**Message complexity.** There are  $f+1$  rounds where  $f < n$ .

The number of messages is at most  $O(n^2)$  in each round.

The total number of messages is thus  $O((f+1) n^2)$ .

**Communication complexity.** Each message contains up to  $n$  values, and each value is represented with  $b$  bits, the communication complexity is  $O(b (f+1) n^3)$  bits.

**Time complexity.** There are  $f+1$  rounds, in each of these rounds, messages can be exchanged in parallel. This leads to a time complexity of  $O(f+1)$  message delay.

# Consensus with Byzantine failures

# The Byzantine Consensus Problem

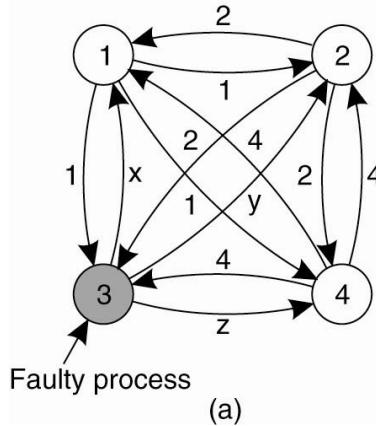
In the consensus problem, the processes propose values and have to agree by deciding one among these values

Consensus Definition 2:

1. Agreement: No two correct processes decide differently
2. Termination: Every correct process eventually decides
3. Validity: if all correct processes propose the same value, then this value is decided moreover, if all processes are correct and some process decides  $v$ , then  $v$  was proposed by some process

# Byzantine Tolerant Consensus

- Assumptions
  - Synchronous communication using unicast
  - Byzantine failures
  - A maximum of  $f < n/3$  processes can fail among  $n$



1 Got(1, 2, x, 4)  
2 Got(1, 2, y, 4)  
3 Got(1, 2, 3, 4)  
4 Got(1, 2, z, 4)

1 Got (1, 2, y, 4)	2 Got (1, 2, x, 4)	4 Got (1, 2, x, 4)
(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)

(b)

(c)

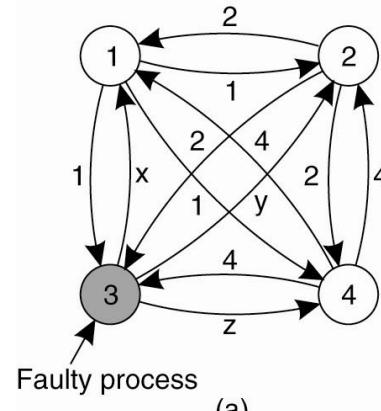
# Byzantine Tolerant Consensus

- Example of solution for  $n=4$  and  $f=1$

- (a) every process sends its input to all others
  - (b) the collected vectors ( $i^{\text{th}}$  coordinate received from  $p_i$ )
  - (c) every process sends the collected vector and obtains these 3 new vectors
- A majority vector is computed using for each coordinate the one appearing a majority of times (or  $\perp$  if there is none)

- $(1, 2, \perp, 4)$

1 Got(1, 2, x, 4)	1 Got (1, 2, y, 4)	2 Got (1, 2, x, 4)	4 Got (1, 2, x, 4)
2 Got(1, 2, y, 4)	(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
3 Got(1, 2, 3, 4)	(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)
4 Got(1, 2, z, 4)			



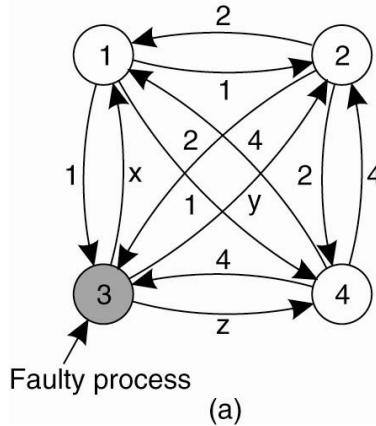
(b)

Year 2023

(c)

# Byzantine Tolerant Consensus

- Assumptions
  - Synchronous communication using unicast
  - Byzantine failures
  - A maximum of  $f < n/3$  processes can fail among  $n$



1 Got(1, 2, x, 4)  
2 Got(1, 2, y, 4)  
3 Got(1, 2, 3, 4)  
4 Got(1, 2, z, 4)

1 Got (1, 2, y, 4)	2 Got (1, 2, x, 4)	4 Got (1, 2, x, 4)
(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)

(b)

(c)

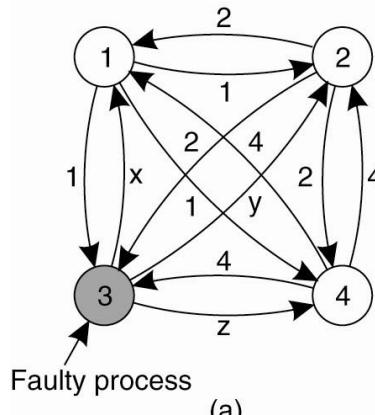
# Byzantine Tolerant Consensus

- Example of solution for  $n=4$  and  $f=1$

- (a) every process sends its input to all others
  - (b) the collected vectors ( $i^{\text{th}}$  coordinate received from  $p_i$ )
  - (c) every process sends the collected vector and obtains these 3 new vectors
- A majority vector is computed using for each coordinate the one appearing a majority of times (or  $\perp$  if there is none)

- $(1, 2, \perp, 4)$

1 Got(1, 2, x, 4)	1 Got $(1, 2, y, 4)$	2 Got $(1, 2, x, 4)$	4 Got $(1, 2, x, 4)$
2 Got(1, 2, y, 4)	(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
3 Got(1, 2, 3, 4)	(a, b, c, d)	(e, f, g, h)	(1, 2, y, 4)
4 Got(1, 2, z, 4)	(1, 2, z, 4)	(1, 2, z, 4)	(i, j, k, l)



(c)

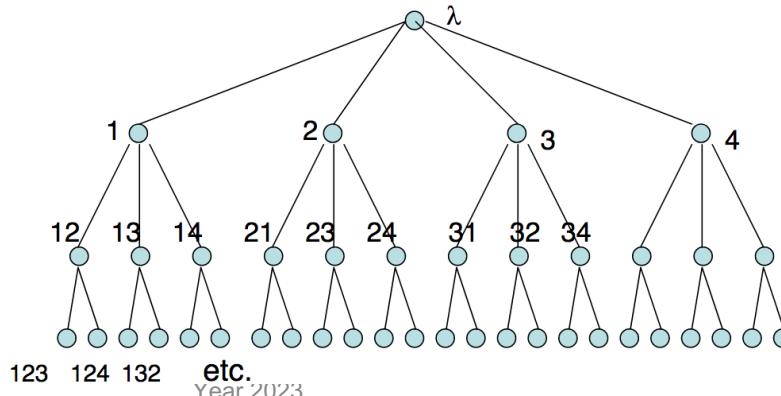
(b)

Year 2023

# Generalization: Exponential Information Gathering

## EIG Tree

- $T_{n,f}$  for  $n$  processes and  $f$  failures:
- $f+2$  levels
- Each path from root to leaf corresponds to a string of  $f+1$  distinct processes names
- Example  $T_{4,1}$

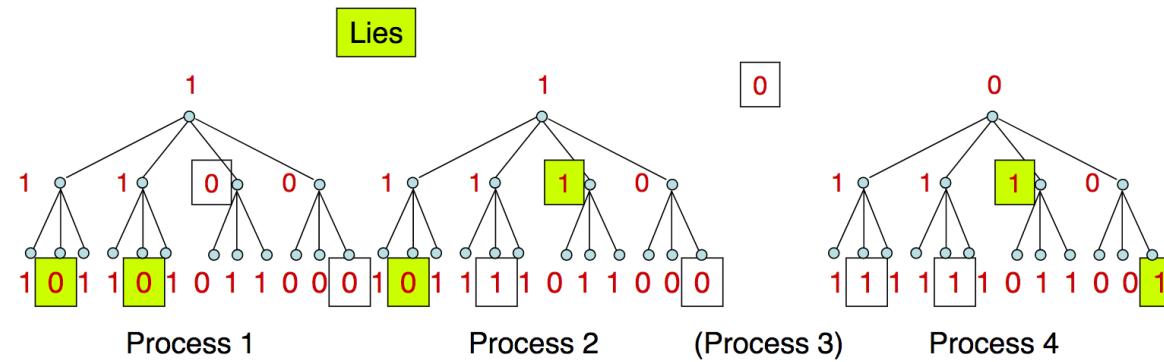
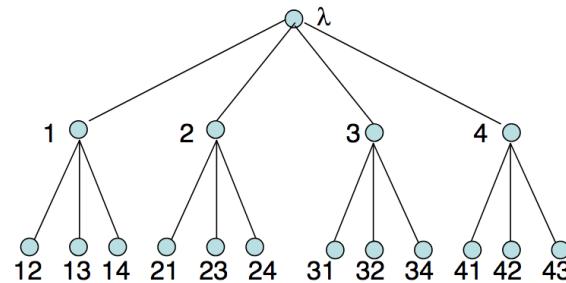


# Generalization: Exponential Information Gathering

- Each process  $i$  uses the same EIG tree,  $T_{n,f}$ .
- Decorates nodes of the tree with values  $V$ , level by level.
- Initially: Decorate root with  $i$ 's input value.
- Round  $r \geq 1$ :
  - Send all level  $r-1$  decorations for nodes whose labels don't include  $i$ , to everyone.
    - Including yourself---simulate locally.
  - Use received messages to decorate level  $r$  nodes---to determine label, append sender's id at the end.
  - If no message received, use  $\perp$ .
- The decoration for node  $(i_1, i_2, i_3, \dots, i_k)$  in  $i$ 's tree is the value  $v$  such that  $(i_k \text{ told } i) \text{ that } (i_{k-1} \text{ told } i_k) \text{ that } \dots \text{ that } (i_1 \text{ told } i_2) \text{ that } i_1 \text{ 's initial value was } v$ .
- Decision rule for stopping case:
  - Trivial
  - Let  $W = \text{set of all values decorating the local EIG tree}$ .
  - If  $|W| = 1$  decide that value, else default  $v_0$ .

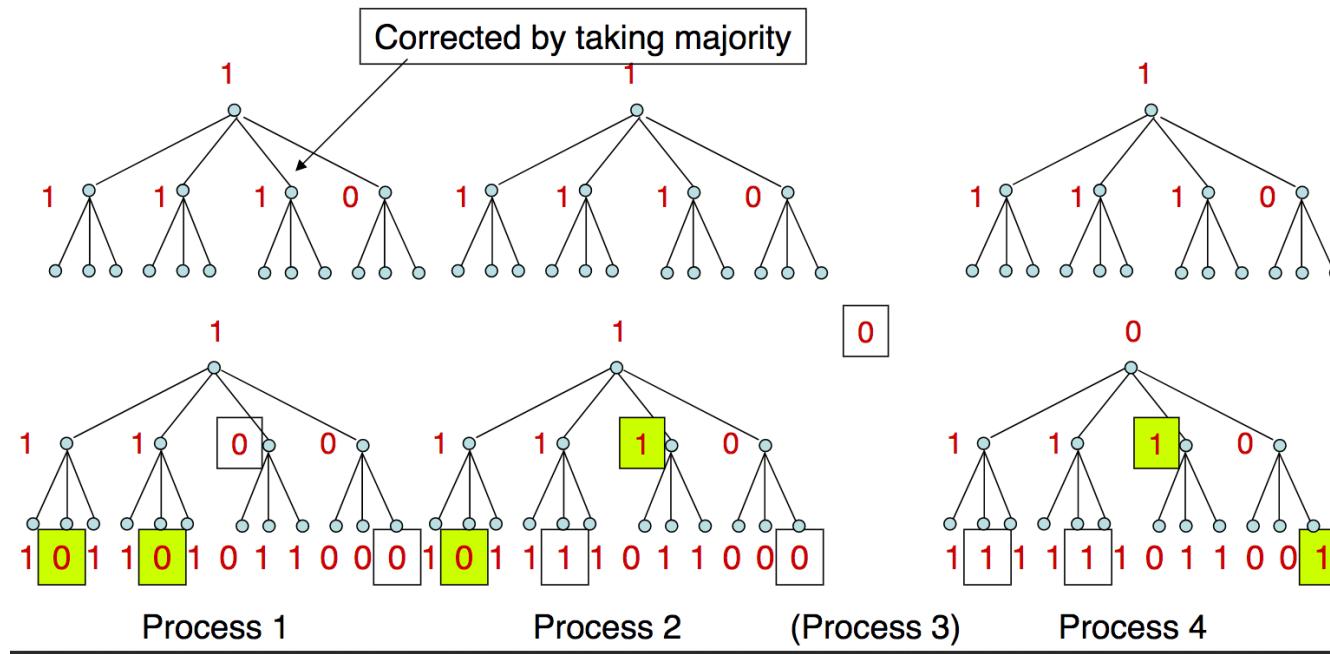
# Example EIG n=4, f=1

- $T_{4,1}$ :
- Consider a possible execution in which p3 is faulty.
- Initial values 1 1 0 0
- Round 1
- Round 2



## Example EIG n=4, f=1 (con't)

- Now calculate newvals, bottom-up, choosing majority values,  $v_0 = 0$  if no majority.



# Complexity of Byzantine Tolerant Consensus?

# Complexity of Byzantine Tolerant Consensus (EIG)?

**Message complexity.** There are  $f+1$  rounds where  $f < n$ .

The number of messages is at most  $O(n^2)$  in each round.

The total number of messages is thus  $O((f+1) n^2)$ .

**Communication complexity.** **Exponential** in the number of failures  $f$ ,  $O(b n^{f+1})$  bits.

**Time complexity.** There are  $f+1$  rounds, in each of these rounds, messages can be exchanged in parallel. This leads to a time complexity of  $O(f+1)$  message delays.

# Message losses

---

# Message loss

## Cause

- Networks are in general **unreliable**
- Messages can be **lost** (never been delivered even if sent)
- Examples:
  - A server receives too many requests simultaneously so it cannot treat all
  - A router drops the message because its queue is full

Message losses may impact the computation of a distributed system

---

# Message loss

## Coordinated Attack Problem



- Constraints of the problem
  - Two armies, each led by a general on separate mountains surrounding a battlefield (distributed system)
  - Can only communicate via messengers (message passing)
  - Messengers can be killed before reaching destination (message losses)
- Goal: they want to coordinate an attack
  - If they attack at different times, they both die
  - If they attack at the same time, they win

# Message loss

## Coordinated Attack Problem (con't)

- There is no protocols to make sure they will win!



12h!



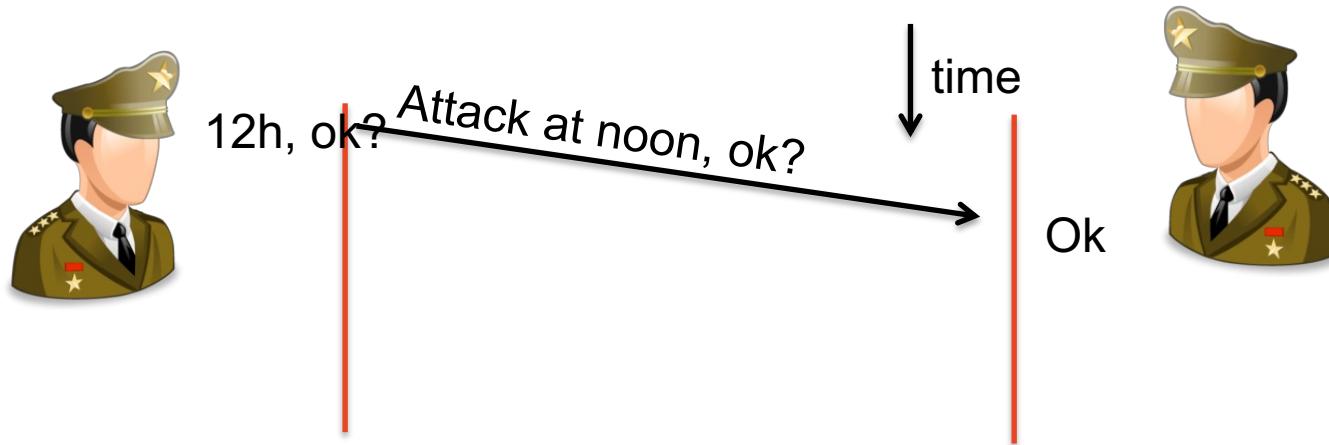
↓ time



# Message loss

## Coordinated Attack Problem (con't)

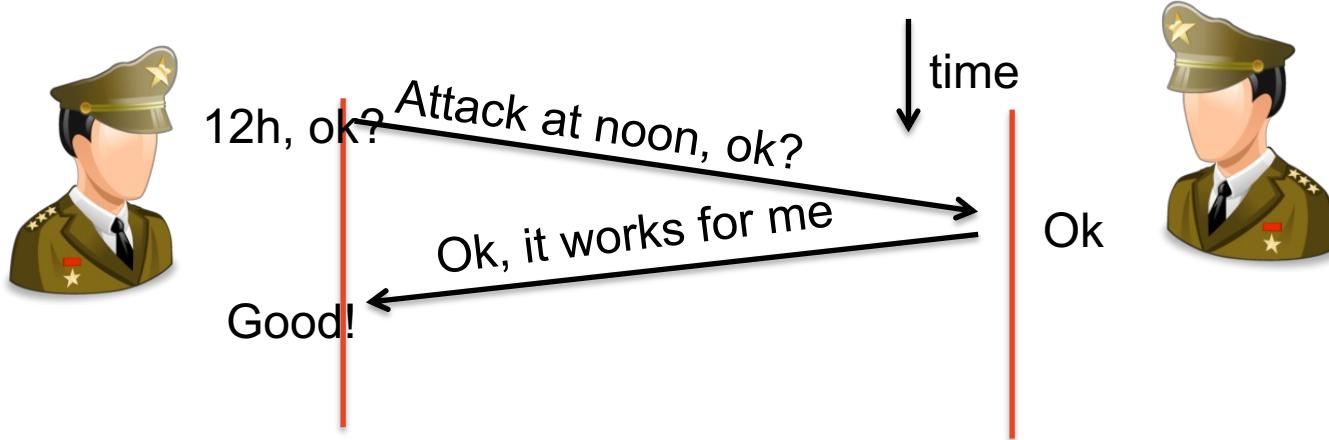
- There is no protocols to make sure they will win!



# Message loss

## Coordinated Attack Problem (con't)

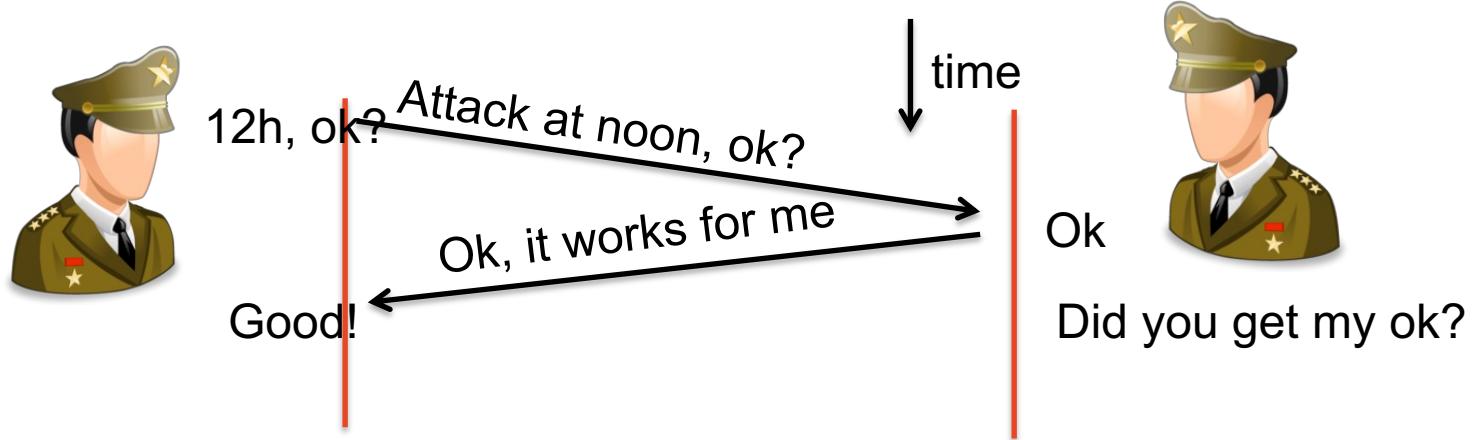
- There is no protocols to make sure they will win!



# Message loss

## Coordinated Attack Problem (con't)

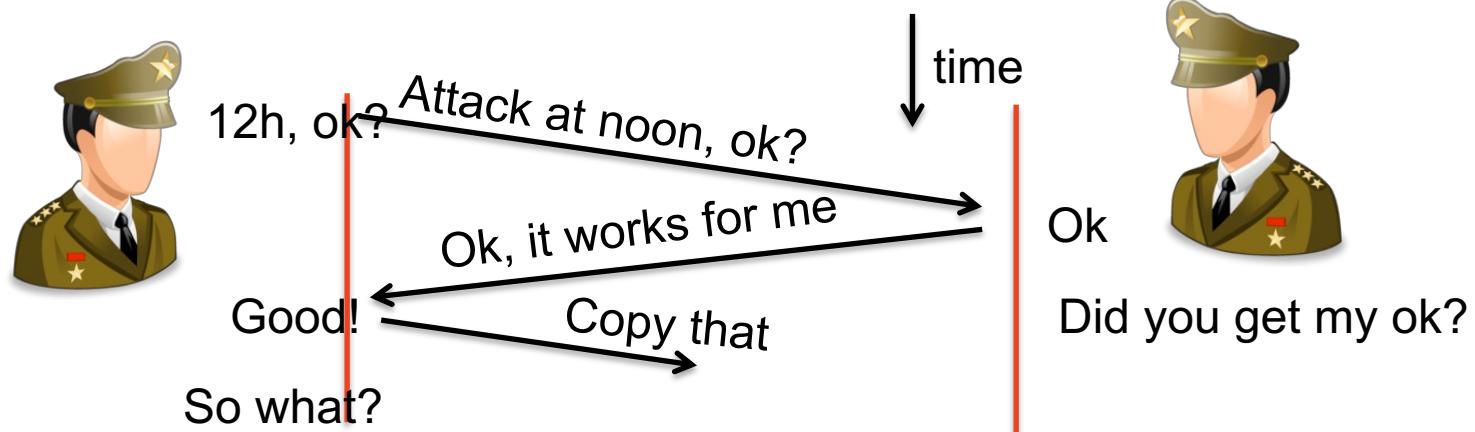
- There is no protocols to make sure they will win!



# Message loss

## Coordinated Attack Problem (con't)

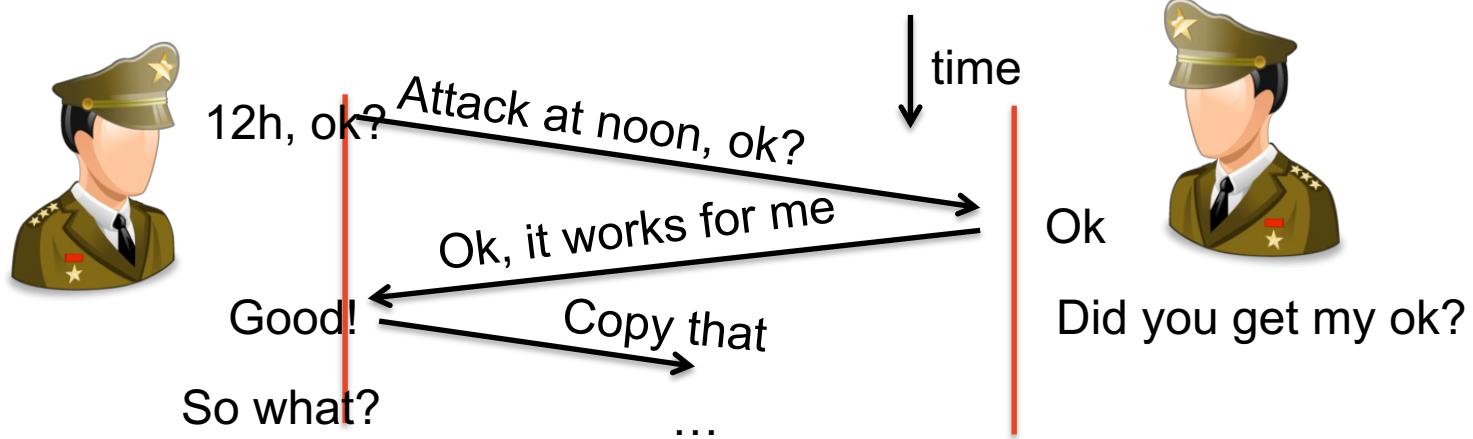
- There is no protocols to make sure they will win!



# Message loss

## Coordinated Attack Problem (con't)

- There is no protocols to make sure they will win!



# Theorem

There is no algorithm that ensures that the generals attack simultaneously

## Proof:

1. Consider the algorithm that sends the fewest messages
2. It still works if last message lost
3. So just don't send it (messengers' union happy!)
4. But now we have a shorter protocol!
5. Contradicting #1

**Fundamental limitation:** We need an unbounded number of messages, otherwise it is possible that no attack takes place!

---

# Message loss

## Analogy in networking

- Constraints of the problem
  - Two remote entities of a distributed system
  - Can only communicate through messages
  - The network is unreliable: messages can be dropped
- Goal: they want to make sure to do something simultaneously

This is impossible, even if all messages go through

# Asynchronous consensus

# Theorem

---

There is no algorithm that solves consensus in an asynchronous system with even one failure

## Proof:

There is no way to distinguish between a delayed message and a failure  
Termination requires this distinction

**Workaround:** some randomized consensus algorithms terminate in expected number of message delays.

[FLP85] Fischer M. J., Lynch N. A., and Paterson M. S., Impossibility of distributed consensus with one faulty process. *Journal of the ACM*, 32(2):374-382 (1985).

---

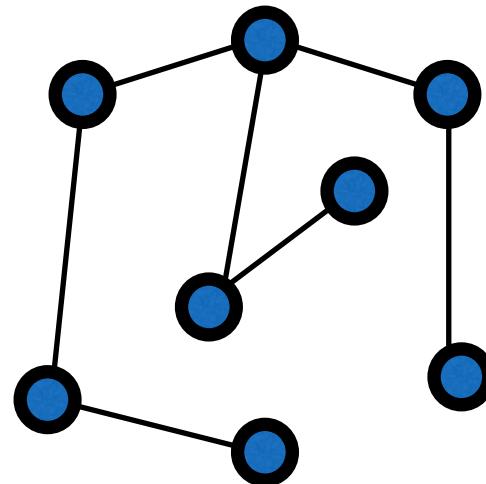
# Security

# Balance attacks

---

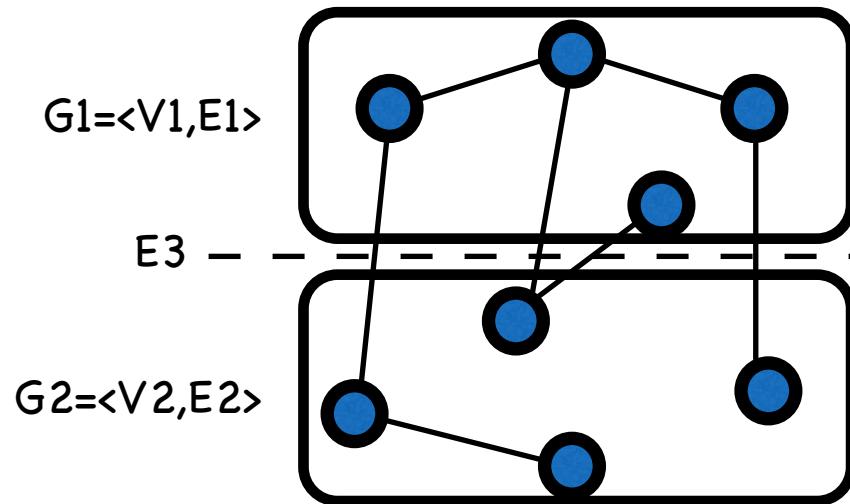
## The Balance Attack

Consider a communication graph  $G = \langle V, E \rangle$



## The Balance Attack (con't)

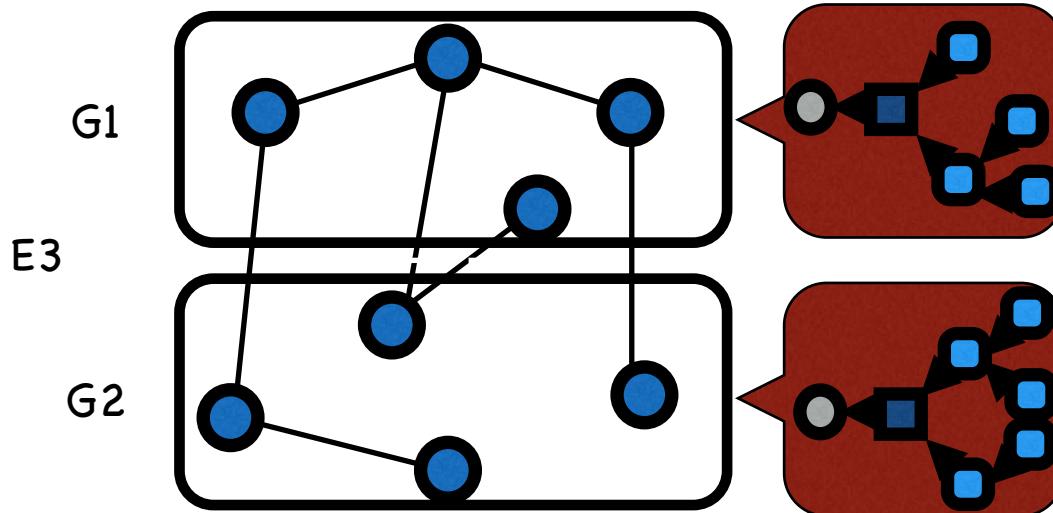
Consider a communication graph  $G = \langle V, E \rangle$



Let  $G_1$  and  $G_2$  be two subgraphs of  $G$  with the same mining power separated by  $E_3$  such that  $E = E_1 \cup E_2 \cup E_3$

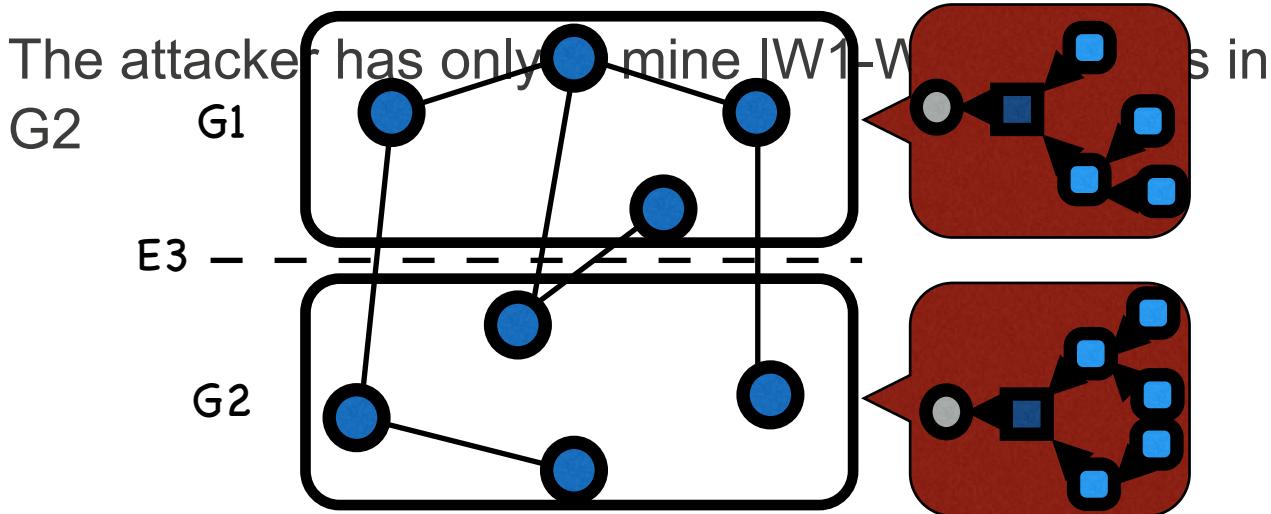
## The Balance Attack (con't)

Let an attacker issue  $t_A$  in  $G_1$  and delay links  $E_3$



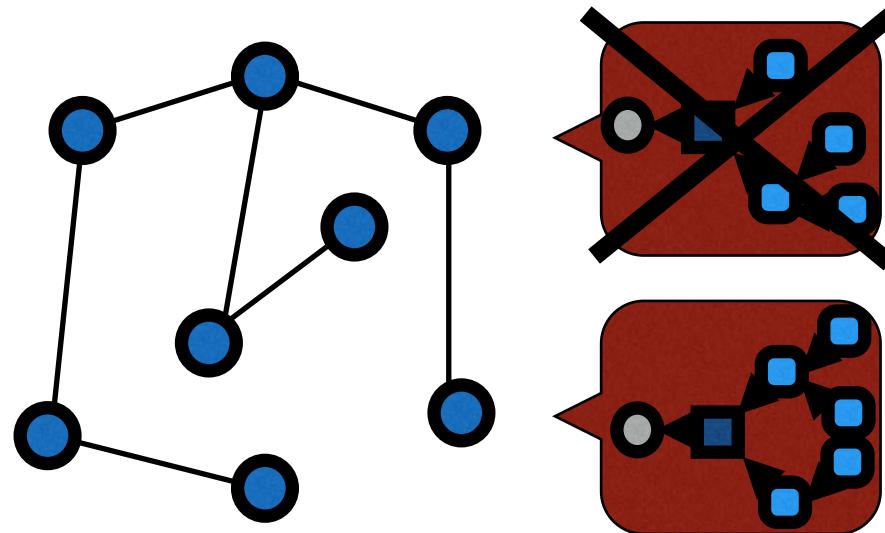
Then the two subgraphs end up with branches of similar lengths or weights  $W_1$  and  $W_2$

## The Balance Attack (con't)



## The Balance Attack (con't)

When the attacker stops delaying the messages



$t_A$  is discarded, allowing double spending with  $t_{A'}$

# R3 Ethereum Tesnet

# R3 Ethereum Testnet



# R3 Ethereum Testnet

○	r3	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	2.4 MH/s	0	0	#376,358	032235b4...7a879e73	6,341,489,326,541			
○	bnkasa	Geth/v1.4.5-stable/linux/go1.5.1	4 ms	1.3 MH/s	0	0	#375,704	80a4a9fe...b68d66b8	6,317,896,659,713			
○	cba	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.3 MH/s	1	0	#376,776	4b17aae1...e0b7da4f	6,358,160,274,447			
○	bony	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.3 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	bbva	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.3 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	cs	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	baml	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	ntrs	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.2 MH/s	0	0	#375,569	34316d78...300084964	6,313,742,215,265			
○	ubs	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	nordea	Geth/v1.4.5-stable/linux/go1.5.1	5 ms	1.2 MH/s	5	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	ms	Geth/v1.4.5-stable/linux/go1.5.1	7 ms	1.2 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	intesa	Geth/v1.4.5-stable/linux/go1.5.1	4 ms	1.2 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	cibc	Geth/v1.4.5-stable/linux/go1.5.1	8 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	r3lab - utils	Geth/v1.4.5-stable/linux/go1.5.1	1 ms	1.1 MH/s	8	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	r3lab	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	344 KH/s	3	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	rbc		6 ms		9	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	danske		2 ms		1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	bmo		6 ms		1	0	#376,776	4b17aae1...e0b7da4f	6,358,160,274,447			
○	stt		1 ms		0	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228			
○	macq		3 ms		0	0	#319,864	9076a734...bc5c0180	3,845,547,329,094			
○	nab		8 ms		5	0	#355,464	b7ec6fde...6db75642	5,453,765,160,440			
○	nomura											

20MH/s

# My miner

4-core Intel Core i5-6500, 3.2GHz

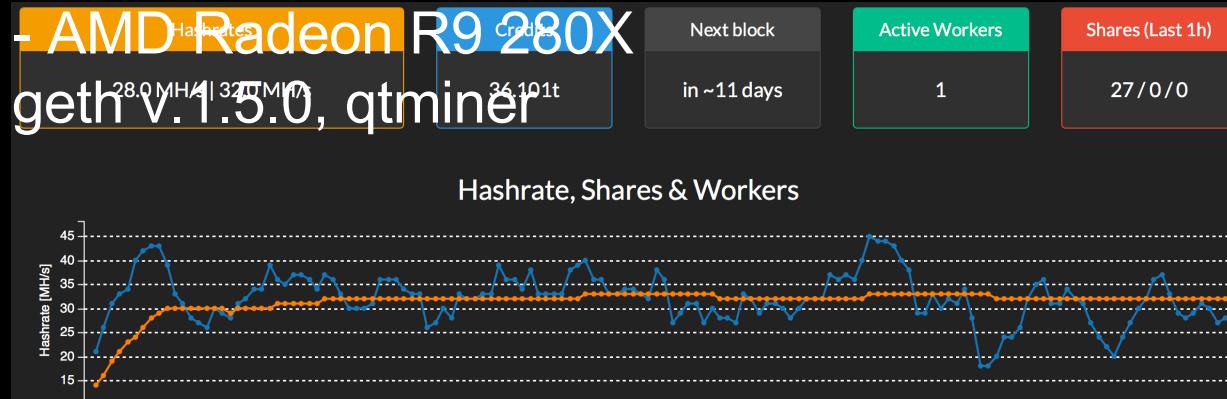
Ubuntu 14.04

2 GPUs:

- AMD Radeon R9 290X

- AMD Radeon R9 280X

geth v.1.5.0, qtminer



# Analysis of the Balance Attack

## Analysis

---

As there is no other strategy than **brute force** to solve the **crypto-puzzle**, each subgraph performs, during delay  $\tau$ , a series of  $n$  independent Bernoulli trials returning:

- 1 with probability  $p = (1-\rho)t/(2d)$  and
- 0 otherwise

The numbers  $X_1$  and  $X_2$  of blocks have a binomial distribution with mean:

$$\mu_c = np = \frac{(1 - \rho)t\tau}{2d}.$$

...where  $\rho$  is the probability of success for each trial.

## Analysis

---

We can upper-bound the deviation of the number of blocks mined  $X_i$  ( $i \in \{1, 2\}$ ) from their mean using Chernoff bounds [Motwani and P. Raghavan 1995] by

$$\Pr[|X_i - \mu_c| < \delta\mu_c] > 1 - 2e^{-\frac{\delta^2}{3}\mu_c}.$$

Let  $\Delta = |X_1 - X_2|$  be the difference of the number of blocks mined in each subgraph

## Analysis

Observe that the probability that these random variables are within a  $\pm\delta\mu_c$  is lower than the probability that their difference  $\Delta$  is upper-

b)  $\Pr[\Delta < 2\delta\mu_c] > \left(1 - 2e^{-\frac{\delta^2}{3}\mu_c}\right)^2.$

2nd Bernoulli inequality gives us:

$$\Pr[\Delta < 2\delta\mu_c] > 1 - 4e^{-\frac{\delta^2}{3}\mu_c}.$$

## Analysis

By choosing  $\delta = 2\rho/(1-\rho)$  we obtain that the expectation of the number of blocks mined by the malicious miner  $\mu_m$  is strictly greater than 1 with probability

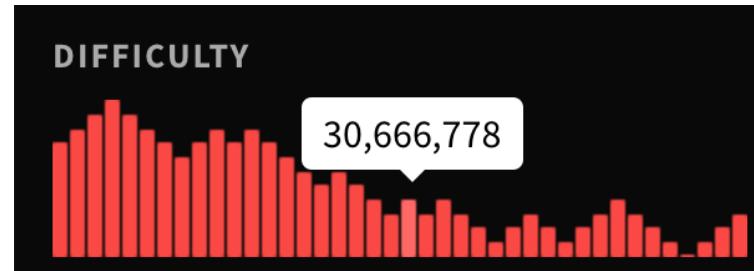
$$\Pr[\Delta < \mu_m] > 1 - 4e^{-\frac{4\rho^2 \mu_c}{3(1-\rho)^2}}.$$

With communication delay  $\tau \geq \frac{(1-\rho)6d \log(\frac{4}{\epsilon})}{4\rho^2 t}$

$$\Pr[\Delta < \mu_m] > 1 - \epsilon.$$

## Example with Banks

Consider an Ethereum/GHOST private chain with difficulty **30MH**...



## Example with Banks

Consider an Ethereum/GHOST private chain with difficulty 30MH and a mining power of **20MH/s**.

○ r3	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	2.4 MH/s	0	0	#376,558	032235d4...7a879e73	6,541,489,326,541	
○ bnkasa	Geth/v1.4.5-stable/linux/go1.5.1	4 ms	1.3 MH/s	0	0	#375,704	80a4a9fe...b68d66b8	6,317,896,659,713	
○ cba	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.3 MH/s	1	0	#376,776	4b17aae1...e0b7da4f	6,358,160,274,447	
○ bony	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.3 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ bbva	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.3 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ cs	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ bamt	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ ntrs	Geth/v1.4.5-stable/linux/go1.5.1	3 ms	1.2 MH/s	0	0	#375,569	34316d78...30084964	6,313,42,215,265	
○ ubs	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ nordea	Geth/v1.4.5-stable/linux/go1.5.1	5 ms	1.2 MH/s	5	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ ms	Geth/v1.4.5-stable/linux/go1.5.1	7 ms	1.2 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ intesa	Geth/v1.4.5-stable/linux/go1.5.1	4 ms	1.2 MH/s	2	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ cibc	Geth/v1.4.5-stable/linux/go1.5.1	8 ms	1.2 MH/s	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ r3lab + utils	Geth/v1.4.5-stable/linux/go1.5.1	1 ms	1.1 MH/s	8	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ r3lab	Geth/v1.4.5-stable/linux/go1.5.1	6 ms	344 KH/s	3	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ rbc		6 ms	✗	9	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ danske		2 ms	✗	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ bmo		2 ms	✗	1	0	#376,776	4b17aae1...e0b7da4f	6,358,160,274,447	
○ stt		6 ms	✗	1	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	
○ macq		1 ms	✗	0	0	#319,864	9076a734...bc5c0180	3,845,547,329,094	
○ nab		3 ms	✗	0	0	#355,464	b7ec6fde...6db75642	5,453,765,160,440	
○ nomura		8 ms	✗	5	0	#382,162	7f3dcc5c...532f9e3f	6,530,921,096,228	

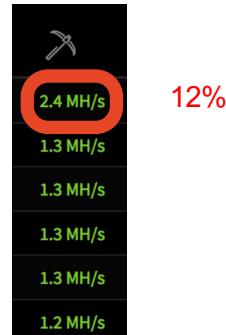
~20MH/s

Year 2023

## Example with Banks

Consider an Ethereum/GHOST private chain with difficulty 30MH and a mining power of 20MH/s.

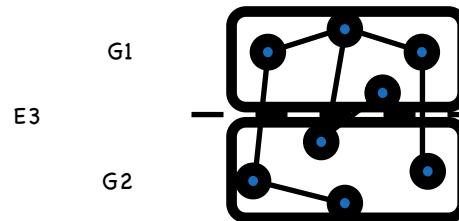
Assume the attacker has 12% of the mining power, i.e., 2.4MH/s, so he can mine, in expectation, 96 blocks during 20 minutes...



12%

## Example with Banks

Consider an Ethereum/GHOST private chain with difficulty 30MH and a mining power of 20MH/s. Assume the attacker has 12% of the mining power, i.e., 2.4MH/s, so he can mine, in expectation, 96 blocks during 20 minutes and that he can select E3 such that **each subgraph has a mining power of 8.8MH/s.**



## Example with Banks

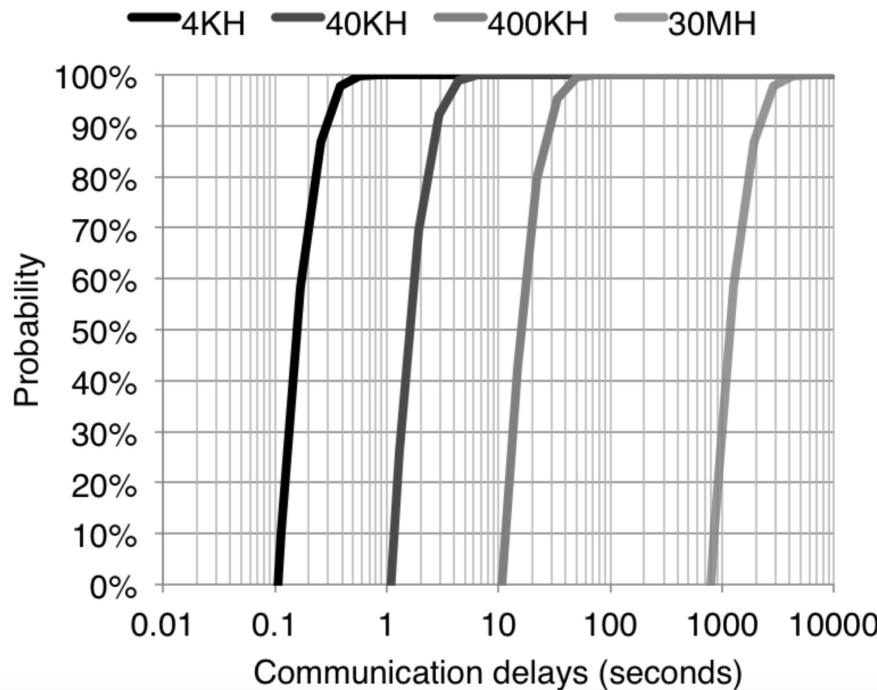
Consider an Ethereum/GHOST private chain with difficulty 30MH and a mining power of 20MH/s.

Assume the attacker has 12% of the mining power, i.e., 2.4MH/s, so he can mine, in expectation, 96 blocks during 20 minutes and that he can select E3 such that each subgraph has a mining power of 8.8MH/s.

Let us select  $\delta=95/(2\mu_c) \approx 0.13$  for the difference  $\Delta$  to be lower than 96, which happens with probability

$$P[\Delta < 96] > 52\%$$

## Probability of Attack Success per Difficulties



Number of nodes	50
Number of miners	15
Total mining power (MH/s)	20
Mining power of the most powerful miner (MH/s)	2.4
Difficulty (MH)	30

[NG17] Natoli C., Gramoli V. The Balance or Why Proof-of-Work is Ill-Suited for Consortium Blockchains. DSN 2017.  
The University of Sydney  
May 2020

# Synchrony

*"This is not at all novel. This basically just points out that proof of work has a synchrony assumption."*

Vitalik Buterin



# Synchrony

*"This is not at all novel. This basically just points out that proof of work has a synchrony assumption."*

Vitalik Buterin

The screenshot shows a news article from news.com.au. The header features the news.com.au logo with a colorful circular icon, a pink banner reading "The news in colour", and a search bar. Below the banner, there are navigation links for National, World, Lifestyle, Travel, Entertainment, Technology (highlighted in red), Finance, and Sport. A large, bold, black headline reads: "Aussie internet users warned to expect slow speeds for at least six weeks after undersea cables damaged". A sub-headline below it states: "AUSTRALIAN web users will be plagued by slow internet speeds for at least six weeks after a number of undersea cables were damaged." At the bottom right, there is a sidebar with an advertisement for "Mylease" and a timestamp "SEPTEMBER 1, 2017 6:18".

The news.com.au website features a prominent pink banner at the top with the slogan "The news in colour". Below the banner, the navigation menu includes categories like National, World, Lifestyle, Travel, Entertainment, Technology, Finance, and Sport. A large, bold headline dominates the page: "Aussie internet users warned to expect slow speeds for at least six weeks after undersea cables damaged". A sub-headline provides more detail: "AUSTRALIAN web users will be plagued by slow internet speeds for at least six weeks after a number of undersea cables were damaged". In the bottom right corner, there is a small sidebar with an advertisement for "Mylease" and a timestamp indicating the article was published on September 1, 2017, at 6:18.

# Feasibility of the Balance Attack

## Network attacks

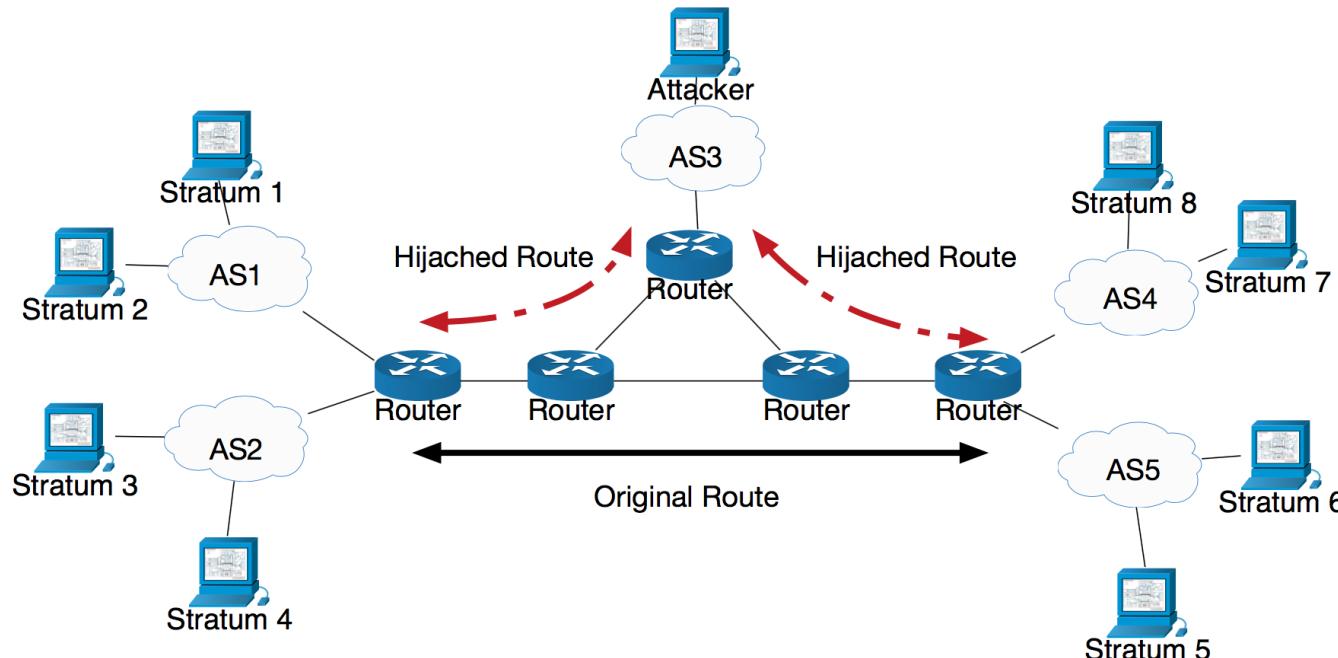
The Balance Attack requires network delays:

- Man-in-the-middle attack
- DoS attack
- BGP hijacking attack
- ARP spoofing attack, ...

Ethereum can run as:

- a public chain
- a consortium blockchain
- a private blockchain

# BGP Hijacking Attack



# Ethereum mining pools

Power (%)	Pool Name	Stratum Servers	Location	ASN	AS Owner
27.02	f2pool	eth.f2pool.com	Hangzhou, China	37963	Alibaba (China) Technology Co., Ltd.
23.76	Ethermine	us1.ethermine.org	Montreal, Canada	16276	OVH SAS
		us2.ethermine.org	California, US	63949	Linode, LLC
		eu1.ethermine.org	France	16276	OVH SAS
		eu2.ethermine.org	France	16276	OVH SAS
		asia1.ethermine.org	Singapore	16276	OVH SAS
9.73	miningpoolhub	us-east.ethash-hub.miningpoolhub.com	Georgia, US	63949	Linode, LLC
		europe.ethash-hub.miningpoolhub.com	Hesse, Germany	63949	Linode, LLC
		asia.ethash-hub.miningpoolhub.com	Tokyo, Japan	63949	Linode, LLC
9.7	Nanopool	eth-eu1.nanopool.org	France	16276	OVH SAS
		eth-eu2.nanopool.org	France or Italy	16276	OVH SAS
		eth-asia1.nanopool.org	Singapore	16276	OVH SAS
		eth-us-east1.nanopool.org	Montreal, Canada	16276	OVH SAS
		eth-us-west1.nanopool.org	California, US	20473	Choopa, LLC
9.12	ethfans	guangdong-pool.ethfans.org	Fujian, China	4134	No.31,Jin-rong Street
		huabei-pool.ethfans.org	Fujian, China	4134	No.31,Jin-rong Street
6.24	DwarfPool	eth-eu.dwarfpool.com	France	16276	OVH SAS
		eth-us.dwarfpool.com	Montreal, Canada	16276	OVH SAS
		eth-us2.dwarfpool.com	Las Vegas, US	53667	FranTech Solutions
		eth-ru.dwarfpool.com	France	16276	OVH SAS
		eth-asia.dwarfpool.com	Taiwan	59253	Leaseweb Asia Pacific pte. ltd.
		eth-cn.dwarfpool.com	Shanghai, China	37963	Alibaba (China) Technology Co., Ltd.
		eth-cn2.dwarfpool.com	Beijing, China	37963	Alibaba (China) Technology Co., Ltd.
		eth-sg.dwarfpool.com	Singapore	59253	Leaseweb Asia Pacific pte. ltd.
		eth-au.dwarfpool.com	Melbourne, Australia	38880	Micron21 Melbourne Australia Datacentre
		eth-ru2.dwarfpool.com	Moscow, Russia	42632	MnogoByte LLC
		eth-hk.dwarfpool.com	Hong Kong	45102	Alibaba (China) Technology Co., Ltd.
		eth-br.dwarfpool.com	Sao Paulo, Brazil	262287	Maxhost Hospedagem de Sites Ltda
4.45	BW	eth-ar.dwarfpool.com	Rosario, Argentina	27823	Dattatec.com
		ether.bw.com	Wuhan, China	58563	CHINANET Hubei province network
3.34	Ethpool	us1.ethpool.org	Montreal, Canada	16276	OVH SAS
		us2.ethpool.org	Montreal, Canada	16276	OVH SAS
		eu1.ethpool.org	France	16276	OVH SAS
		asia1.ethpool.org	Singapore	16276	OVH SAS
1.83	Coinotron	coinotron.com	Poland	51290	HOSTTEAM-AS
0.88	Poolgpu	eth.poolgpu.com	Hangzhou, China	37963	Alibaba (China) Technology Co., Ltd.
96.07	Total				

## Ethereum pools power is sufficient to double spend

Adversary	Mining power of members in an adversary subgroup (%) The rest of subgroup	Mining power of members in a victim subgroup (%)	Difference between two subgroups	Attack Success rate (%)
27.02	23.76, 6.24, 3.34, 0.88	9.73, 9.7, 9.12, 4.45, 1.83	26.41	76.7
23.76	27.02, 6.24, 1.83, 0.88	9.73, 9.7, 9.12, 4.45, 3.34	23.39	63.3
9.73	27.02, 9.12, 4.45, 1.83, 0.88	23.76, 9.7, 6.24, 3.34	9.99	56.7
9.7	27.02, 9.12, 4.45, 1.83, 0.88	23.76, 9.73, 6.24, 3.34	9.93	43.3
9.12	27.02, 9.7, 4.45, 1.83, 0.88	23.76, 9.73, 6.24, 3.34	9.93	43.3
6.24	27.02, 9.7, 4.45, 3.34	23.76, 9.73, 9.12, 1.83, 0.88	5.43	40
4.45	27.02, 9.7, 6.24, 1.83, 0.88	23.76, 9.73, 9.12, 3.34	4.17	40
3.34	27.02, 9.7, 6.24, 1.83, 0.88	23.76, 9.73, 9.12, 4.45	1.95	43.3
1.83	27.02, 9.7, 6.24, 3.34, 0.88	23.76, 9.73, 9.12, 4.45	1.95	33.3
0.88	27.02, 9.7, 6.24, 3.34, 1.83	23.76, 9.73, 9.12, 4.45	1.95	26.7

Attack success with the mining power of the 10 largest mining pools of Ethereum at the end of July 2017

## On the difficulty of partitioning mining pools

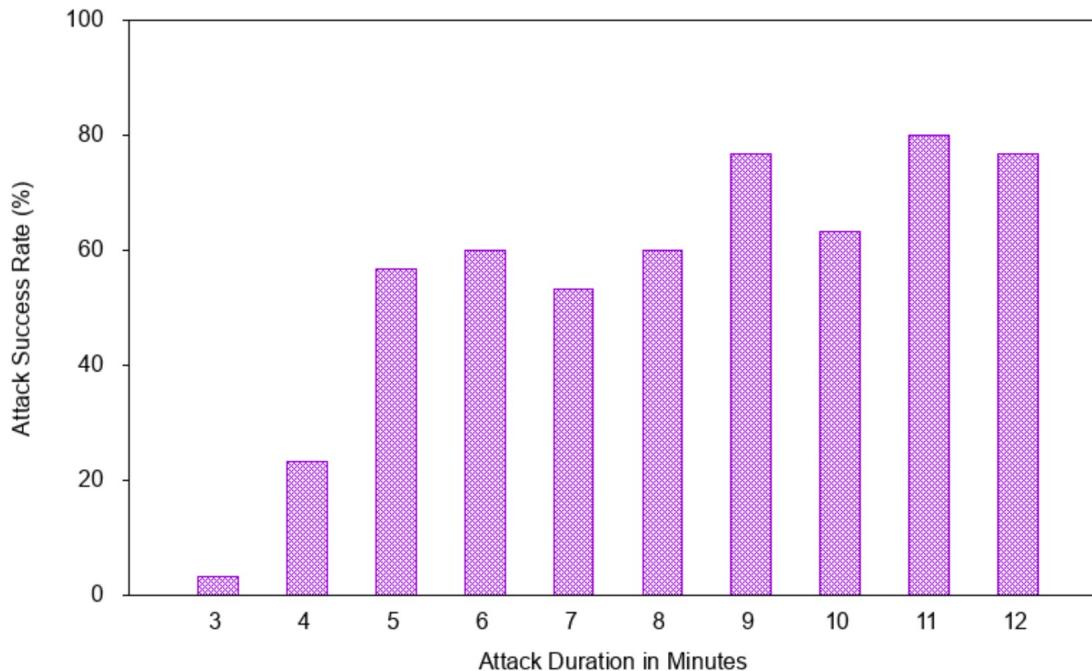
Mining pools often operate multiple stratum servers

These servers are in same regions and different ASes

Disconnected miners simply try to reconnect to the next stratum server

This makes it hard to partition the network as the miner need to target multiple ASes

## Feasibility of the attack on consortium Ethereum



[EGJ18] Ekparinya P., Gramoli V., Jourjon G. Impact of Man-in-the-Middle Attacks on Ethereum. *IEEE SRDS 2018*.  
The University of Sydney  
Sep. 2023

# Quiz

---

# Question 1

How to make sure that a transaction is committed in Ethereum or Bitcoin?

---

# Scalability

# Leader-based Consensus

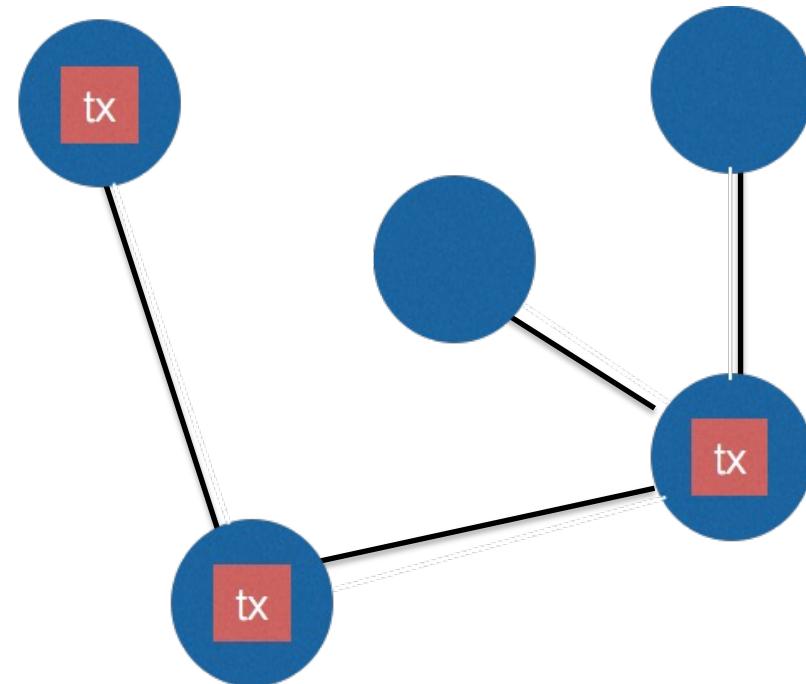
# Leader-based consensus in blockchains

The process that plays the role of the leader changes over time.

It's hard for a client (e.g. wallet) to know which process is the leader to send a request to.

Easier to:

- send transaction requests to the processes that will run the consensus
- send their request to the processes that are the closest geographically



# Leader vs. leaderless consensus algorithms

Consider a consensus algorithm that decides a *set of transactions* of  $B$  bits

We want to compare the performance of:

- A leader-based solution
- A leaderless solution

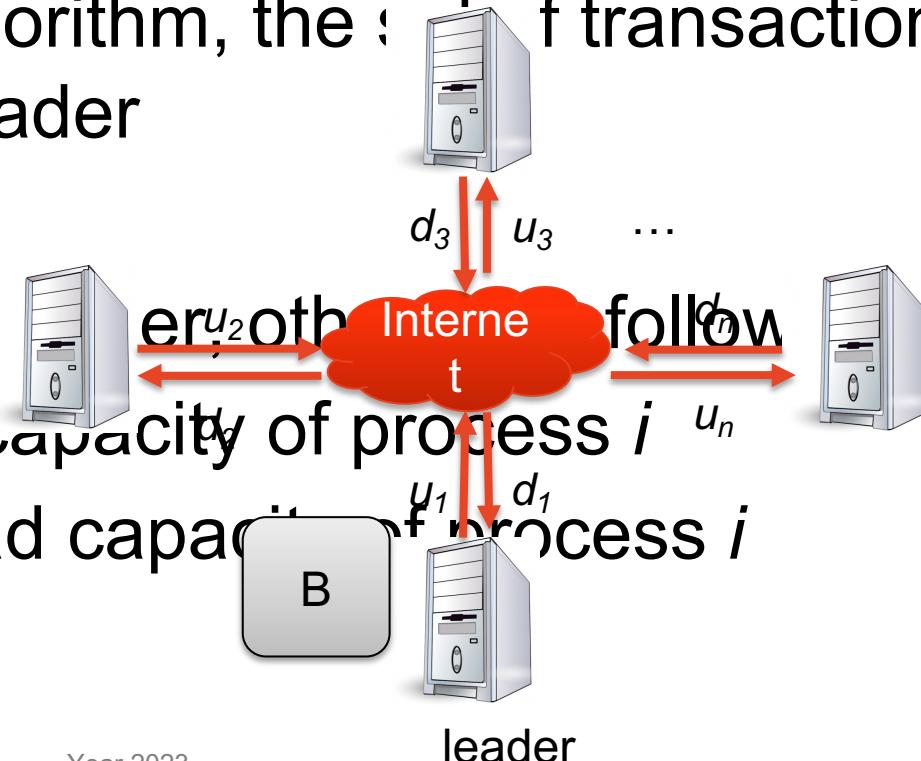
In a leader-based algorithm, the set of transactions is proposed by the leader

In a leaderless algorithm, the set of transactions is proposed by all the participating processes

# Block exchange time in leader-based algorithms

In a leader-based algorithm, the set of transactions is proposed by the leader

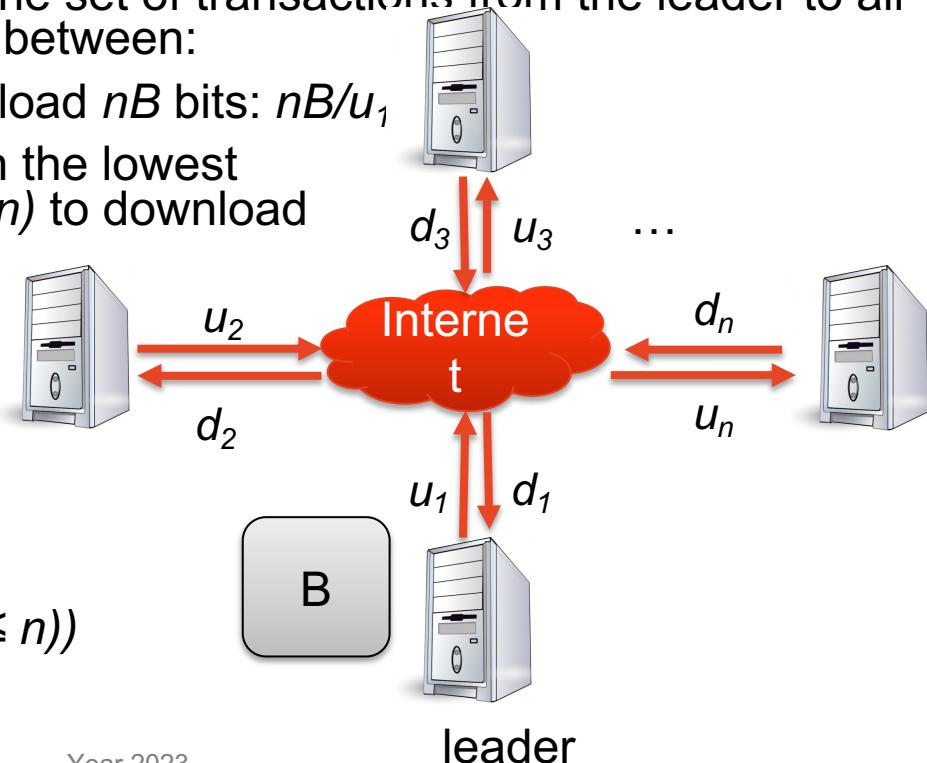
Let process 1 be the leader  
Let  $u_i$  be the upload capacity of process  $i$   
Let  $d_i$  be the download capacity of process  $i$



# Block exchange time in leader-based algorithms (con't)

The time  $P$  it takes to exchange the set of transactions from the leader to all other processes is the maximum between:

- The time for the leader to upload  $nB$  bits:  $nB/u_1$ ,
- The time for the follower with the lowest download rate  $\min(d_i: 1 \leq i \leq n)$  to download  $B$  bits:  $B/\min(d_i: 1 \leq i \leq n)$



$$\Rightarrow P = \max(nB/u_1, B/\min(d_i: 1 \leq i \leq n)) \\ = O(nB)$$

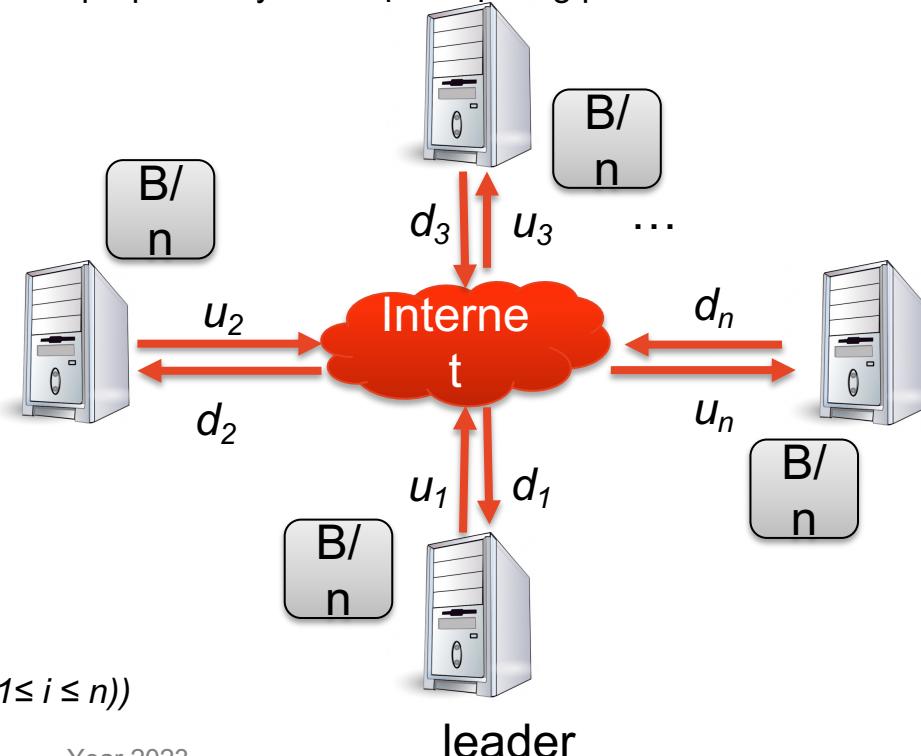
# Block exchange time in leaderless algorithms

In a leaderless algorithm, the set of transactions is proposed by all the participating processes

Consider that each process has a portion  $B/n$  of the total set of transactions

The time it takes to exchange these sets of transactions is the max of:

- The time for the slowest process to upload  $B/n$  bits
- The time for the slowest process to download  $B$  bits



$$\Rightarrow P = \max((B/n)/\min(u_i : 1 \leq i \leq n), (B/\min(d_i : 1 \leq i \leq n))$$

The University of Sydney  
=  $O(B)$

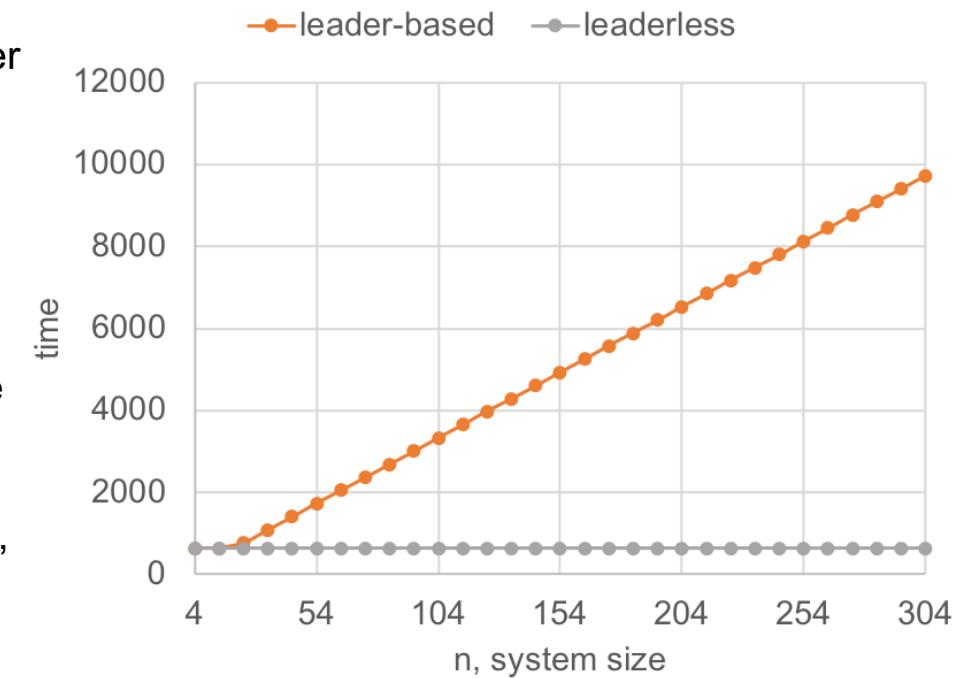
# Leader vs. leaderless block exchange time

Consider that the leader has 10 times higher upload rate than other processes

And that every process download 2 times faster than it uploads

The time it takes for all processes to be aware of the set of transactions is the same on a small network

When the number of processes reaches 21, then the leaderless algorithm starts being faster and faster



# The Set Consensus Problem

# Traditional Consensus does not Scale

Two reasons why traditional consensus does not scale:

1. The leader tries to impose its block, disallowing other nodes from proposing blocks
2. The time to propagate the block to the followers is limited by:
  - the upload capacity of the leader and
  - the lowest download capacity of the followers

# Defining Blockchain Consensus

Instead of proposing blocks, processes propose sets of transactions that will be included in the block:

1. All processes propose a set of transactions
2. They exchange these sets among each other
3. They decide a block that:
  - Is not necessarily what any process has proposed
  - Can result from the union of several proposed sets of transactions

# Blockchain Consensus

A set of transactions is *non-conflicting* if it does not contain conflicting transactions.

Assuming that each correct node proposes a proposal  $s$ , the Set Byzantine Consensus (SBC) problem is for each of them to decide on a set in such a way that the following properties are satisfied:

- SBC-Termination: every correct node eventually decides a set of transactions;
- SBC-Agreement: no two correct nodes decide on different sets of transactions;
- SBC-Validity: a decided set of transactions is a non-conflicting set of valid transactions taken from the union of the proposed sets; and if all nodes are correct and propose a common valid non-conflicting set of transactions, then this subset is the decided set.

# Democratic Byzantine Fault Tolerance

# Democratic Byzantine Fault Tolerance (DBFT)

- est: local current estimate of the value to decide
- r: local round number, initially 0
- bin-values[]: array of binary values for all rounds
- b, auxiliary binary value
- $\text{values}_i$ , auxiliary set of values
- The algorithm uses two message types
  - EST[r]() used at round r to BV-bcast est.
  - AUX[r]() used at round r to broadcast bin-values[r]

# DBFT (con't)

## BV-broadcast [JACM'15]

- BV-obligation: if  $t+1$  correct BV-bcast  $v$ , then  $v$  is eventually added to the set bin-values of all correct processes
- BV-justification: if  $p_i$  is correct and  $v$  in  $\text{bin-values}_i$ , then  $v$  was BV-bcast by a correct process
- BV-uniformity: if  $v$  is added to  $\text{bin-values}_i$  of correct  $p_i$ , then eventually  $v$  will be in  $\text{bin-values}_j$  for all correct  $p_j$
- BV-termination: eventually  $\text{bin-values}$  of correct  $p_i$  is not empty

# Binary Consensus

```
1. est ← v
2. r ← 0
3. while (true) {
4.   r ← r+1
5.   BV-bcast EST[r](est)
6.   wait until (bin-values[r] ≠ ∅)
7.   bcast AUX[r](bin-values[r])
8.   wait until (messages AUX[r](bvalp1), ..., AUX[r](bvalp(n-t)) received
9.     from (n-t) distinct processes p(x), 1 ≤ x≤n-t, and their content is such that:
10.    ∃ a non-empty set valuesi such that
11.      (i) valuesi ⊆ bin-values[r]i and
12.          (ii) valuesi = U b_valx
13.    b ← r mod 2
14.    if (valuesi = {v}) {
15.      est ← v
16.      if (v=b) then decide(v) and exit (if not done yet)
17.    } else {
18.      est ← b
19.    }
20. }
```

# Binary Consensus (con't)

```
1. est ← v
2. r ← 0
3. while (true) {
4.   r ← r+1
5.   BV-bcast EST[r](est)
6.   wait until (bin-values[r] ≠ ∅)
7.   bcast AUX[r](bin-values[r])
8.   wait until (messages AUX[r](bvalp1), ..., AUX[r](bvalpn-t) from (n-t) distinct processes p(x), 1 ≤ x ≤ n-t, and bvalpi is such that:
9.     ∃ a non-empty set valuesi such that
10.       (i) valuesi ⊆ bin-values[r]i and
11.           (ii) valuesi = U b_valx
12.       b ← r mod 2
13.       if (valuesi = {v}) {
14.         est ← v
15.         if (v=b) then decide(v) and exit (if not done yet)
16.       } else {
17.         est ← b
18.       }
19.     }
20. }
```

Filter out  
values  
proposed  
only by  
Byzantine  
processes

# Agreement proof

Lemma 1: If at the beginning of a round, all correct processes have the same estimate, they never change their estimate thereafter.

Proof sketch. Assume all correct have the same estimate  $v$  at round  $r$ . BV-Obligation and BV-Justification =>  $\text{bin\_values}[r] = \{v\}$ . Hence, at every correct we have values =  $\{\bar{v}\}$ , so that est becomes  $v$ .

# Agreement proof (con't)

Lemma 2: Let  $p_i$  and  $p_j$  be correct. If their values are singletons, then they are the same.

Proof sketch. If a correct process has values  $= \{v\}$  then it received  $\text{AUX}[r]\{v\}$  from  $n-t$  distinct processes and  $t+1$  correct. For two correct processes to have  $w$  and  $v$  as this singleton, it would mean they received these distinct values from  $n-t$  processes each. As  $(n-t)+(t+1) > n$ , one correct process must have sent the same value to both, and we have  $v=w$ .

# Agreement proof (con't)

Theorem [Agreement]: No two correct processes decide different values.

Proof sketch. Let  $r$  be the 1st round where a correct process decides, hence  $\text{values}[r] = \{v=r \bmod 2\}$ . If another correct process decides  $w$  in the same round, then  $v=w$  by Lemma 2. Let  $p_j$  be a correct that does not decide in round  $r$ . By Lemma 2,  $\text{values}_j \neq \{w\}$  so  $\text{values}_j = \{0,1\}$ . Thus  $\text{est}_j = v$ . All correct have thus the same estimate in round  $r+1$ , and stick to it (Lemma 1).

# Safe version of DBFT

BC, binary consensus instance

```
mv-propose(v) { // similar to [PODC '94]
    RB-bcast VAL(v) // reliable broadcast [Bra'87]
    repeat if ( $\exists k : \text{proposals}[k] \neq \perp \wedge \text{BC}[k].\text{binpropose}()$  not invoked)
        BC[k].binpropose(1)
    until ( $\exists l : \text{bin-decisions}[l] = 1$ )
    for each k such that BC[k].binpropose() not invoked:
        BC[k].binpropose(0)
    wait until ( $\wedge 1 \leq x \leq n \text{ bin-decisions}[x] \neq \perp$ )
    j = min{x : bin-decisions[j] = 1}
    wait until proposals[j]  $\neq \perp$ 
    decide( $\cup_{j: \text{bin-decisions}[j] = 1} \text{proposals}[j]$ )
}
```

```
when val(v) is RB-delivered from pj do // reliable broadcast delivery
    if valid(v) then
        proposals_i[j]  $\leftarrow v$ ;
        BV-deliver b-val[1](1) to BC[j]
```

```
when BC[k].binpropose() returns b do bin-decisions[k]  $\leftarrow b$ 
```

Spawn multiple binary cons. instances concurrently  
Use binary decisions as a bitmask

# Redbelly Blockchain

---

# The Transaction Verification Cost

We saw how to leverage bandwidth for the consensus solution to scale

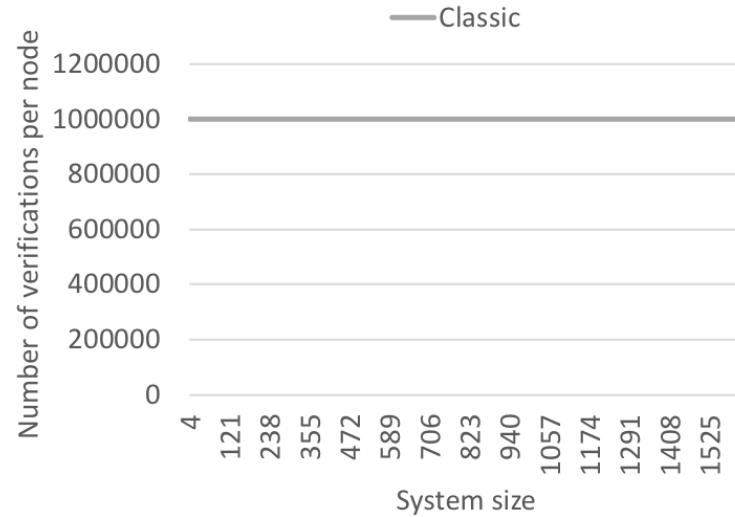
The Set Byzantine Consensus problem is well-suited for blockchains but requires a verification of signatures to check that a transaction is valid

This verification typically uses a public-key cryptosystem, that is CPU intensive and costly.

We will see how to lower the cost of verification to build a scalable blockchain

# CPU waste

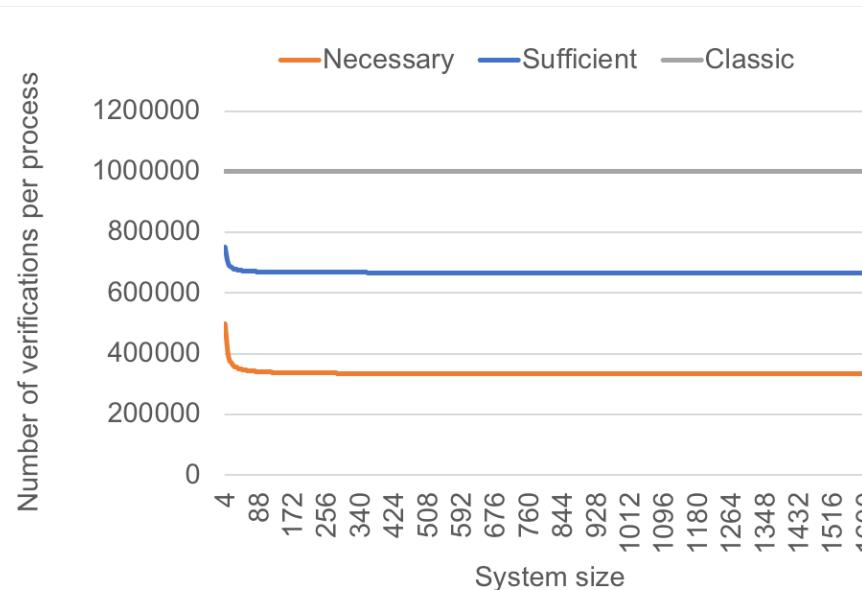
In traditional blockchains all miners verify all transactions which limits scalability



This means that as the number of miners grow, the number of verification grows proportionally

# How Sharding can help Scaling

For example, if only  $k$  nodes out of the  $n$  nodes verify the same set of transactions, then another set of  $k$  nodes can verify other transactions in parallel, hence saving time



The more the system grows, the more CPU resources we bring to the system and the faster the system verifies all transactions. This is called *verification sharding*.

# Red Belly Blockchain

It is a community blockchain [VG18]

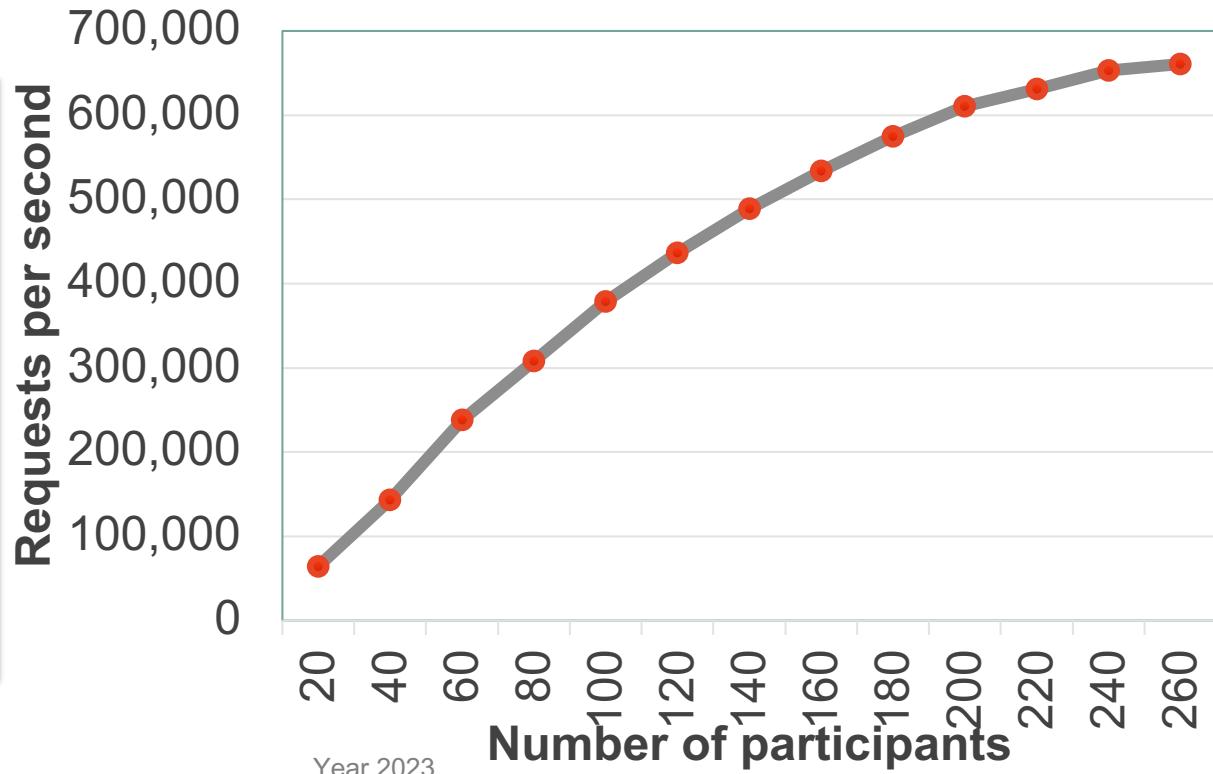
- It differs from a consortium blockchain by not pre-determining the participants but by having a set of participants that changes at runtime.
- It differs from a public blockchain by not incentivizing all participants to mine the same block

As more nodes join the set of participants, the performance increases.

# Red Belly Blockchain

Performance scales with the amount of participating resources [Gra17]

Amazon EC2 instances  
One availability zone in US  
System size from 4 to 260  
18 HT cores per node  
60 GB memory  
2 Gbps  
Verification of ECDSA sig.  
 $t = 6$  failures  
TCP+SSL



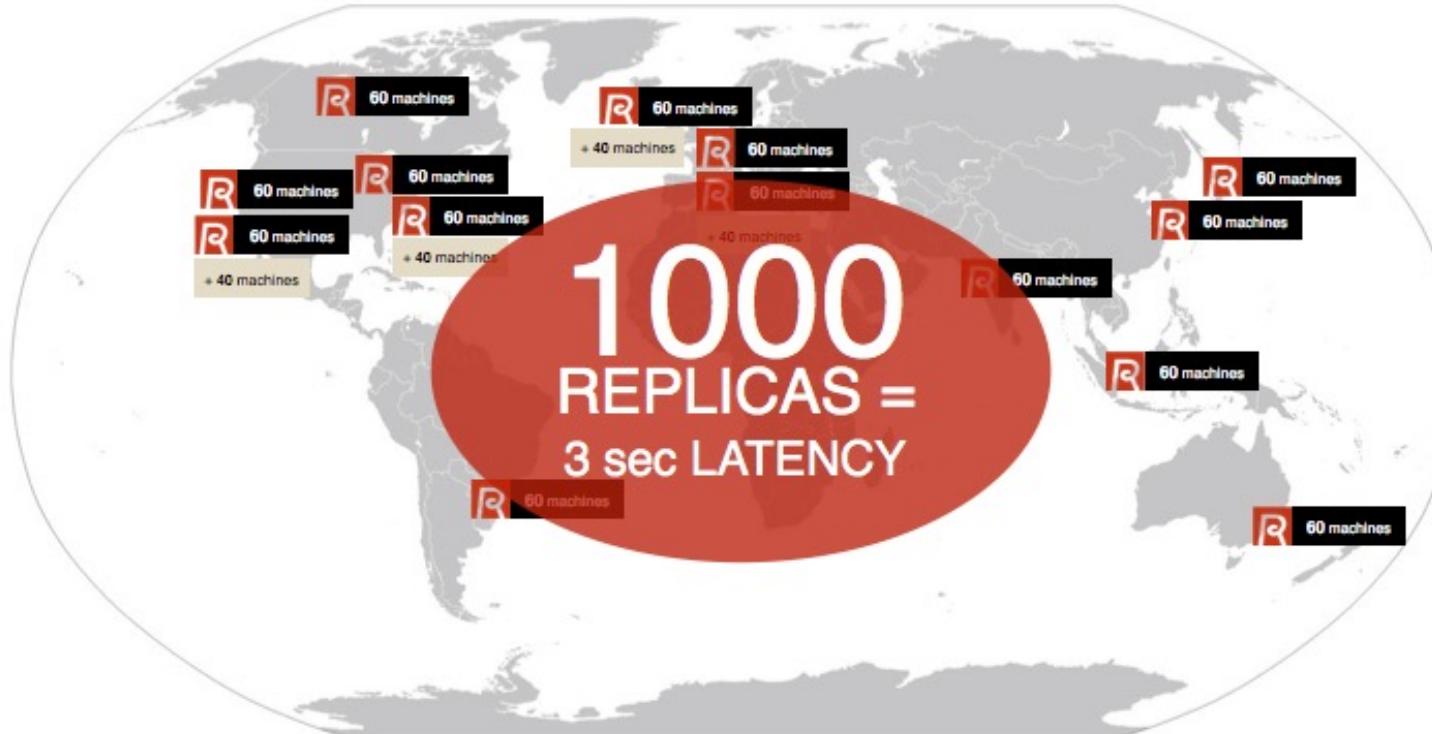
# Red Belly Blockchain

Scalable to 1000 machines from 11 countries in 4 continents



# Red Belly Blockchain

Fast: transactions executed within 3 seconds



# Conclusion

The Red Belly Blockchain is secure and scalable.

It solves the Set Byzantine Consensus problem using the Democratic BFT algorithm.

Because it does not use any leader, it leverages the bandwidth

It also shards the verification of signatures so that every participant does not need to verify every transaction.

# Theorem

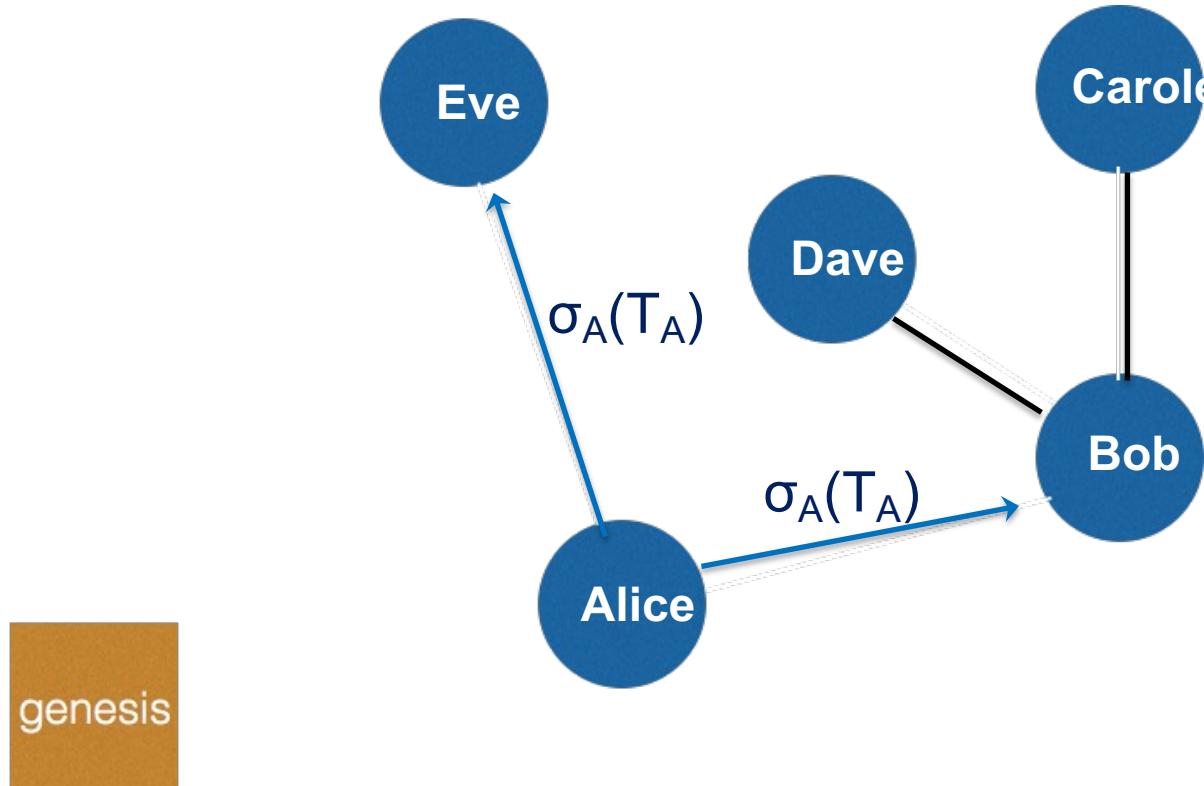
There is no solution to the Byzantine consensus if  $f=1$  and  $n=3$ .

## Proof.

- Consider execution E1 where p1 and p3 are correct and p2 Byzantine. Process p1 proposes 1 to p2 and p3, and p3 reports back to p2 but p3 says that p1 proposed 0. As p1 and p3 are correct, it follows that we need to decide 1 (due to validity).
- Consider a similar execution E2 but where p1 is Byzantine and sends different values to p2 and p3 that are correct. We need to decide 1 (due to agreement).
- Consider E3 where p1 sends 0 and is correct while p3 reports 1 back to p2. We need to decide 0 (due to validity).
- At the end of the execution, p2 cannot distinguish E3 from E2 yet it should decide differently.

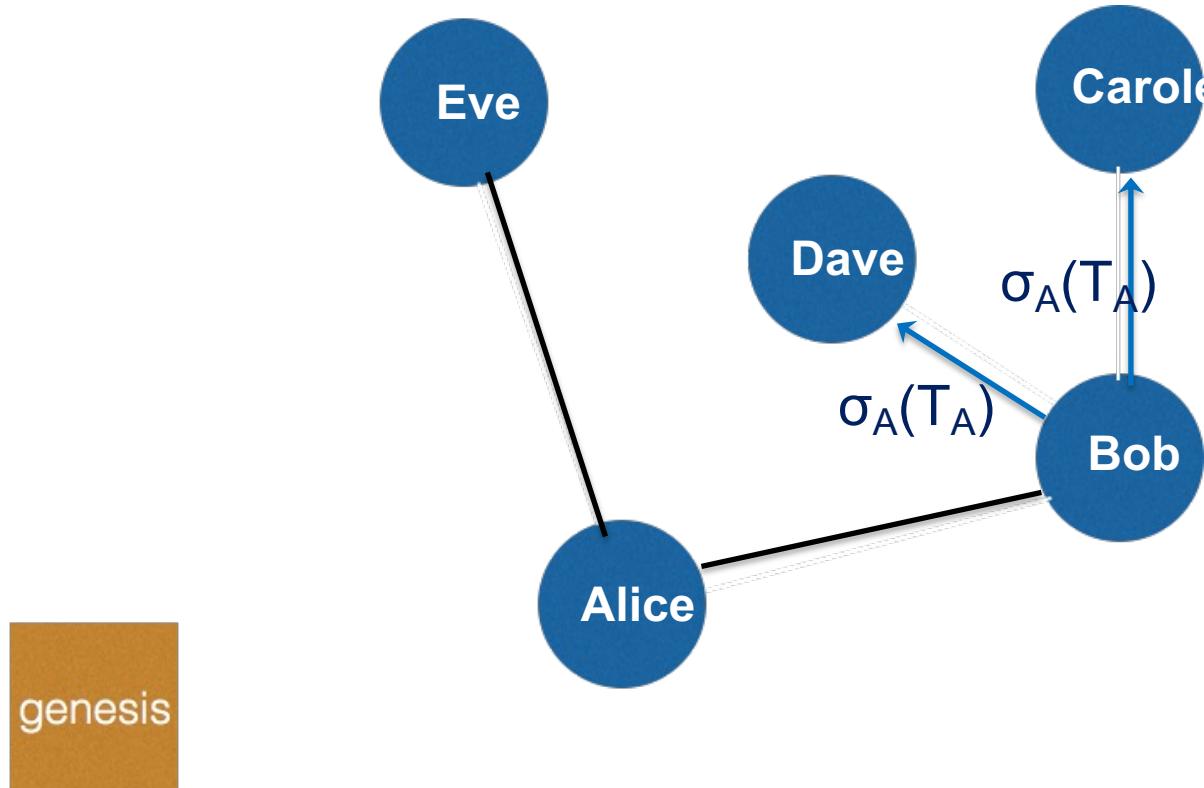
# Smart Redbelly Blockchain

# Transactions are validated twice



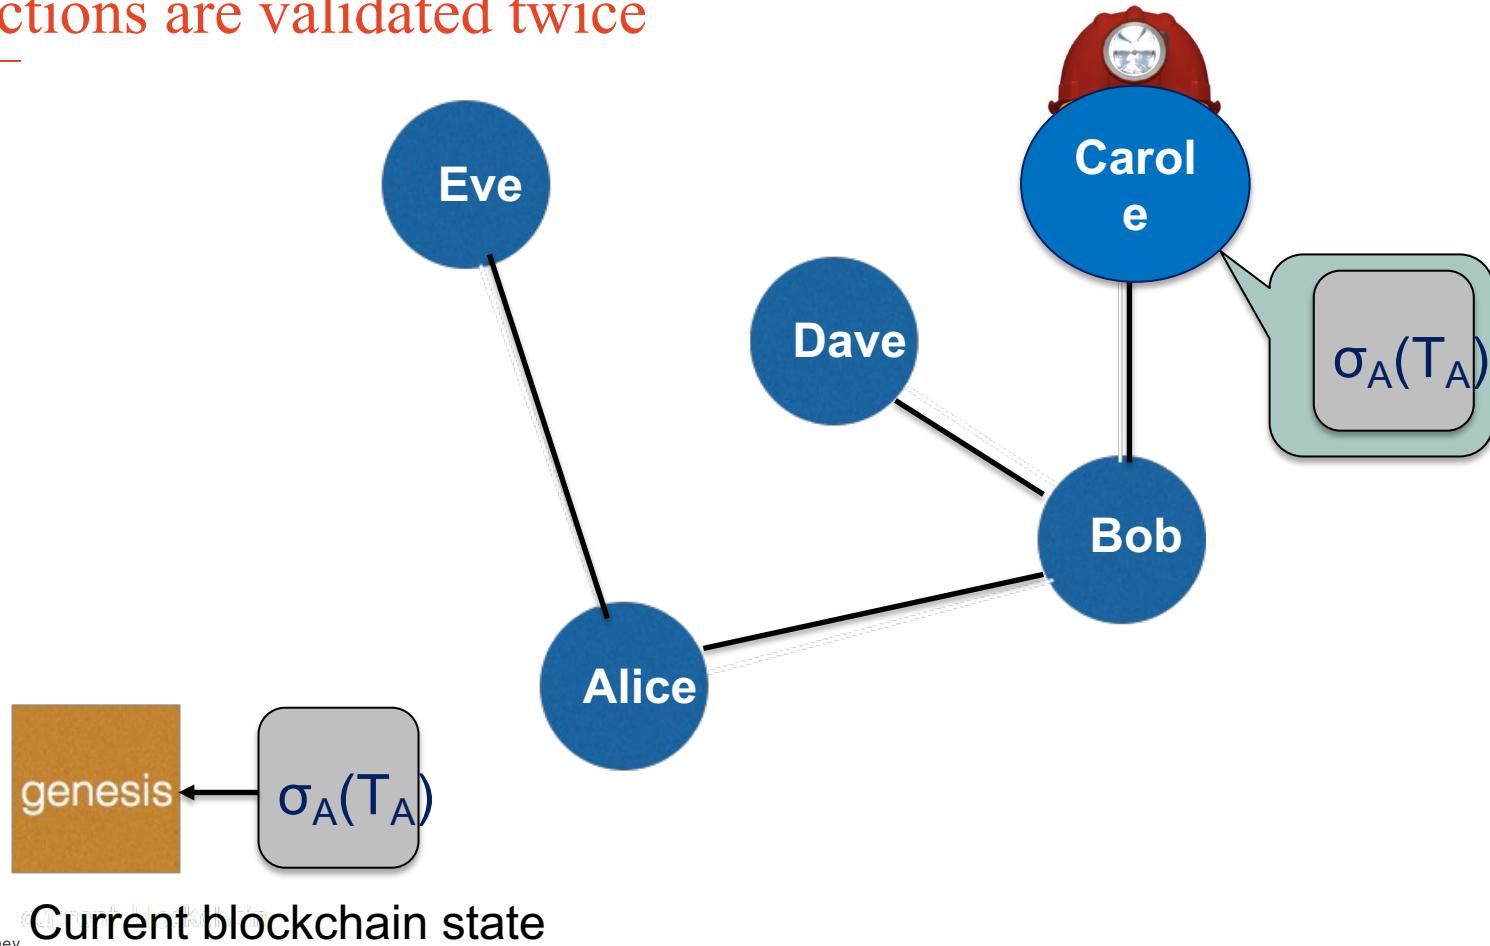
Current blockchain state

# Transactions are validated twice

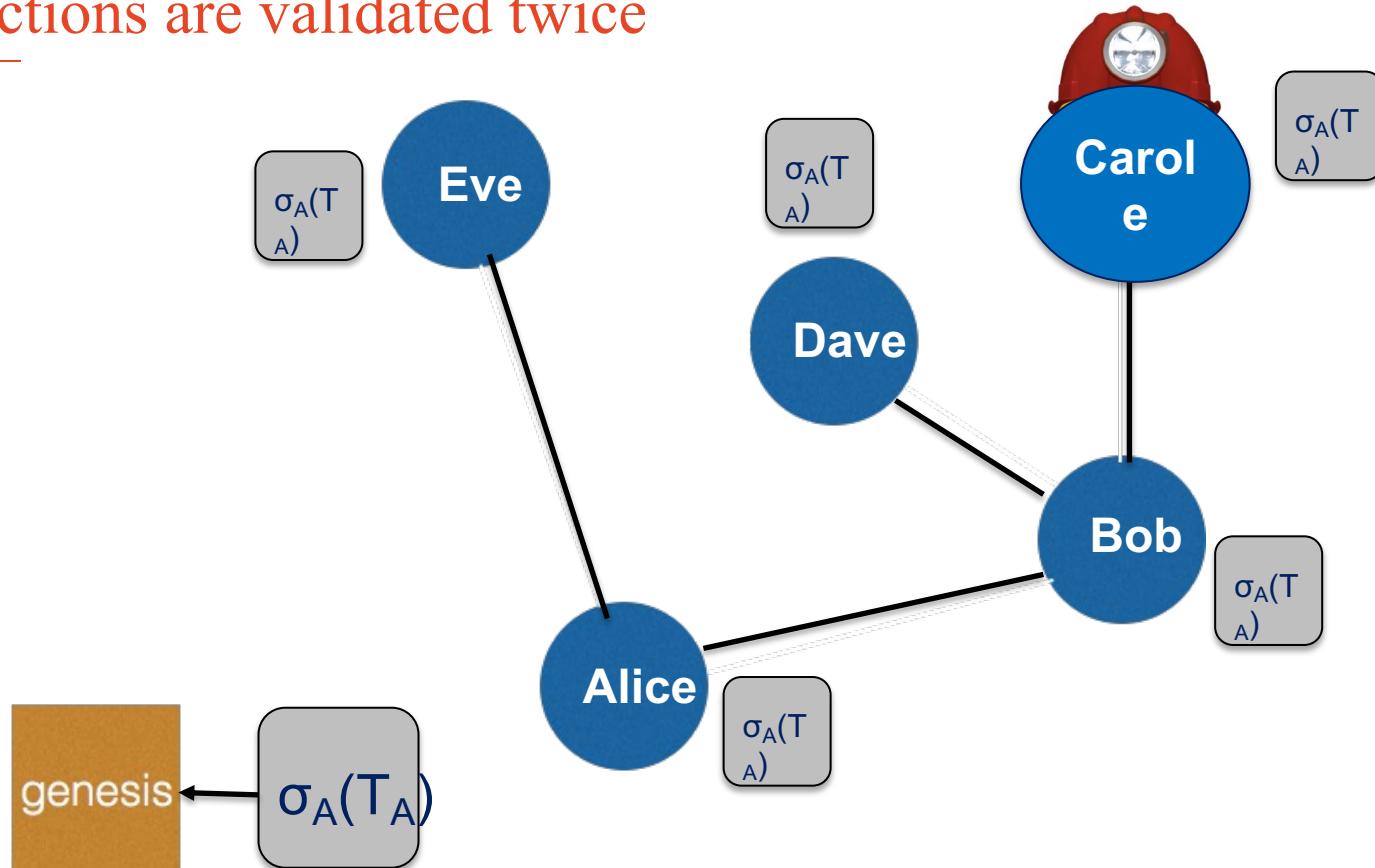


Current blockchain state

# Transactions are validated twice

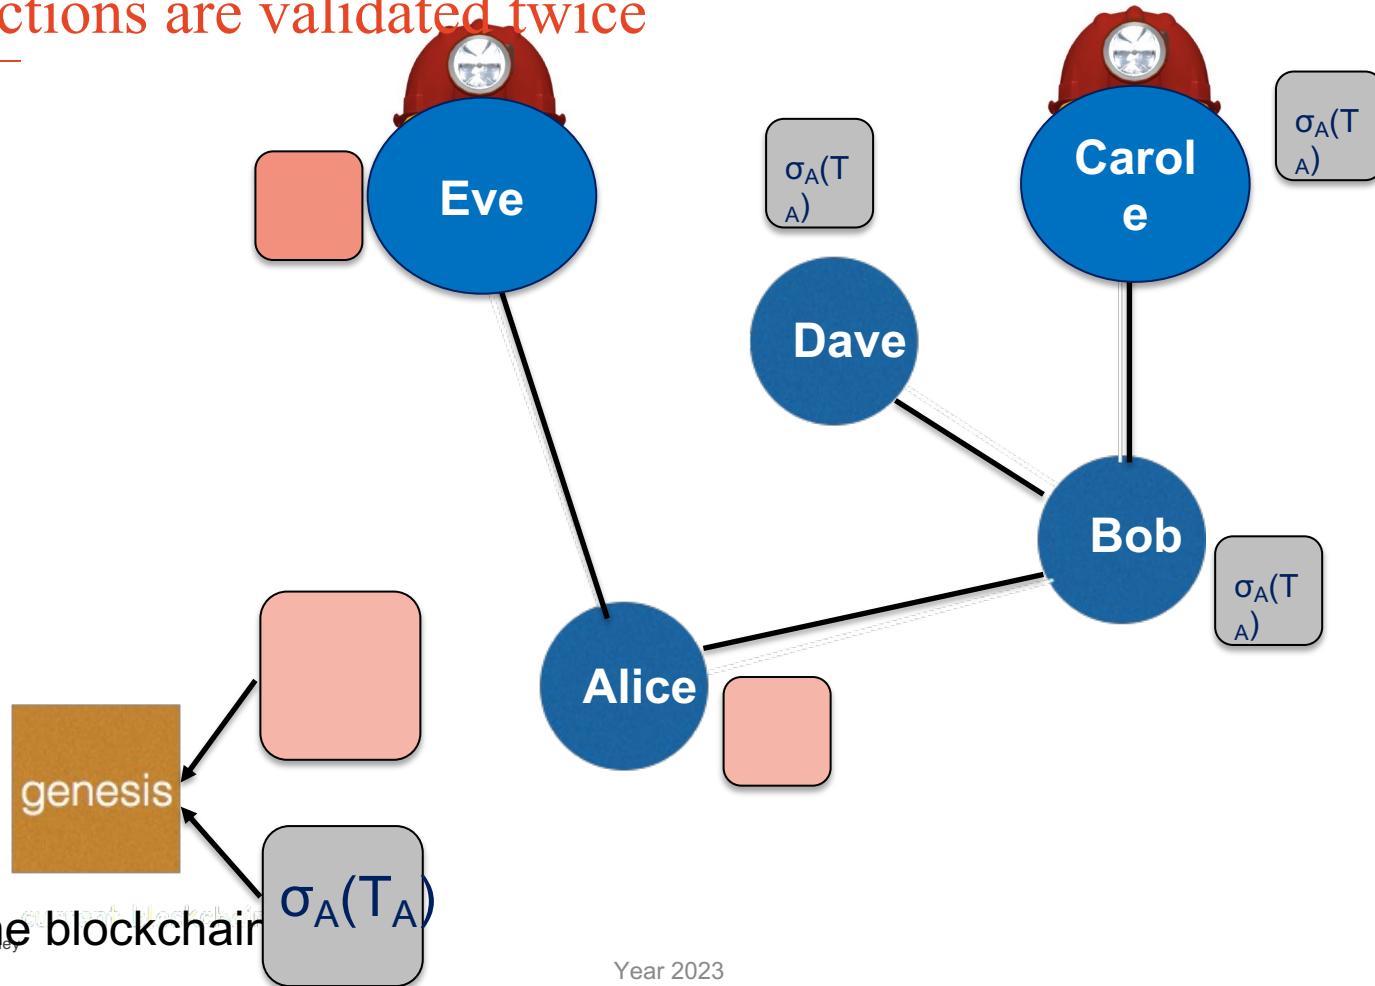


# Transactions are validated twice

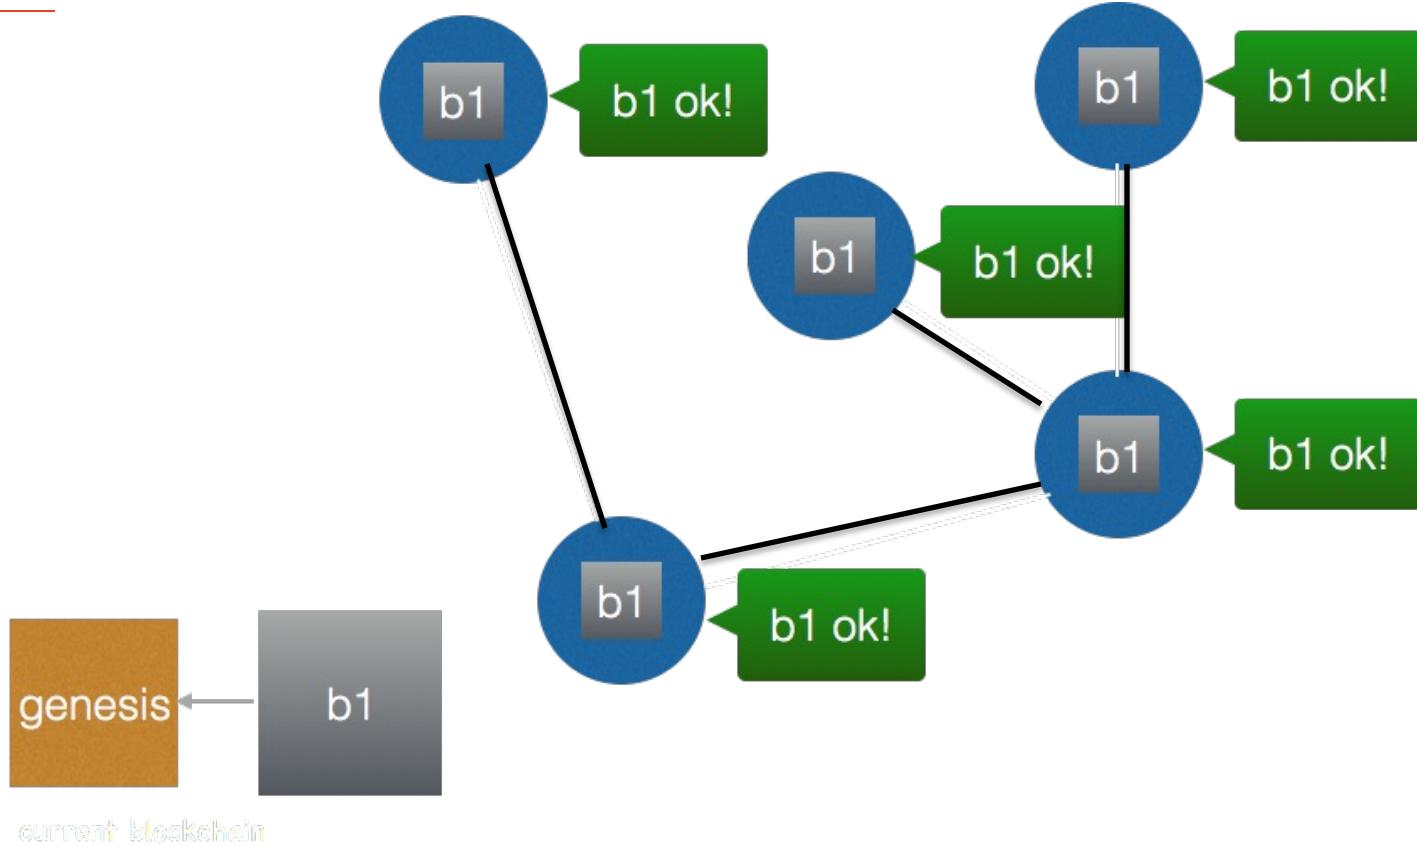


Current blockchain state

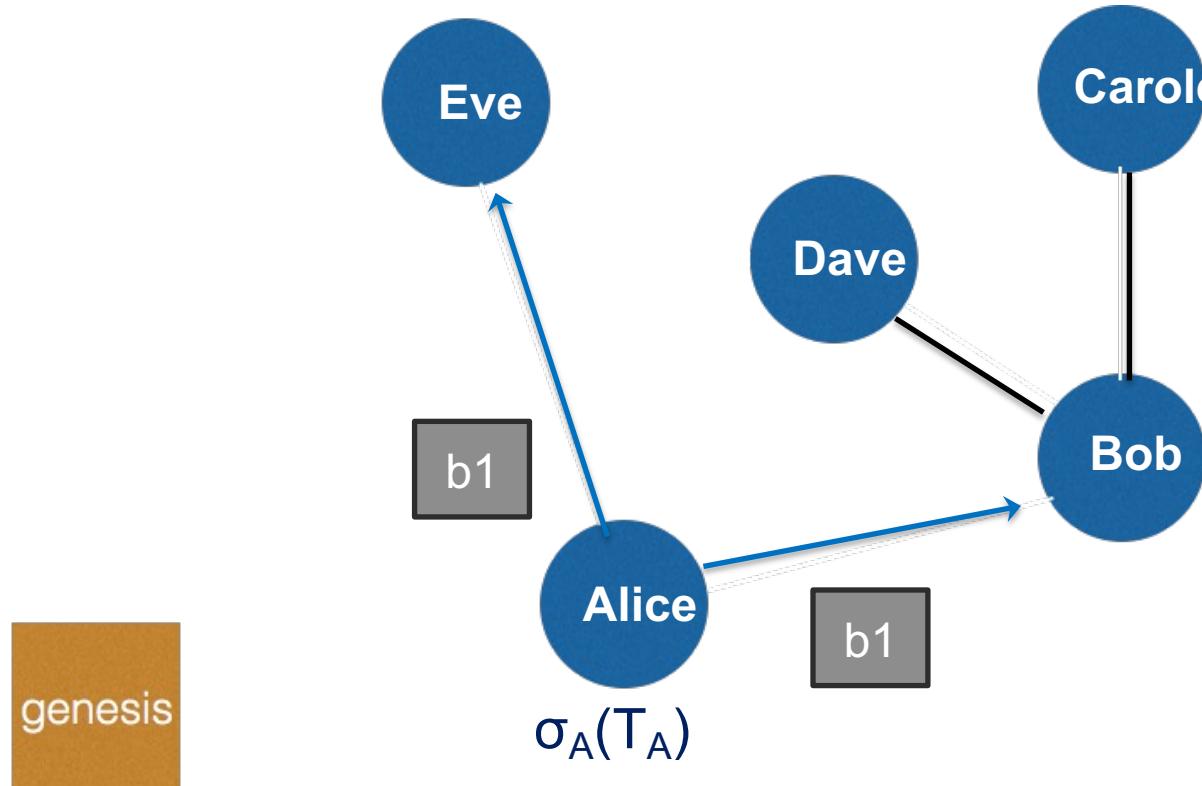
# Transactions are validated twice



# Transactions are validated twice

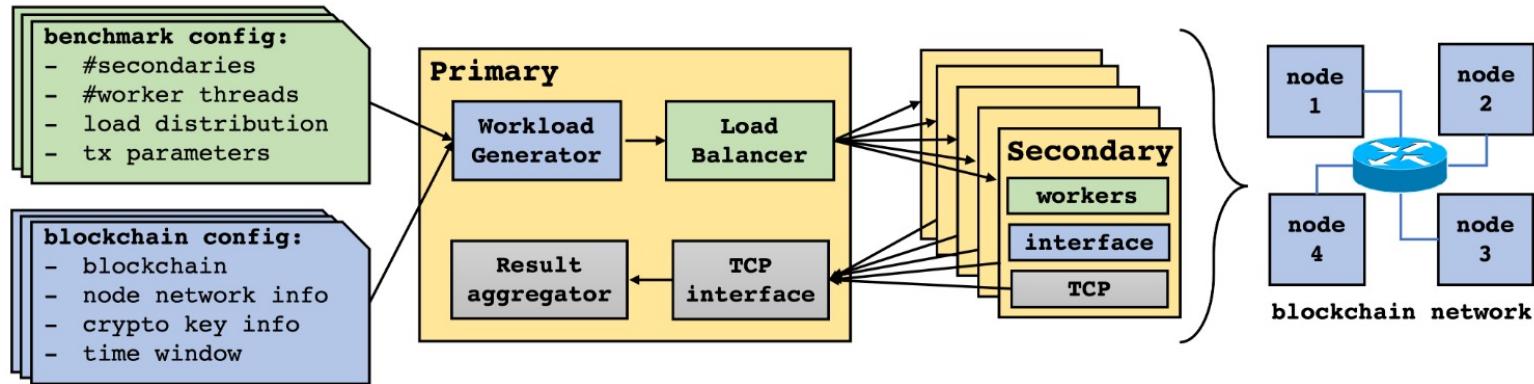


# Smart Redbelly Blockchain validates less

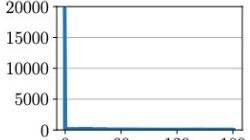
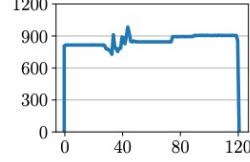
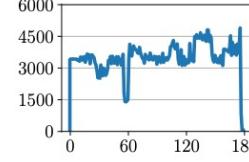
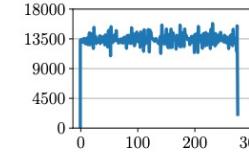
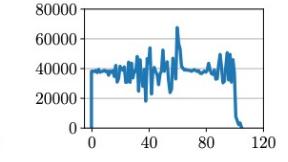


Current blockchain state

# Diablo benchmarking framework



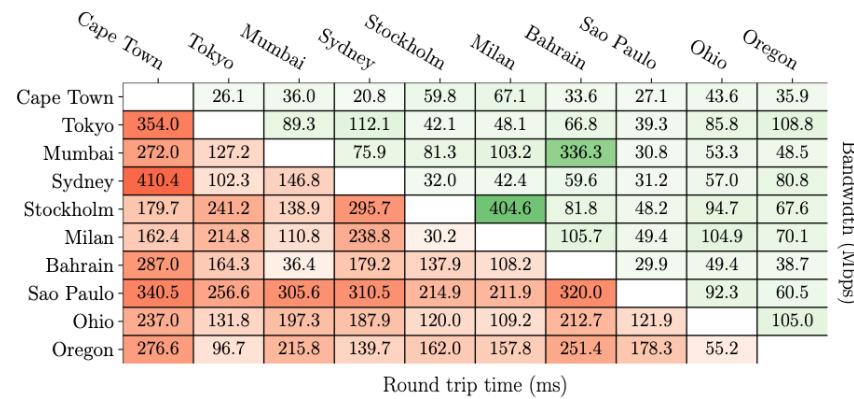
# Diablo benchmarking framework

DApp	Exchange	Mobility service	Web service	Gaming	Video sharing
Workload					
Source trace Characteristics	NASDAQ Burst	Uber Compute intensive	FIFA Contended	Dota 2 High sending rate	YouTube Very high sending rate

**Table 2.** Decentralized applications (DApps) used as DIABLO benchmarks and their associated workload based on real traces. Each graph shows the number of submitted transactions (y-axis) per second (x-axis).

# Diablo benchmarking framework

Configuration	Blockchain nodes			Regions
	number	#vCPUs	memory	
datacenter	10	36	72 GiB	Ohio
testnet	10	4	8 GiB	Ohio
devnet	10	4	8 GiB	all
community	200	4	8 GiB	all
consortium	200	8	16 GiB	all



# Diablo benchmarking framework

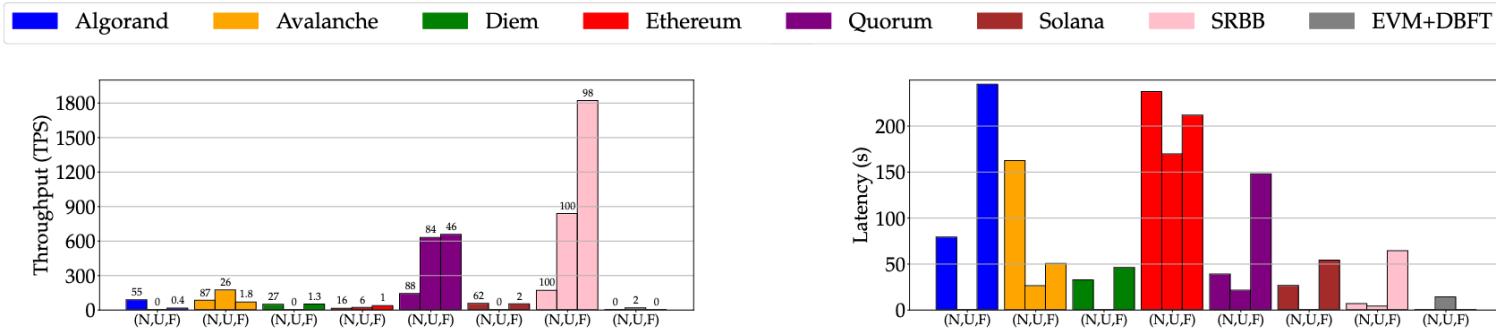


Fig. 2: Throughput (y-axis) and commit percentage (top of the bar) for NASDAQ, Uber and FIFA workloads (i.e., (N,U,F) is NASDAQ, Uber and FIFA)

Fig. 3: Latency (y-axis) for NASDAQ, Uber and FIFA workloads (i.e., (N,U,F) is NASDAQ, Uber and FIFA)

**VINCENT GRAMOLI** A/Prof.  
THE UNIVERSITY OF SYDNEY  
Rm 417, Computer Science Building | The University of Sydney | NSW | 2006  
+61 2 9036 9270

<https://gramoli.github.io>

 <https://twitter.com/VincentGramoli>

 <https://www.linkedin.com/in/vincent-gramoli-206148a/>

 <https://scholar.google.com.au/citations?user=Fv7ZfN8AAAAJ>



THE UNIVERSITY OF  
**SYDNEY**

# Backup

---

# A guide to using our University PowerPoint template

**Please read the following 3 slides, prior to building your PowerPoint presentation.**

Within this presentation there are both sample presentation slides AND a full range of templates built into the master slides for your use. There is a range of traditional bullet pointed slides, title slides and chapter break slide templates, and also graphic slides to help you to make your presentation less text heavy and more visually engaging.

It is recommended that upon designing your PowerPoint presentation you:

- Start by writing out your talking points.
- Utilise presenter notes. It's important not to read off the slide and speak directly to your audience throughout your PowerPoint presentation.
- Minimise the amount of text on each slide, and use images, icons and simple data/statistics to engage your audience.
- Keep your design consistent throughout. Read the notes on the following 2 slides detailing colour and font style.
- Delete the slides from the presentation that you don't need, and add slides that you do need by going to: Home > New Slide > Drop down arrow

## Design and layout

In the following presentation we have included examples of presentation slides.

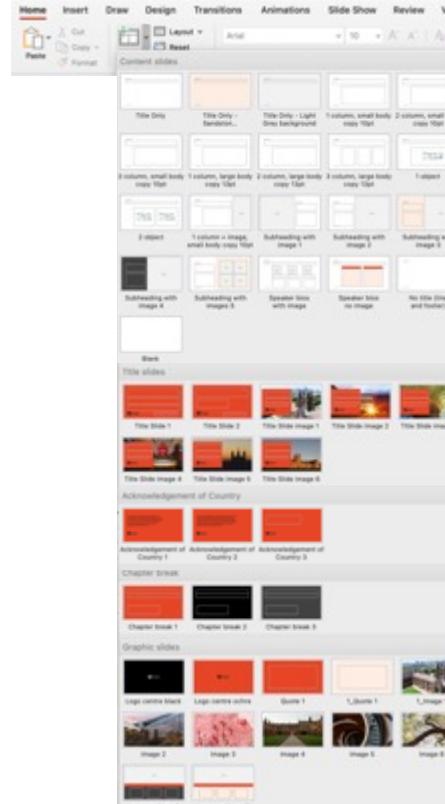
There are many more slide templates available.

These can be viewed and added to your presentation by going to:

[Home > New Slide > Drop down arrow](#)

Templates have been grouped into categories:

- Title slides
- Acknowledgement of Country
- Chapter break slides
- Graphic slides
- Content slides



## Slide transitions and animation

Our slide transition style is to keep things simple. Our preference is no transition style (ie 'None') as transitions can distract your audience from you and your content. Simple animation (ie 'Appear') can add value when used to build the story telling of your presentation, ie Step 1, Step 2 etc.

## Colours

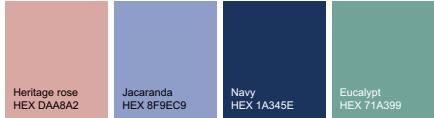
Primary colour palette



Secondary colour palette



Tertiary colour palette



These colours are preloaded into this PowerPoint colour palette.

**Tertiary colours are only for use in graphs and tables.** Only use the Tertiary colour palette after the Primary and Secondary colours have already been used on the same slide.

You can use tones of the colour palette.

## Graphs and table

Our style is to keep graphs and tables simple with the use of the solid colours available in the colour palette, or mono tones of the Ochre colour. Do not use shading or patterns.

## Fonts

The University primary fonts for use in PowerPoint are:

Arial Regular

*Plus a dash of Times Italic*

All body copy is Charcoal colour (not black).  
Do not use any other fonts.

## Bullet points

Our bullet point style is to use a 'dash' as a bullet point, rather than the tradition solid circle. This can be change by:

- Format > Bullet and numbering > Customise > Add in '-' > Press 'OK'.

## Logo

If you need to use the University logo, please copy the appropriate version from here:

Standard



THE UNIVERSITY OF  
SYDNEY

Reversed



USER NOTE:

There are more Title Slide options available to add to your presentation by going to:  
Home > New Slide > Drop down arrow

# Heading sits here

## 1-2 lines of copy 28pt

*Presentation subheading  
can go over 2 lines 28pt*

Presented by 14pt



THE UNIVERSITY OF  
**SYDNEY**

#### USER NOTE:

- To add your own image to this Title Slide
1. Select the sky background image
  2. Right click (or Control click for Mac users)
  3. Select 'Change image'
  4. Navigate to the picture you want, select it, and then select Insert

# Heading sits here 28pt 1-2 lines

Intro text 14pt, to sit here.

2–3 lines of copy. Am quo'xnkaln ium aute  
fuidestel eces, Cas cus aus opubli.

[sydney.edu.au/xxxxx](http://sydney.edu.au/xxxxx)



THE UNIVERSITY OF  
**SYDNEY**

USER NOTE:

It is mandatory to include an 'Acknowledgement of Country' in your presentation. Please include this slide in your presentation.

We recognise and pay respect to the Elders and communities – past, present, and emerging – of the lands that the University of Sydney's campuses stand on. For thousands of years they have shared and exchanged knowledges across innumerable generations for the benefit of all.

USER NOTE:  
It is optional to include  
this additional slide.

The University of Sydney's Camperdown Campus sits on the lands of the Gadigal people with campuses, teaching and research facilities on the lands of the Gamaraygal, Dharug, Wangal, Darkinyung, Burramadagal, Dharawal, Gandangara, GAMILARAAY, Barkindji, Bundjalung, Wiradjuri, Ngunawal, Gureng Gureng, and Gagadju peoples.



USER NOTE:

Alternatively you can write your own  
'Acknowledgement of Country', specific to  
the land you are presenting from.



THE UNIVERSITY OF  
**SYDNEY**

USER NOTE:

There are more Chapter Break Slides options available to add to your presentation by going to:  
Home > New Slide > Drop down arrow

# Chapter break headline 28pt

*1-2 lines of copy*

Presented by 14pt

USER NOTE:

There are more Graphic Slide  
image options available to add to your  
presentation by going to:  
Home > New Slide > Drop down arrow

A better place to work  
*and a place that works better*



USER NOTE:

There are more Content Slide options available to add to your presentation by going to:  
Home > New Slide > Drop down arrow

---

# Agenda or program

## Add date of presentation

- Morning session
  - Details of morning session 1
  - Details of morning session 2
  - Details of morning session 3
- Lunch break
- Afternoon session
  - Details of afternoon session 1
  - Details of afternoon session 2
  - Details of afternoon session 3
- Close

---

# Title and 2 column text box

## Small body copy

### **Subheading, Arial 10pt**

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur? Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus, ut facessitae liciam faceptat re, nos nos etus ut modicia valorrovid quatemp osaperu ptatur sam, corum que evenimolupi eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut.

Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur? Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus, ut facessitae liciam faceptat re, nos nos etus ut modicia valorrovid quatemp osaperu ptatur sam, corum que evenimolupi:

Add bullet point. Eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut faccull aturepe llesequo tem et mo offic totatem cum fugit quodio volor apitibusam expla quam dis aceatio samustism quatiistio. Namuscia venes vellaces aut omnient officab orepell oriam, eaquunt.

Add bullet point. Eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut faccull aturepe llesequo tem et mo offic totatem cum fugit quodio volor apitibusam expla quam dis aceatio samustism quatiistio. Namuscia venes vellaces aut omnient officab orepell oriam, eaquunt.

---

# Title and 3 column text box

## Small body copy

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur?

Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus, ut facessitiae liciam faceptat re, nos nos etus ut modicia valorrovid quatemp osaperu ptatur sam, corum que evenimolupi eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut.

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur?

Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus, ut facessitiae liciam faceptat re, nos nos etus ut modicia valorrovid quatemp osaperu ptatur sam, corum que evenimolupi eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut.

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur?

Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus,.

### Research implication

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie

Small body copy 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitae pra volupie ndempor atis quias acilit latemperibus idesti non entur?

Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus.

---

# Title and 3 column text box

## Large body copy

### **Subheading large 13pt**

Body copy large 13pt.  
Body copy text to go here,  
herelmaximus explab iunt  
qui assit quam ut facimpori  
rest aliique odit, solut quidi  
te voluptam iduscium  
rerem ut que ratatur  
simusae, solut quidi te  
voluptam iduscium rerem  
ut que ratatur simusae.

### **Subheading large 13pt**

Body copy large 13pt.  
Body copy text to go here,  
herelmaximus explab iunt  
qui assit quam ut facimpori  
rest aliique odit, solut quidi  
te voluptam iduscium  
rerem ut que ratatur  
simusae, solut quidi te  
voluptam iduscium rerem  
ut que ratatur simusae.

### **Subheading large 13pt**

Body copy large 13pt.  
Body copy text to go here,  
herelmaximus explab iunt  
qui assit quam ut facimpori  
rest aliique odit, solut quidi  
te voluptam iduscium  
rerem ut que ratatur  
simusae, solut quidi te  
voluptam iduscium rerem  
ut que ratatur simusae.

“Insert quote here.  
Unti cuptasitae velluptatur suntem as  
derempo rposten dignisqui omnissum  
eiur, earchit ab ipidici llorum,  
coreriae quia.”

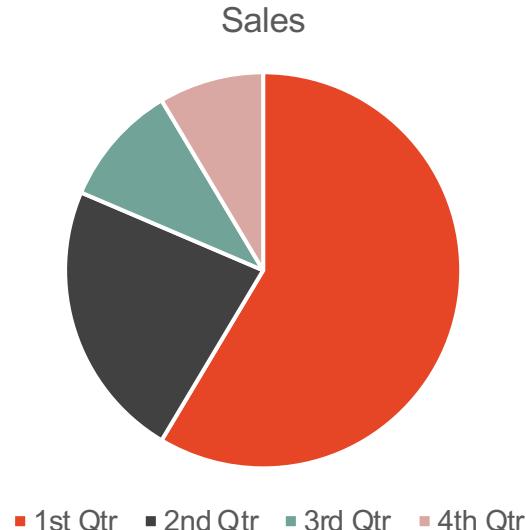
USER NOTE:  
When adding in a chart, use the default  
chart style 'Style 1'.

---

# Title and 2 objects

## Full colour graph

- When adding in a chart, use the default chart style 'Style 1'.
- This will colour your graph using the University full colour themes
- Small body copy 10pt. Uas estrum fuga.  
Et fugia cusaepernam, siminto.  
As sequis eatibus. Dolesti oresequae  
eat ressitiae pra volupie ndempor atis quias acili



---

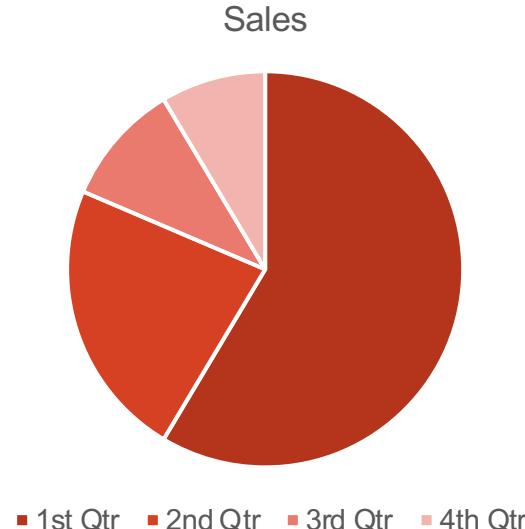
#### USER NOTE:

When adding in a chart, use the default chart style 'Style 1'. Go to 'Change colours' (to the left of the graph styles) and choose Monochromatic.

## Title and 2 objects

### Mono colour graph

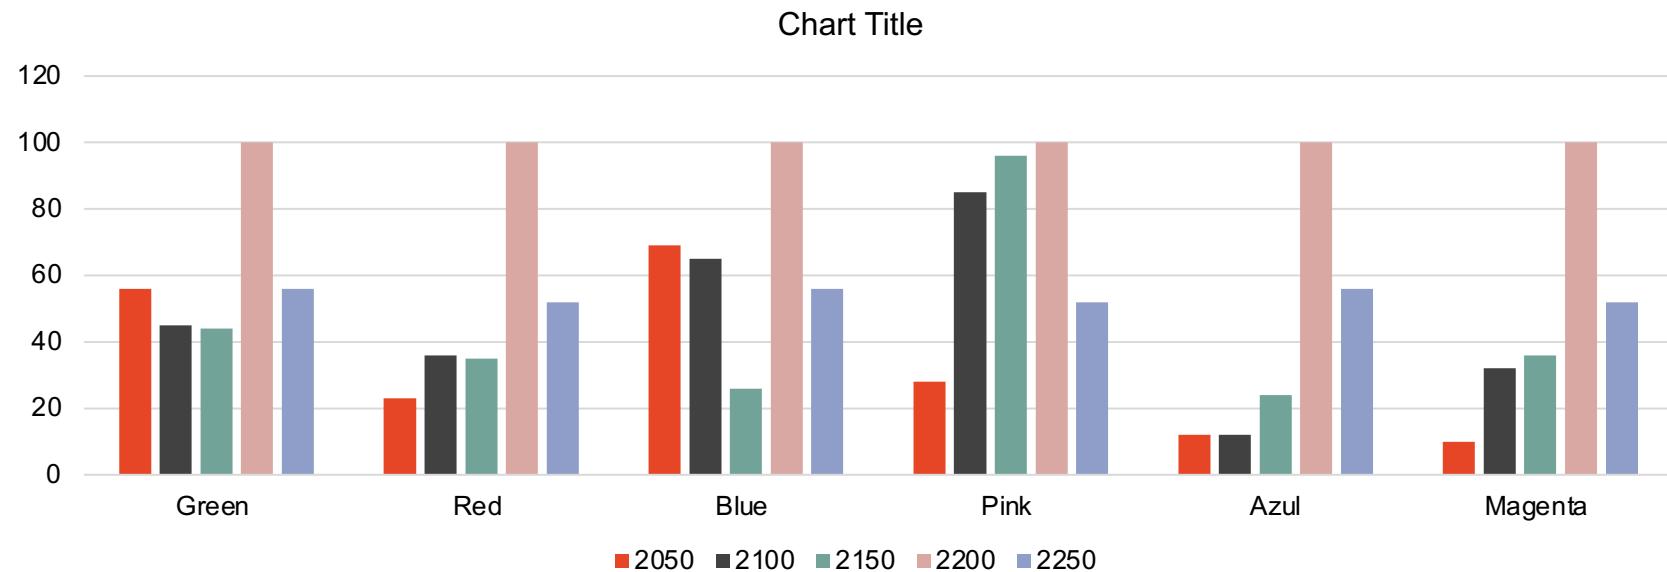
- When adding in a chart, use the default chart style 'Style 1'.
- This will colour your graph using the University full colour themes
- Alternatively, you can go to 'Change colours' (to the left of the graph styles) and choose Monochromatic.
- Small body copy 10pt. Uas estrum fuga.  
Et fugia cusaepernam, siminto.  
As sequis eatibus. Dolesti oresequae  
eat ressitiae pra volupie ndempor atis quias acili.



USER NOTE:  
When adding in a chart, use the default  
chart style 'Style 1'.

# Title and 1 object

## Full colour graph

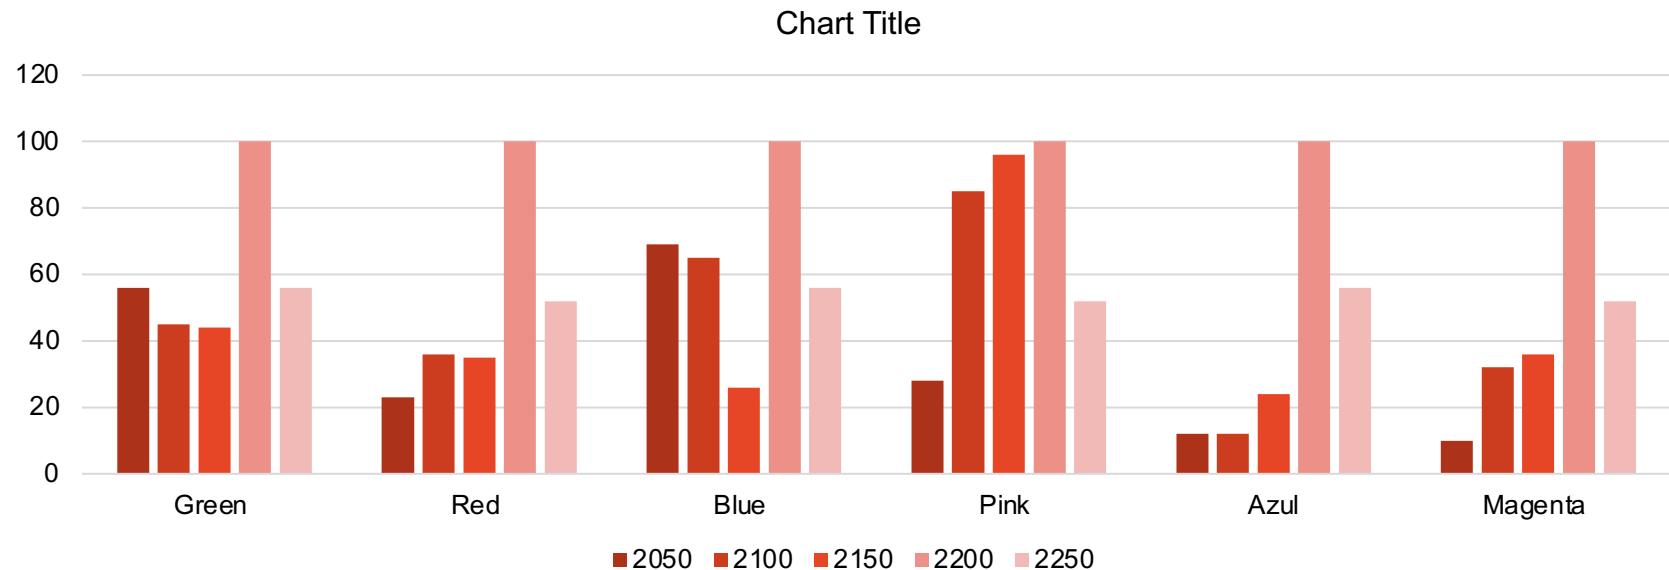


---

## USER NOTE:

When adding in a chart, use the default chart style 'Style 1'. Go to 'Change colours' (to the left of the graph styles) and choose Monochromatic.

# Title and 1 object Mono colour graph





Body copy 10pt. Ribusciist ut qui dire est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Etfugia cusaepernam, siminto eatur sae.

Body copy 10pt. Ribusciist ut qui dire est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Etfugia cusaepernam, siminto eatur sae.

Body copy 10pt. Ribusciist ut qui dire est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Etfugia cusaepernam, siminto eatur sae.

# Welcome Week (Title + 4 images)



Left to right: Welcome Week 2022, flags on Eastern Avenue, aerial view of Welcome Week activities, international students arriving at Sydney airport.



---

## Our student-focus education is *transformational*

We will be the best Australian university for teaching and learning. Our students will attest to the transformational impact a Sydney education has on their lives, whenever and wherever they learn

---

## *Our community thrives through diversity*

Social equity is at the heart of our past and our future.  
By 2032 our community will be more diverse and  
inclusive, helping us all to shape a positive future.



---

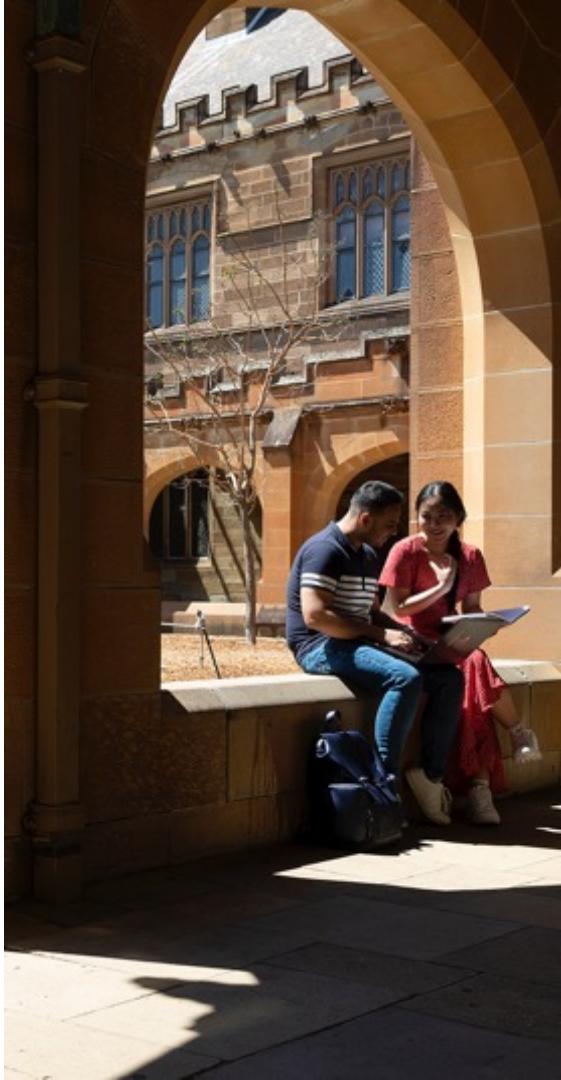
# Title, text box and image

## Small body copy

Body copy text to go here 10pt. Uas estrum fuga. Et fugia cusaepernam, siminto eAs sequis eatibus. Dolesti oresequae eat ressitaе pra volupie ndempor

Iatis quias acilit latemperibus idesti non entur? Ritatem sincte rerum voluptincto quam qui sit, que vid qui a senderi quam, ide nobis acea sus, ut faccessitiae liciam faceptat re, nos nos etus ut modicia volorrovid quatemp osaperu ptatur sam, corum que evenimolupid eos dempos is pa pedisci bla dolessimus quatur, ut quia nobis aut escit ant ut faccull aturepe llesequo tem et mo offic totatem cum fugit quodio valor, apitibusam expla quam dis aceatio samustisim quatiistio. Namuscia venes vellaces aut omnient officab orepell oriam, eaquunt. Ro et aut eos atiae. Ut videlitae nonestiorrum ipsae cuptamet lantion sequae sim quis dem evel inimi, veles aut re volupta.

Inctet et, si blanturiam quassitemped quo ius.Otatem. Orestia es as eos earchil lorepudi berio bearior uptatis ma nos quis abor maionse aut est ut et, que plam in core velitio coreperat velende nis venisi alitatem.



---

## Subheading 4

### Title, text box and image

#### Large body copy

Body copy large 13pt. Ribusciist ut qui dire est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is dolumquas estrum fuga. Et fugia cusaepernam, siminto eatur sae.



---

# Speaker bios with image



**First Last Name 10pt**

*Title times 10pt*

Body copy small 10pt. Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Et fugia cusaepernam.



**First Last Name 10pt**

*Title times 10pt*

Body copy small 10pt. Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Et fugia cusaepernam.



**First Last Name 10pt**

*Title times 10pt*

Body copy small 10pt. Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rfernam ius, nume aut dis dolorro to magnitates et aut quam volorro et que num is columquas estrum fuga. Et fugia cusaepernam.

---

# Icon style



Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernam ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.



Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernam ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.



Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernam ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.

---

# Infographic data

100

Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.

95

Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.

23

Body copy small 10pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius, nume aut dis  
dolorro to magnitates et aut  
quam volorro et que.

#### USER NOTE:

This layout uses a custom table.

You can edit this table to keep this style.

Alternatively use the table style

'Medium Style 1 – Accent 1" as it is most similar

# Table layout option

## Text can run over second line

*“Large quote or key finding could go here. 13pt Times italic. Genihillupti is ut quatiam, sunt remquo inciis quiatistrum facea eosamet viti ommosandae. Et vid quas ventibus nosto bearum et antur? Quia sim sitae restem voluptibus a vendes debis volorest ma samusci llaccaest, unio molest, etur, veniendam architati.”*



Interesting data detail, 7pt.  
Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rferniam ius.



Interesting data detail, 7pt.  
Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rferniam ius.



Interesting data detail, 7pt.  
Ribusciist ut qui di re est eostorem is dit iuremolut ipis aut molupta quam, sintium quidem hicabor emperfe rferniam ius.

Facebook	Twitter	Pinterest	YouTube	LinkedIn	Outcome
Text here	Text here	Text here	Date	Date	
Text here	Text here	Text here	Number views	Number views	
Text here	Text here	Text here	Text here	Text here	

Facebook	Twitter	Pinterest	YouTube	LinkedIn	Outcome
Text here	Text here	Text here	Date	Date	
Text here	Text here	Text here	Number views	Number views	
Text here	Text here	Text here	Text here	Text here	

USER NOTE:

This layout uses a custom table.

You can edit this table to keep this style.

Alternatively use the table style  
'Medium Style 1 – Accent 1" as it is most  
similar

# Table layout option

## Text can run over second line

### Key highlights



Interesting data detail, 7pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius.



Interesting data detail, 7pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius.



Interesting data detail, 7pt.  
Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis  
aut molupta quam, sintium  
quidem hicabor emperfe  
rfernам ius.

Facebook	Twitter	Pinterest	YouTube	LinkedIn	Outcome
Text here	Text here	Text here	Date	Date	
Text here	Text here	Text here	Number views	Number views	
Text here	Text here	Text here	Text here	Text here	
Text here	Text here	Text here	Text here	Text here	

*“Large quote or key finding could go here. 13pt Time Italic.  
Genihillupti is ut quatiam, sunt remquo inciis quiatistrum facea  
eosamet viti ommosandae. Et vid quas ventibus nosto bearum et  
antur? Quia sim sitae restem voluptibus vendes debis volorest  
ma samusci llaccaest, untio molest, etur, veniendam architati.”*

USER NOTE:

This layout uses a custom table.

You can edit this table to keep this style.

Alternatively use the table style

'Medium Style 1 – Accent 1" as it is most similar

# Table layout option

## Text can run over second line

Facebook	Twitter	Pinterest	YouTube	LinkedIn	Outcome
Text here	Text here	Text here	Date	Date	
Text here	Text here	Text here	Number views	Number views	
Text here	Text here	Text here	Text here	Text here	
Text here	Text here	Text here	Text here	Text here	
Text here	Text here	Text here	Text here	Text here	
Text here	Text here	Text here	Number views	Number views	
Text here	Text here	Text here	Text here	Text here	
Text here	Text here	Text here	Text here	Text here	

USER NOTE:  
This layout uses a custom table.  
You can edit this table to keep this style.

# Example of timeline

Steps	January	February	March	April
Detail 1				
Detail 2				
Detail 3				
Detail 4				
Detail 5				
Detail 6				

---

USER NOTE:

This layout uses a custom graphic.  
You can edit this graphic to keep this style.

# Flow chart style 1



**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.



**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.



**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.



**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.



**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

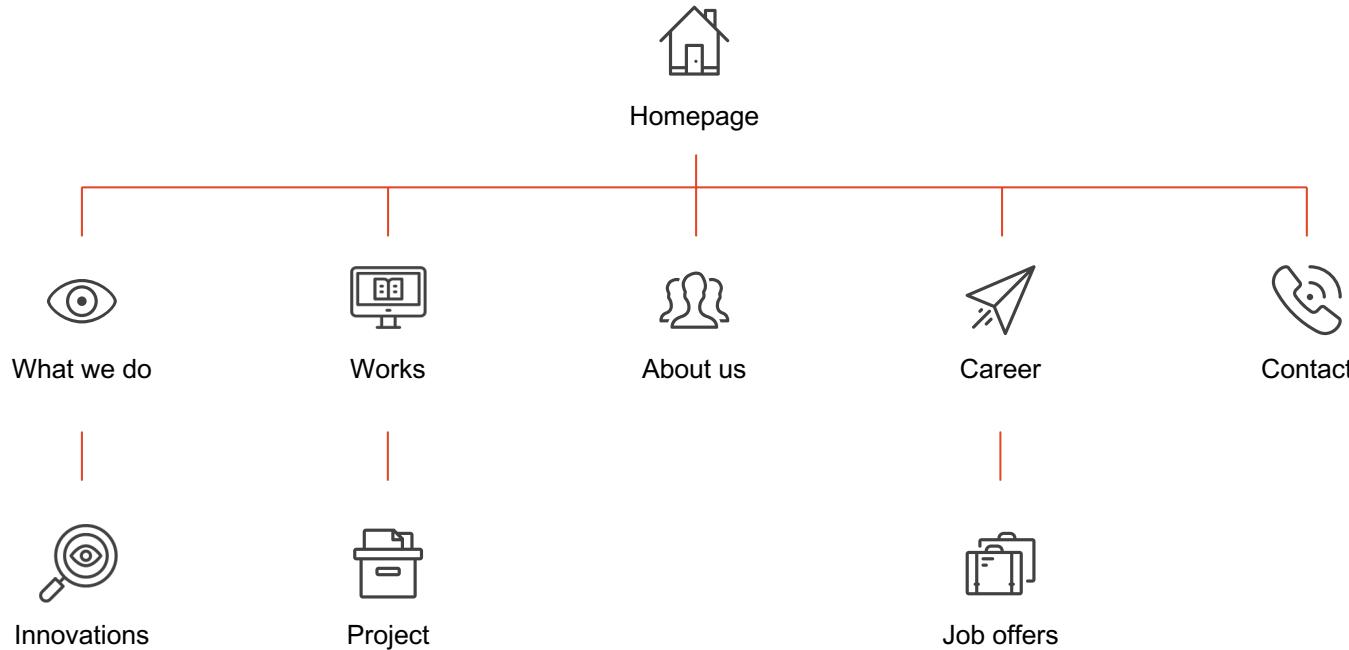
100

Body copy small 10pt. Ribusciist ut qui di re est  
eostorem is dit iuremolut ipis aut molupta quam,  
sintium quidem hicabor emperfe rfernam ius.

USER NOTE:

This layout uses a custom graphic.  
You can edit this graphic to keep this style.

# Flow chart style 2



USER NOTE:  
This layout uses a custom graphic.  
You can edit this graphic to keep this style.

# Flow chart style 3



Text 1

Text 2

Text 3

Text 4

Text 5

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

---

## USER NOTE:

This layout uses a custom graphic.  
You can edit this graphic to keep this style.

# Flow chart style 4



---

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

**CONTENT TITLE**  
Add body copy 10pt.  
Add body copy 10pt.

USER NOTE:

Icons can be changed to white  
using 'Shape fill' for use on Ochre  
or Charcoal background

# Icon style

## Communication



## Education



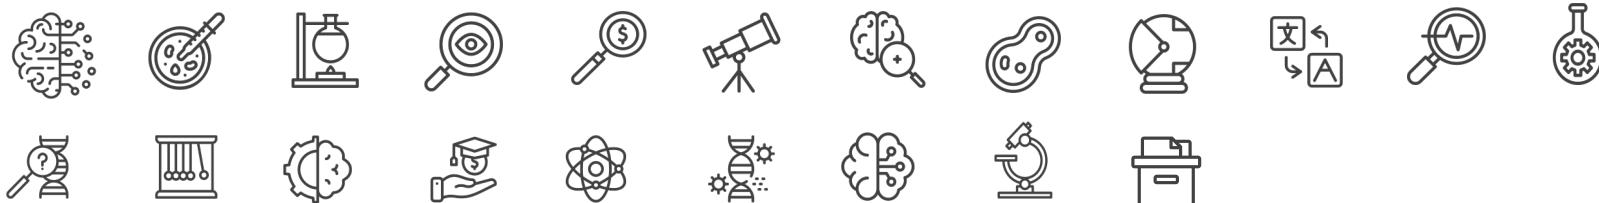
USER NOTE:

Icons can be changed to white  
using 'Shape fill' for use on Ochre  
or Charcoal background

## Marketing



## Research



## Other icons

Our icon style is to use a thin line style of icon. Additional icons can be sourced at:

<https://www.streamlinehq.com/icons/streamline-mini-line>

<https://thenounproject.com/pricing/#icons>

USER NOTE:

Delete the social media icons  
below that you do not require and  
align the icons you need neatly

**Insert team/discipline area**

Name Surname Title

+61 2 1234 5678

**[sydney.edu.au/xxxx](#)**

